
HIP Documentation

Release 6.3.42134

Advanced Micro Devices, Inc.

Mar 13, 2025

INSTALL

1	Overview	3
2	Install HIP	5
2.1	Prerequisites	5
2.2	Installation	5
2.3	Verify your installation	6
3	Build HIP from source	7
3.1	Prerequisites	7
3.2	Building the HIP runtime	7
3.3	Build HIP tests	10
3.4	Run HIP	11
4	Understanding the HIP programming model	13
4.1	RDNA & CDNA Architecture Summary	13
4.2	Single Instruction Multiple Threads	15
4.3	Heterogeneous Programming	15
5	Programming model reference	17
5.1	Threading Model	17
5.1.1	Inherent Thread Model	17
5.1.2	Cooperative Groups Thread Model	18
5.2	Memory Model	18
5.3	Execution Model	19
5.3.1	Host-side execution	19
5.3.2	Device-side execution	20
5.3.3	Kernel launch	20
6	Hardware Implementation	21
6.1	Compute Units	21
6.1.1	SIMD	22
6.1.2	Vector Cache	22
6.1.3	Local Data Share	22
6.1.4	Scalar Unit	22
6.2	CDNA Architecture	22
6.3	RDNA Architecture	23
6.4	Shader Engines	23
7	AMD compute language runtime (CLR)	25
7.1	Project Organisation	25
7.2	How to build/install	25

7.2.1	Prerequisites	25
7.2.2	Linux	25
7.2.3	Test	26
7.2.4	Release notes	26
8	HIP Programming Manual	27
8.1	Host Memory	27
8.1.1	Introduction	27
8.1.2	Memory allocation flags	27
8.1.3	Numa-aware host memory allocation	28
8.1.4	Coherency Controls	28
8.1.5	Visibility of Zero-Copy Host Memory	29
8.1.6	hipEventSynchronize	29
8.1.7	Summary and Recommendations	29
8.1.8	Managed memory allocation	30
8.1.9	HIP Stream Memory Operations	30
8.2	Direct Dispatch	30
8.3	HIP Runtime Compilation	31
8.4	HIP Graph	31
8.5	Device-Side Malloc	31
8.6	Use of Per-thread default stream	31
8.7	Use of Long Double Type	32
8.8	Use of _Float16 Type	32
8.9	FMA and contractions	32
8.10	Math functions with special rounding modes	32
8.11	Creating Static Libraries	32
9	HIP Porting Guide	35
9.1	Porting a New CUDA Project	35
9.1.1	General Tips	35
9.1.2	Scanning existing CUDA code to scope the porting effort	35
9.1.3	Converting a project “in-place”	36
9.1.4	Library Equivalents	37
9.2	Distinguishing Compiler Modes	37
9.2.1	Identifying HIP Target Platform	37
9.2.2	Identifying the Compiler: hip-clang or NVCC	38
9.2.3	Identifying Current Compilation Pass: Host or Device	38
9.2.4	Compiler Defines: Summary	39
9.3	Identifying Architecture Features	39
9.3.1	HIP_ARCH Defines	39
9.3.2	Device-Architecture Properties	40
9.3.3	Table of Architecture Properties	40
9.4	Finding HIP	41
9.5	Identifying HIP Runtime	42
9.6	hipLaunchKernelGGL	42
9.7	Compiler Options	42
9.7.1	Compiler options supported on AMD platforms	42
9.8	Linking Issues	43
9.8.1	Linking With hipcc	43
9.8.2	-lm Option	43
9.9	Linking Code With Other Compilers	43
9.9.1	libc++ and libstdc++	43
9.9.2	HIP Headers (hip_runtime.h, hip_runtime_api.h)	44
9.9.3	Using a Standard C++ Compiler	44

9.9.3.1	<code>cuda.h</code>	44
9.9.4	Choosing HIP File Extensions	44
9.10	Workarounds	45
9.10.1	<code>warpSize</code>	45
9.10.2	Kernel launch with group size > 256	45
9.11	<code>memcpyToSymbol</code>	45
9.12	<code>CU_POINTER_ATTRIBUTE_MEMORY_TYPE</code>	46
9.13	<code>threadfence_system</code>	47
9.13.1	Textures and Cache Control	47
9.14	More Tips	48
9.14.1	HIP Logging	48
9.14.2	Debugging <code>hipcc</code>	49
9.14.3	Editor Highlighting	49
10	Porting CUDA Driver API	51
10.1	Introduction to the CUDA Driver and Runtime APIs	51
10.1.1	<code>cuModule</code> API	51
10.1.2	<code>cuCtx</code> API	52
10.2	HIP Module and Ctx APIs	52
10.2.1	<code>hipModule</code> API	52
10.2.2	<code>hipCtx</code> API	53
10.2.3	hipify translation of CUDA Driver API	53
10.2.3.1	Address Spaces	53
10.2.3.2	Using <code>hipModuleLaunchKernel</code>	53
10.2.3.3	Additional Information	53
10.2.4	<code>hip-clang</code> Implementation Notes	53
10.2.4.1	<code>.hip_fatbin</code>	53
10.2.4.2	Initialization and Termination Functions	54
10.2.4.3	Kernel Launching	54
10.2.5	NVCC Implementation Notes	54
10.2.5.1	Interoperation between HIP and CUDA Driver	54
10.2.5.2	Compilation Options	55
10.3	HIP Module and Texture Driver API	57
11	Programming for HIP Runtime Compiler (RTC)	59
11.1	Compilation APIs	59
11.2	HIPRTC specific options	63
11.2.1	Bitcode	64
11.2.2	CU Mode vs WGP mode	64
11.3	Linker APIs	64
11.3.1	Example	65
11.3.1.1	Note	65
11.3.2	Input Types	66
11.3.3	Backward Compatibility of LLVM Bitcode/IR	66
11.3.4	Link Options	66
11.4	Error Handling	67
11.5	HIPRTC General APIs	67
11.6	Lowered Names (Mangled Names)	68
11.6.1	Note	68
11.6.2	Example	68
11.7	Versioning	69
11.8	HIP header support	70
11.9	Deprecation notice	70

12 Performance Guidelines	71
12.1 Parallel execution	71
12.1.1 Application level	71
12.1.2 Device level	71
12.1.3 Multiprocessor level	72
12.2 Memory optimization	72
12.2.1 Data Transfer	72
12.2.2 Device Memory Access	73
12.3 Optimization for maximum instruction throughput	73
12.3.1 Arithmetic instructions	74
12.3.2 Control flow instructions	74
12.3.3 Synchronization	74
12.4 Minimizing memory thrashing	75
13 Debugging with HIP	77
13.1 Tracing	77
13.2 Debugging	79
13.2.1 Debugging HIP applications	79
13.3 Useful environment variables	81
13.3.1 Kernel enqueue serialization	81
13.3.2 Making device visible	81
13.3.3 Dump code object	81
13.3.4 HSA-related environment variables (Linux)	82
13.3.5 HIP environment variable summary	82
13.4 General debugging tips	84
14 Logging HIP activity	85
14.1 Logging level	85
14.2 Logging mask	86
14.3 Logging command	86
14.4 Logging examples	87
15 Frequently asked questions	91
15.1 What APIs and features does HIP support?	91
15.2 What is not supported?	91
15.2.1 Runtime/Driver API features	91
15.2.2 Kernel language features	92
15.3 Is HIP a drop-in replacement for CUDA?	92
15.4 What specific version of CUDA does HIP support?	92
15.5 What libraries does HIP support?	93
15.6 How does HIP compare with OpenCL?	93
15.7 How does porting CUDA to HIP compare to porting CUDA to OpenCL?	93
15.8 What hardware does HIP support?	94
15.9 Do HIPIFY tools automatically convert all source code?	94
15.10 What is NVCC?	94
15.11 What is HIP-Clang?	94
15.12 Why use HIP rather than supporting CUDA directly?	94
15.13 Can I develop HIP code on an NVIDIA CUDA platform?	95
15.14 Can I develop HIP code on an AMD HIP-Clang platform?	95
15.15 How to use HIP-Clang to build HIP programs?	95
15.16 What is AMD clr?	95
15.17 What is hipothor?	96
15.18 Can I get HIP open source repository for Windows?	96
15.19 Can a HIP binary run on both AMD and NVIDIA platforms?	96

15.20	On HIP-Clang, can I link HIP code with host code compiled with another compiler such as gcc, icc, or clang?	96
15.21	Can HIP API support C style application? What is the difference between C and C++?	96
15.22	Can I install both CUDA SDK and HIP-Clang on the same machine?	97
15.23	HIP detected my platform (HIP-Clang vs NVCC) incorrectly * what should I do?	97
15.24	On CUDA, can I mix CUDA code with HIP code?	98
15.25	How do I trace HIP application flow?	98
15.26	What are the maximum limits of kernel launch parameters?	98
15.27	Are <code>__shfl_*_sync</code> functions supported on HIP platform?	98
15.28	How to create a guard for code that is specific to the host or the GPU?	98
15.29	Why <code>_OpenMP</code> is undefined when compiling with <code>-fopenmp</code> ?	99
15.30	Does the HIP-Clang compiler support extern shared declarations?	99
15.31	I have multiple HIP enabled devices and I am getting an error code <code>hipErrorSharedObjectInitFailed</code> with the message “Error: shared object initialization failed”?	99
15.32	How to use per-thread default stream in HIP?	100
15.33	How to use complex multiplication and division operations?	100
15.34	Can I develop applications with HIP APIs on Windows the same on Linux?	100
15.35	Does HIP support LUID?	101
15.36	How can I know the version of HIP?	101
16	HIP Runtime API Reference	103
16.1	Related Pages	104
16.2	Modules	104
16.3	Data Structures	104
16.3.1	Data Structures	104
16.3.2	Data Structure Index	104
16.3.3	Data Fields	104
16.3.3.1	All	104
16.3.3.1.1	Data Fields	104
16.3.3.1.2	Data Fields	104
16.3.3.1.3	Data Fields	104
16.3.3.1.4	Data Fields	104
16.3.3.1.5	Data Fields	104
16.3.3.1.6	Data Fields	104
16.3.3.1.7	Data Fields	104
16.3.3.1.8	Data Fields	104
16.3.3.1.9	Data Fields	104
16.3.3.1.10	Data Fields	104
16.3.3.1.11	Data Fields	104
16.3.3.1.12	Data Fields	104
16.3.3.1.13	Data Fields	104
16.3.3.1.14	Data Fields	104
16.3.3.1.15	Data Fields	104
16.3.3.1.16	Data Fields	104
16.3.3.1.17	Data Fields	104
16.3.3.1.18	Data Fields	104
16.3.3.1.19	Data Fields	104
16.3.3.1.20	Data Fields	104
16.3.3.1.21	Data Fields	104
16.3.3.1.22	Data Fields	104
16.3.3.1.23	Data Fields	104
16.3.3.1.24	Data Fields	104
16.3.3.2	Variables	104

16.3.3.2.1	Data Fields - Variables	104
16.3.3.2.2	Data Fields - Variables	104
16.3.3.2.3	Data Fields - Variables	104
16.3.3.2.4	Data Fields - Variables	104
16.3.3.2.5	Data Fields - Variables	104
16.3.3.2.6	Data Fields - Variables	104
16.3.3.2.7	Data Fields - Variables	104
16.3.3.2.8	Data Fields - Variables	104
16.3.3.2.9	Data Fields - Variables	104
16.3.3.2.10	Data Fields - Variables	104
16.3.3.2.11	Data Fields - Variables	104
16.3.3.2.12	Data Fields - Variables	104
16.3.3.2.13	Data Fields - Variables	104
16.3.3.2.14	Data Fields - Variables	104
16.3.3.2.15	Data Fields - Variables	104
16.3.3.2.16	Data Fields - Variables	104
16.3.3.2.17	Data Fields - Variables	104
16.3.3.2.18	Data Fields - Variables	104
16.3.3.2.19	Data Fields - Variables	104
16.3.3.2.20	Data Fields - Variables	104
16.3.3.2.21	Data Fields - Variables	104
16.3.3.2.22	Data Fields - Variables	104
16.3.3.2.23	Data Fields - Variables	104
16.3.3.2.24	Data Fields - Variables	104
16.4	Files	104
16.4.1	File List	104
16.4.2	Globals	104
16.4.2.1	All	104
16.4.2.1.1	Globals	104
16.4.2.1.2	Globals	104
16.4.2.1.3	Globals	104
16.4.2.1.4	Globals	104
16.4.2.1.5	Globals	104
16.4.2.1.6	Globals	104
16.4.2.1.7	Globals	104
16.4.2.1.8	Globals	104
16.4.2.1.9	Globals	104
16.4.2.1.10	Globals	104
16.4.2.1.11	Globals	104
16.4.2.1.12	Globals	104
16.4.2.1.13	Globals	104
16.4.2.1.14	Globals	104
16.4.2.1.15	Globals	104
16.4.2.1.16	Globals	104
16.4.2.1.17	Globals	104
16.4.2.1.18	Globals	104
16.4.2.1.19	Globals	104
16.4.2.2	Functions	104
16.4.2.2.1	Globals	104
16.4.2.2.2	Globals	104
16.4.2.2.3	Globals	104
16.4.2.2.4	Globals	104
16.4.2.2.5	Globals	104
16.4.2.2.6	Globals	104

16.4.2.2.7	Globals	104
16.4.2.2.8	Globals	104
16.4.2.2.9	Globals	104
16.4.2.2.10	Globals	104
16.4.2.2.11	Globals	104
16.4.2.2.12	Globals	104
16.4.2.2.13	Globals	104
16.4.2.2.14	Globals	104
16.4.2.2.15	Globals	104
16.4.2.2.16	Globals	104
16.4.2.3	Globals	104
16.4.2.4	Globals	104
16.4.2.5	Globals	104
16.4.2.6	Enumerator	104
16.4.2.6.1	Globals	104
16.4.2.6.2	Globals	104
16.4.2.7	Globals	104
17	C++ Language Extensions	105
17.1	Function-type qualifiers	105
17.1.1	__device__	105
17.1.2	__global__	106
17.1.3	__host__	106
17.2	Calling __global__ functions	106
17.3	Kernel launch example	107
17.4	Variable type qualifiers	108
17.4.1	__constant__	108
17.4.2	__shared__	108
17.4.3	__managed__	108
17.4.4	__restrict__	108
17.5	Built-in variables	108
17.5.1	Coordinate built-ins	108
17.5.2	warpSize	109
17.6	Vector types	109
17.6.1	Short vector types	109
17.6.2	dim3	110
17.7	Memory fence instructions	110
17.8	Synchronization functions	110
17.9	Math functions	110
17.9.1	Single precision mathematical functions	110
17.9.2	Double precision mathematical functions	120
17.9.3	Integer intrinsics	130
17.9.4	Floating-point Intrinsics	132
17.10	Texture functions	134
17.11	Surface functions	135
17.12	Timer functions	135
17.13	Atomic functions	135
17.13.1	Unsafe floating-point atomic RMW operations	137
17.14	Warp cross-lane functions	137
17.14.1	Warp vote and ballot functions	138
17.14.2	Warp match functions	139
17.14.3	Warp shuffle functions	139
17.15	Cooperative groups functions	140
17.16	Warp matrix functions	140

17.17 Independent thread scheduling	141
17.18 Profiler Counter Function	141
17.19 Assert	141
17.20 printf	142
17.21 Device-Side Dynamic Global Memory Allocation	142
17.22 __launch_bounds__	142
17.22.1 Compiler Impact	143
17.22.2 CU and EU Definitions	143
17.22.3 Porting from CUDA __launch_bounds__	143
17.22.4 maxregcount	144
17.23 Asynchronous Functions	144
17.23.1 Memory stream	144
17.23.2 Peer to peer	159
17.23.3 Memory management	161
17.23.4 External Resource Interoperability	190
17.24 Register Keyword	193
17.25 Pragma Unroll	193
17.26 In-Line Assembly	193
17.27 C++ Support	193
17.28 Kernel Compilation	194
17.29 gfx-arch-specific-kernel	194
18 Table Comparing Syntax for Different Compute APIs	195
18.1 Notes	196
19 HSA Runtime API for ROCm	197
20 HIP deprecated runtime API functions	203
20.1 Context management	203
20.2 Memory management	204
20.3 Profiler control	204
20.4 Texture management	204
21 FP8 Numbers	207
21.1 FP8 Definition	207
21.2 HIP Header	207
21.3 Supported Devices	207
21.4 Usage	208
22 License	213
Index	215

The Heterogeneous-computing Interface for Portability (HIP) API is a C++ runtime API and kernel language that lets developers create portable applications for AMD and NVIDIA GPUs from single source code.

For HIP supported AMD GPUs on multiple operating systems, see:

- [Linux system requirements](#)
- [Microsoft Windows system requirements](#)

The CUDA enabled NVIDIA GPUs are supported by HIP. For more information, see [GPU Compute Capability](#).

On the AMD ROCm platform, HIP provides header files and runtime library built on top of HIP-Clang compiler in the repository *Compute Language Runtime (CLR)*, which contains source codes for AMD's compute languages runtimes.

On non-AMD platforms, like NVIDIA, HIP provides header files required to support non-AMD specific back-end implementation in the repository 'hipother', which translates from the HIP runtime APIs to CUDA runtime APIs.

OVERVIEW

Install

- *Installing HIP*
- *Building HIP from source*

Conceptual

- *Understanding the HIP programming model*
- *Programming model reference*
- *Hardware Implementation*
- *AMD compute language runtime (CLR)*

How to

- *Programming Manual*
- *HIP Porting Guide*
- *HIP Porting: Driver API Guide*
- *Programming for HIP Runtime Compiler (RTC)*
- *Performance Guidelines*
- *Debugging with HIP*
- *Logging HIP activity*
- *Frequently asked questions*

Reference

- *HIP Runtime API Reference*
- *C++ language extensions*
- *Comparing Syntax for different APIs*
- *HSA Runtime API for ROCm*
- *List of deprecated APIs*
- *FP8 numbers in HIP*

Tutorial

- *HIP examples*
- *HIP test samples*

Known issues are listed on the [HIP GitHub repository](#).

To contribute features or functions to the HIP project, refer to [Contributing to HIP](#). To contribute to the documentation, refer to [Contributing to ROCm docs](#) page.

You can find licensing information on the [Licensing](#) page.

INSTALL HIP

HIP can be installed on AMD (ROCm with HIP-Clang) and NVIDIA (CUDA with NVCC) platforms.

Note: The version definition for the HIP runtime is different from CUDA. On an AMD platform, the `hipRuntimeGetVersion` function returns the HIP runtime version; on an NVIDIA platform, this function returns the CUDA runtime version.

2.1 Prerequisites

AMD

Refer to the Prerequisites section in the ROCm install guides:

- [rocm-install-on-linux:reference/system-requirements](#)
- [rocm-install-on-windows:reference/system-requirements](#)

NVIDIA

Check the system requirements in the [NVIDIA CUDA Installation Guide](#).

2.2 Installation

AMD

HIP is automatically installed during the ROCm installation. If you haven't yet installed ROCm, you can find installation instructions [here](#):

- [rocm-install-on-linux:index](#)
- [rocm-install-on-windows:index](#)

By default, HIP is installed into `/opt/rocm/hip`.

Note: There is no autodetection for the HIP installation. If you choose to install it somewhere other than the default location, you must set the `HIP_PATH` environment variable as explained in [Build HIP from source](#).

NVIDIA

1. Install the NVIDIA driver.

```
sudo apt-get install ubuntu-drivers-common && sudo ubuntu-drivers autoinstall  
sudo reboot
```

Alternatively, you can download the latest [CUDA Toolkit](#).

2. Install the `hip-runtime-nvidia` and `hip-dev` packages. This installs the CUDA SDK and HIP porting layer.

```
apt-get install hip-runtime-nvidia hip-dev
```

The default paths are:

- CUDA SDK: `/usr/local/cuda`
- HIP: `/opt/rocm/hip`

You can optionally add `/opt/rocm/bin` to your path, which can make it easier to use the tools.

2.3 Verify your installation

Run `hipconfig` in your installation path.

```
/opt/rocm/bin/hipconfig --full
```


BUILD HIP FROM SOURCE

3.1 Prerequisites

HIP code can be developed either on AMD ROCm platform using HIP-Clang compiler, or a CUDA platform with nvcc installed. Before building and running HIP, make sure drivers and prebuilt packages are installed properly on the platform.

You also need to install Python 3, which includes the CppHeaderParser package. Install Python 3 using the following command:

```
apt-get install python3
```

Check and install CppHeaderParser package using the command:

```
pip3 install CppHeaderParser
```

Install ROCm LLVM package using the command:

```
apt-get install rocm-llvm-dev
```

3.2 Building the HIP runtime

Set the repository branch using the variable: ROCM_BRANCH. For example, for ROCm 6.1, use:

```
export ROCM_BRANCH=rocm-6.1.x
```

AMD

1. Get HIP source code.

Note: Starting in ROCm 5.6, CLR is a new repository that includes the former ROCclr, HIPAMD and OpenCL repositories. OpenCL provides headers that ROCclr runtime depends on.

Note: Starting in ROCm 6.1, a new repository `hipother` is added to ROCm, which is branched out from HIP. `hipother` provides files required to support the HIP back-end implementation on some non-AMD platforms, like NVIDIA.

```
git clone -b "$ROCM_BRANCH" https://github.com/ROCm/clr.git
git clone -b "$ROCM_BRANCH" https://github.com/ROCm/hip.git
```

CLR (Compute Language Runtime) repository includes ROCclr, HIPAMD and OpenCL.

ROCclr (ROCm Compute Language Runtime) is a virtual device interface which is defined on the AMD platform. HIP runtime uses ROCclr to interact with different backends.

HIPAMD provides implementation specifically for HIP on the AMD platform.

OpenCL provides headers that ROCclr runtime currently depends on. hipother provides headers and implementation specifically for non-AMD HIP platforms, like NVIDIA.

2. Set the environment variables.

```
export CLR_DIR="$(readlink -f clr)"
export HIP_DIR="$(readlink -f hip)"
```

3. Build HIP.

```
cd "$CLR_DIR"
mkdir -p build; cd build
cmake -DHIP_COMMON_DIR=$HIP_DIR -DHIP_PLATFORM=amd -DCMAKE_PREFIX_PATH="/opt/rocm/" \
↪ -DCMAKE_INSTALL_PREFIX=$PWD/install -DHIP_CATCH_TEST=0 -DCLR_BUILD_HIP=ON -DCLR_
↪ BUILD_OCL=OFF ..

make -j$(nproc)
sudo make install
```

Note: Note, if you don't specify CMAKE_INSTALL_PREFIX, the HIP runtime is installed at <ROCM_PATH>.

By default, release version of HIP is built. If need debug version, you can put the option CMAKE_BUILD_TYPE=Debug in the command line.

Default paths and environment variables:

- **HIP is installed into <ROCM_PATH>.** This can be overridden by setting the **INSTALL_PREFIX** as the **command option**.
environment variable.
- **HSA is in <ROCM_PATH>.** This can be overridden by setting the **HSA_PATH**
environment variable.
- **Clang is in <ROCM_PATH>/llvm/bin.** This can be overridden by setting the
HIP_CLANG_PATH environment variable.
- **The device library is in <ROCM_PATH>/lib.** This can be overridden by setting the
DEVICE_LIB_PATH environment variable.
- **Optionally, you can add <ROCM_PATH>/bin to your PATH, which can make it easier to**
use the tools.
- **Optionally, you can set HIPCC_VERBOSE=7 to output the command line for compilation.**

After you run the `make install` command, HIP is installed to <ROCM_PATH> by default, or \$PWD/install/hip while **INSTALL_PREFIX** is defined.

1. Generate a profiling header after adding/changing a HIP API.

When you add or change a HIP API, you may need to generate a new `hip_prof_str.h` header. This header is used by ROCm tools to track HIP APIs, such as `rocprofiler` and `roctracer`.

To generate the header after your change, use the `hip_prof_gen.py` tool located in `hipamd/src`.

Usage:

```
`hip_prof_gen.py [-v] <input HIP API .h file> <patched srcs path> <previous_
↪output> [<output>]`
```

Flags:

- `-v`: Verbose messages
- `-r`: Process source directory recursively
- `-t`: API types matching check
- `--priv`: Private API check
- `-e`: On error exit mode
- `-p`: `HIP_INIT_API` macro patching mode

Example usage:

```
hip_prof_gen.py -v -p -t --priv <hip>/include/hip/hip_runtime_api.h \
<hipamd>/src <hipamd>/include/hip/amd_detail/hip_prof_str.h \
<hipamd>/include/hip/amd_detail/hip_prof_str.h.new
```

NVIDIA

1. Get the HIP source code.

```
git clone -b "$ROCM_BRANCH" https://github.com/ROCm/clr.git
git clone -b "$ROCM_BRANCH" https://github.com/ROCm/hip.git
git clone -b "$ROCM_BRANCH" https://github.com/ROCm/hipother.git
```

2. Set the environment variables.

```
export CLR_DIR="$(readlink -f clr)"
export HIP_DIR="$(readlink -f hip)"
export HIP_OTHER="$(readlink -f hipother)"
```

3. Build HIP.

```
cd "$CLR_DIR"
mkdir -p build; cd build
cmake -DHIP_COMMON_DIR=$HIP_DIR -DHIP_PLATFORM=nvidia -DCMAKE_INSTALL_PREFIX=$PWD/
↪install -DHIP_CATCH_TEST=0 -DCLR_BUILD_HIP=ON -DCLR_BUILD_OCL=OFF -DHIPNV_DIR=
↪$HIP_OTHER/hipnv ..
make -j$(nproc)
sudo make install
```

3.3 Build HIP tests

AMD

- Build HIP catch tests.

HIP catch tests are separate from the HIP project and use Catch2.

- Get HIP tests source code.

```
git clone -b "$ROCM_BRANCH" https://github.com/ROCm/hip-tests.git
```

- Build HIP tests from source.

```
export HIPTESTS_DIR="$(readlink -f hip-tests)"
cd "$HIPTESTS_DIR"
mkdir -p build; cd build
cmake ../catch -DHIP_PLATFORM=amd -DHIP_PATH=$CLR_DIR/build/
↪install # or any path where HIP is installed; for example: ``/
↪opt/rocm``
make build_tests
ctest # run tests
```

HIP catch tests are built in \$HIPTESTS_DIR/build.

To run any single catch test, use this example:

```
cd $HIPTESTS_DIR/build/catch_tests/unit/texture
./TextureTest
```

- Build a HIP Catch2 standalone test.

```
cd "$HIPTESTS_DIR"
hipcc $HIPTESTS_DIR/catch/unit/memory/hipPointerGetAttributes.cc \
-I ./catch/include ./catch/hipTestMain/standalone_main.cc \
-I ./catch/external/Catch2 -o hipPointerGetAttributes
./hipPointerGetAttributes
...
All tests passed
```

NVIDIA

The commands to build HIP tests on an NVIDIA platform are the same as on an AMD platform. However, you must first set `-DHIP_PLATFORM=nvidia`.

3.4 Run HIP

After installation and building HIP, you can compile your application and run. A simple example is [square sample](#).

UNDERSTANDING THE HIP PROGRAMMING MODEL

The HIP programming model makes it easy to map data-parallel C/C++ algorithms to massively parallel, wide single instruction, multiple data (SIMD) architectures, such as GPUs. A basic understanding of the underlying device architecture helps you make efficient use of HIP and general purpose graphics processing unit (GPGPU) programming in general.

4.1 RDNA & CDNA Architecture Summary

Most GPU architectures, like RDNA and CDNA, have a hierarchical structure. The innermost piece is a SIMD-enabled vector Arithmetic Logical Unit (ALU). In addition to the vector ALUs, most recent GPUs also house matrix ALUs for accelerating algorithms involving matrix multiply-accumulate operations. AMD GPUs also contain scalar ALUs, that can be used to reduce the load on the vector ALU by performing operations which are uniform for all threads of a warp.

A set of ALUs, together with register files, caches and shared memory, comprise a larger block, often referred to as a compute unit (CU), e.g. in OpenCL and AMD block diagrams, or as streaming multiprocessor (SM).

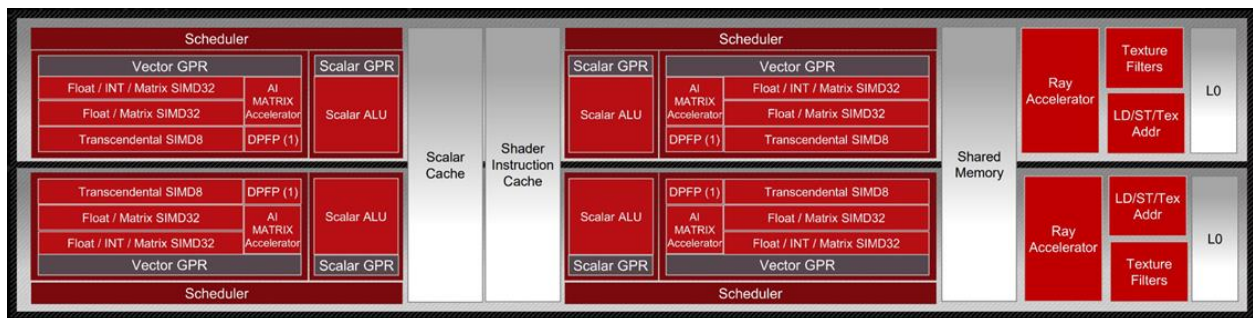


Fig. 4.1: Block Diagram of an RDNA3 Compute Unit.

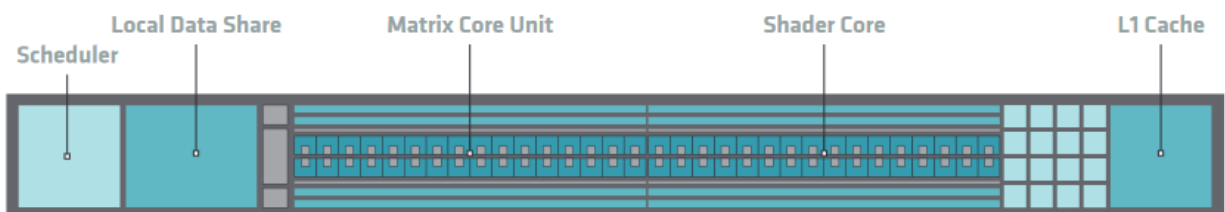


Fig. 4.2: Block Diagram of a CDNA3 Compute Unit.

For implementation in hardware, multiple CUs are grouped together into a Shader Engine or Compute Engine, typically sharing some fixed function units or memory subsystem resources.

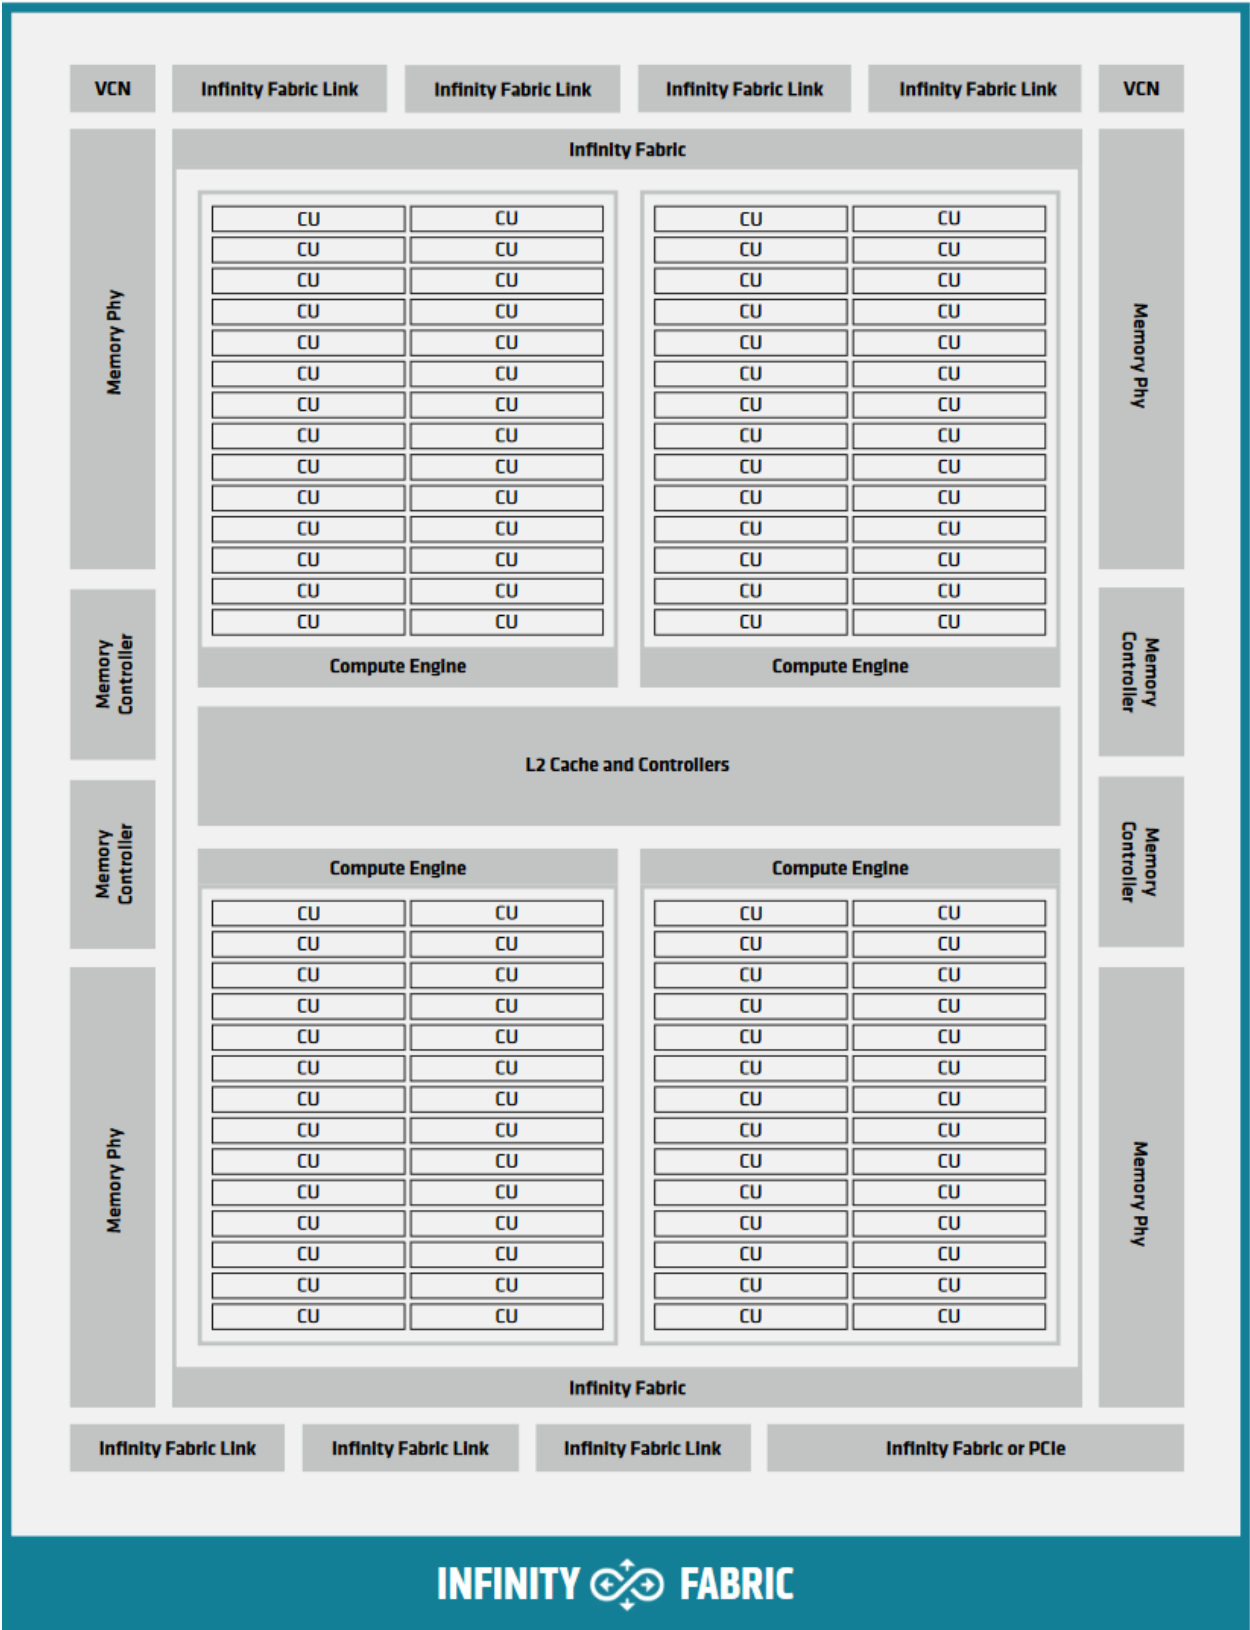


Fig. 4.3: Block Diagram of a CDNA2 Graphics Compute Die.

4.2 Single Instruction Multiple Threads

The single instruction, multiple threads (SIMT) programming model behind the HIP device-side execution is a middle-ground between SMT (Simultaneous Multi-Threading) programming known from multicore CPUs, and SIMD (Single Instruction, Multiple Data) programming mostly known from exploiting relevant instruction sets on CPUs (for example SSE/AVX/Neon).

A HIP device compiler maps SIMT code written in HIP C++ to an inherently SIMD architecture (like GPUs). This is done by scalarizing the entire kernel and issuing the scalar instructions of multiple kernel instances to each of the SIMD engine lanes, rather than exploiting data parallelism within a single instance of a kernel and spreading identical instructions over the available SIMD engines.

Consider the following kernel:

```
__global__ void k(float4* a, const float4* b)
{
    int tid = threadIdx.x;
    int bid = blockIdx.x;
    int dim = blockDim.x;

    a[tid] += (tid + bid - dim) * b[tid];
}
```

The incoming four-vector of floating-point values `b` is multiplied by a scalar and then added element-wise to the four-vector floating-point values of `a`. On modern SIMD-capable architectures the four-vector ops are expected to compile to a single SIMD instruction. GPU execution of this kernel however will typically look the following:

Fig. 4.4: Instruction flow of the sample SIMT program.

In HIP, lanes of a SIMD architecture are fed by mapping threads of a SIMT execution, one thread down each lane of a SIMD engine. Execution parallelism usually isn't exploited from the width of the built-in vector types, but via the thread id constants `threadIdx.x`, `blockIdx.x`, etc. For more details, refer to *Inherent Thread Model*.

4.3 Heterogeneous Programming

The HIP programming model assumes two execution contexts. One is referred to as *host* while compute kernels execute on a *device*. These contexts have different capabilities, therefore slightly different rules apply. The *host* execution is defined by the C++ abstract machine, while *device* execution follows the HIP model, primarily defined by SIMT. These execution contexts in code are signified by the `__host__` and `__device__` decorators. There are a few key differences between the two:

- The C++ abstract machine assumes a unified memory address space, meaning that one can always access any given address in memory (assuming the absence of data races). HIP however introduces several memory namespaces, an address from one means nothing in another. Moreover, not all address spaces are accessible from all contexts.

If one were to look at *Block Diagram of a CDNA2 Graphics Compute Die*, and inside the *Block Diagram of a CDNA3 Compute Unit*, every Compute Unit has an instance of storage backing the namespace `__shared__`. Even if the host were to have access to these regions of memory, the performance benefits of the segmented memory subsystem are supported by the inability of asynchronous access from the host.

- Not all C++ language features map cleanly to typical device architectures, some are very expensive (meaning slow) to implement on GPU devices, therefore they are forbidden in device contexts to avoid users tapping into features that unexpectedly decimate their program's performance. Offload devices targeted by HIP aren't general

purpose devices, at least not in the sense a CPU is. HIP focuses on data parallel computations and as such caters to throughput optimized architectures, such as GPUs or accelerators derived from GPU architectures.

- Asynchrony is at the forefront of the HIP API. Computations launched on the device execute asynchronously with respect to the host, and it is the user's responsibility to synchronize their data dispatch/fetch with computations on the device. HIP does perform implicit synchronization on occasions, more advanced than other APIs such as OpenCL or SYCL, in which the responsibility of synchronization mostly depends on the user.

PROGRAMMING MODEL REFERENCE

HIP defines a model for mapping single instruction, multiple threads (SIMT) programs onto various architectures, primarily GPUs. While the model may be expressed in most imperative languages, (for example Python via PyHIP) this document will focus on the original C/C++ API of HIP.

5.1 Threading Model

The SIMT nature of HIP is captured by the ability to execute user-provided device programs, expressed as single-source C/C++ functions or sources compiled online/offline to binaries in bulk.

Multiple instances of the device program (or kernel) are called threads and may execute in parallel. All uniquely identified by a set of integral values, or thread IDs. The set of integers identifying a thread relate to the hierarchy in which threads execute.

5.1.1 Inherent Thread Model

The thread hierarchy inherent to how AMD GPUs operate is depicted in [Fig. 5.1](#).

Fig. 5.1: Hierarchy of thread groups.

Warp

The innermost grouping is called a warp, or a wavefront in ISA terms. A warp is the most tightly coupled groups of threads, both physically and logically.

When referring to threads inside a warp, they may be called lanes, and the integral value identifying them the lane ID. Lane IDs aren't queried like other thread IDs, but are user-calculated. As a consequence they are only as multidimensional as the user interprets the calculated values to be.

The size of a warp is architecture dependent and always fixed. Warps are signified by the set of communication primitives at their disposal, as discussed in [Warp cross-lane functions](#).

Block

The middle grouping is called a block or thread block. The defining feature of a block is that all threads in a block will share an instance of memory which they may use to share data or synchronize with one another.

The size of a block is user-configurable but is maximized by the queryable capabilities of the executing hardware. The unique ID of the thread within a block is 3-dimensional as provided by the API. When linearizing thread IDs within a block, assume the "fast index" being dimension x , followed by the y and z dimensions.

Grid

The outermost grouping is called a grid. A grid manifests as a single dispatch of kernels for execution. The

unique ID of each block within a grid is 3-dimensional, as provided by the API and is queryable by every thread within the block.

5.1.2 Cooperative Groups Thread Model

The Cooperative Groups API introduces new APIs to launch, group, subdivide, synchronize and identify threads, as well as some predefined group-collective algorithms, but most importantly a matching threading model to think in terms of. It relaxes some restrictions of the *Inherent Thread Model* imposed by the strict 1:1 mapping of architectural details to the programming model.

The rich set of APIs introduced by Cooperative Groups allow the programmer to define their own set of thread groups which may fit their user-cases better than those defined by the hardware. The set of implicit groups by kernel launch parameters are still available.

The thread hierarchy abstraction of Cooperative Groups manifest as depicted in [Fig. 5.2](#).

Fig. 5.2: Cooperative group thread hierarchy.

Multi Grid

An abstraction of potentially multiple simultaneous launches of the same kernel over multiple devices. Grids inside a multi device kernel launch need not be of uniform size, thus allowing taking into account different device capabilities and preferences.

Deprecated since version 5.0.

Grid

Same as the *Inherent Thread Model* Grid entity. The ability to synchronize over a grid requires the kernel to be launched using the Cooperative Groups API.

Block

Same as the *Inherent Thread Model* Block entity.

Note: Explicit warp-level thread handling is absent from the Cooperative Groups API. In order to exploit the known hardware SIMD width on which built-in functionality translates to simpler logic, one may use the group partitioning part of the API, such as `tiled_partition`.

5.2 Memory Model

The hierarchy of threads introduced by *Inherent Thread Model* is induced by the memory subsystem of GPUs. [Fig. 5.3](#) summarizes that memory namespaces and how they relate to the various levels of the threading model.

Fig. 5.3: Memory hierarchy.

Local or per-thread memory

Read-write storage only visible to the threads defining the given variables, also called per-thread memory. The size of a block for a given kernel, the number of concurrent warps are limited by local memory usage. This relates to an important aspect: occupancy. This is the default memory namespace.

Shared memory

Read-write storage visible to all the threads in a given block.

Global

Read-write storage visible to all threads in a given grid. There are specialized versions of global memory with different usage semantics which are typically backed by the same hardware storing global.

Constant

Read-only storage visible to all threads in a given grid. It is a limited segment of global with queryable size.

Texture

Read-only storage visible to all threads in a given grid and accessible through additional APIs.

Surface

A read-write version of texture memory.

5.3 Execution Model

HIP programs consist of two distinct scopes:

- The host-side API running on the host processor. There are two APIs available:
 - The HIP runtime API which enables use of the single-source programming model.
 - The HIP driver API which sits at a lower level and most importantly differs by removing some facilities provided by the runtime API, most importantly around kernel launching and argument setting. It is geared towards implementing abstractions atop, such as the runtime API itself.
- The device-side kernels running on GPUs. Both the host and the device-side APIs have synchronous and asynchronous functions in them.

5.3.1 Host-side execution

The part of the host-side API which deals with device management and their queries are synchronous. All asynchronous APIs, such as kernel execution, data movement and potentially data allocation/freeing all happen in the context of device streams.

Streams are FIFO buffers of commands to execute relating to a given device. Commands which enqueue tasks on a stream all return promptly and the command is executed asynchronously. All side effects of a command on a stream are visible to all subsequent commands on the same stream. Multiple streams may point to the same device and those streams may be fed from multiple concurrent host-side threads. Execution on multiple streams may be concurrent but isn't required to be.

Asynchronous APIs involving a stream all return a stream event which may be used to synchronize the execution of multiple streams. A user may enqueue a barrier onto a stream referencing an event. The barrier will block until the command related to the event does not complete, at which point all side effects of the command shall be visible to commands following the barrier, even if those side effects manifest on different devices.

Streams also support executing user-defined functions as callbacks on the host. The stream will not launch subsequent commands until the callback completes.

5.3.2 Device-side execution

The SIMT programming model behind the HIP device-side execution is a middle-ground between SMT (Simultaneous Multi-Threading) programming known from multicore CPUs, and SIMD (Single Instruction, Multiple Data) programming mostly known from exploiting relevant instruction sets on CPUs (for example SSE/AVX/Neon).

A HIP device compiler maps our SIMT code written in HIP C++ to an inherently SIMD architecture (like GPUs) not by exploiting data parallelism within a single instance of a kernel and spreading identical instructions over the SIMD engines at hand, but by scalarizing the entire kernel and issuing the scalar instructions of multiple kernel instances to each of the SIMD engine lanes.

5.3.3 Kernel launch

Kernels may be launched in multiple ways all with different syntaxes and intended use-cases.

- Using the triple-chevron `<<<. . .>>>` operator on a `__global__` annotated function.
- Using `hipLaunchKernelGGL()` on a `__global__` annotated function.

Tip: This name by default is a macro expanding to triple-chevron. In cases where language syntax extensions are undesirable, or where launching templated and/or overloaded kernel functions define the `HIP_TEMPLATE_KERNEL_LAUNCH` preprocessor macro before including the HIP headers to turn it into a templated function.

- Using the launch APIs supporting the triple-chevron syntax directly.

Caution: These APIs are intended to be used/generated by tools such as the HIP compiler itself and not intended towards end-user code. Should you be writing a tool having to launch device code using HIP, consider using these over the alternatives.

HARDWARE IMPLEMENTATION

This chapter describes the typical hardware implementation of GPUs supported by HIP, and how the *Inherent Thread Model* maps to the hardware.

6.1 Compute Units

The basic building block of a GPU is a compute unit (CU), also known as streaming multiprocessor (SM) on NVIDIA GPUs. The thread blocks making up a grid are scheduled for execution on CUs. Each block is assigned to an individual CU, and a CU can accommodate several blocks. Depending on their resource usage up to thousands of threads can reside on a CU.

CUs contain an array of processing elements, referred to as vector ALU (VALU), that execute the actual instructions of the threads according to the *SIMT model*, together with the necessary registers and caches.

The threads are executed in groupings called warps. The amount of threads making up a warp is architecture dependent. On AMD GPUs the warp size is commonly 64 threads, except in RDNA architectures which can utilize a warp size of 32 or 64 respectively. The warp size of supported AMD GPUs is listed in the `rocm:reference/gpu-arch-specs`. NVIDIA GPUs have a warp size of 32.

In contrast to CPUs, GPUs generally do not employ complex cache structures or control logic, like branch prediction or out-of-order execution, but instead rely on massive hardware multithreading to hide latency.

Context switching between warps residing on a CU incurs no overhead, as the context for the warps is stored on the CU and does not need to be fetched from memory. If there are not enough free registers to accommodate all warps of a block, the block can not be scheduled to that CU and it has to wait until other blocks finish execution.

The amount of warps that can reside concurrently on a CU, known as occupancy, is determined by the warp's resource usage of registers and shared memory.

Fig. 6.1: An AMD Graphics Core Next (GCN) CU. The CDNA and RDNA CUs are based on variations of the GCN CU.

On AMD GCN GPUs the basic structure of a CU is:

- four Single Instruction Multiple Data units (SIMDs)
- a vector cache
- a local data share
- and a scalar unit

6.1.1 SIMD

A SIMD consists of a VALU, that executes the instruction of a warp, together with a register file, that provides the registers warps.

The size of the warp is inherently related to the width of the vector ALU of the SIMD. On GCN compute units the width of the VALU is 16, so a warp can be issued to a SIMD every 4 cycles. Since a CU has 4 SIMDs it issues one warp per cycle. The instructions of a warp are effectively executed in lock-step.

A SIMD always executes the same instruction for the whole VALU. If the control flow of a warp diverges, the performance is decreased, as the results for the threads that do not participate in that branch have to be masked out, and the instructions of the other branch have to be executed in the same way. The best performance can therefore be achieved when thread divergence is kept to a warp level, i.e. when all threads in a warp take the same execution path.

6.1.2 Vector Cache

The usage of cache on a GPU differs from that on a CPU, as there is less cache available per thread. Its main purpose is to coalesce memory accesses of the warps in order to reduce the amount of accesses to device memory, and make that memory available for other warps that currently reside on the compute unit, that also need to load those values.

6.1.3 Local Data Share

The local data share is memory that is accessible to all threads within a block. Its latency and bandwidth is comparable to that of the vector cache. It can be used to share memory between the threads in a block, or as a software managed cache.

6.1.4 Scalar Unit

The scalar unit performs instructions that are uniform within a warp. It thereby improves efficiency and reduces the pressure on the vector ALUs and the vector register file.

6.2 CDNA Architecture

The general structure of CUs stays mostly as it is in GCN architectures. The most prominent change is the addition of matrix ALUs, which can greatly improve the performance of algorithms involving matrix multiply-accumulate operations for int8, float16, bfloat16 or float32.

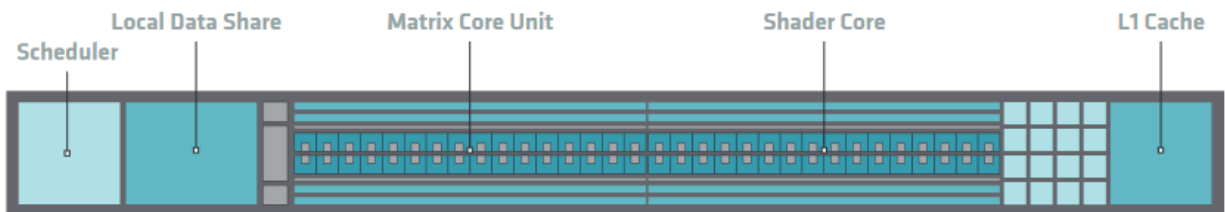


Fig. 6.2: Block Diagram of a CDNA3 Compute Unit.

6.3 RDNA Architecture

RDNA makes a fundamental change to CU design, by changing the size of a warp to 32 threads. This is done by effectively combining two GCN5 SIMDs, creating a VALU of width 32, so that a whole warp can be issued in one cycle. The CU is also replaced by the work group processor (WGP), which encompasses two CUs. For backwards compatibility the WGP can also run in wave64 mode, in which it issues a warp of size 64 in two cycles.

It also adds an extra layer of cache to the WGP, shared by the CUs within it. This cache is referred to as L1 cache, promoting the per-CU cache to an L0 cache.

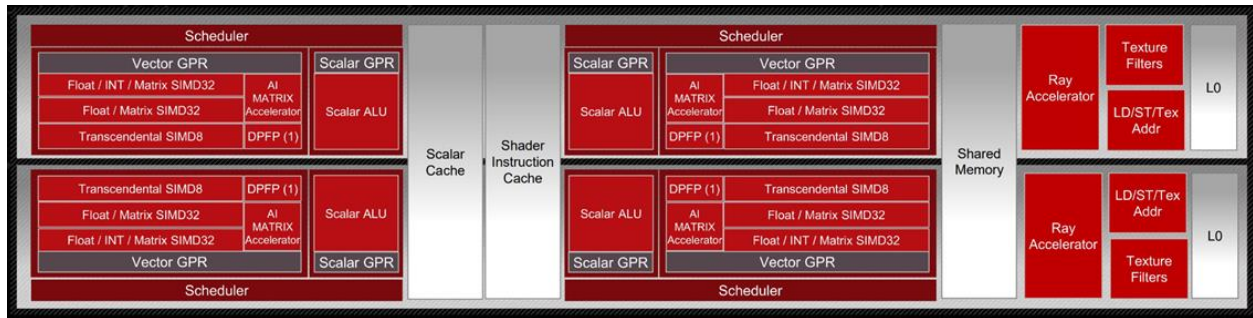


Fig. 6.3: Block Diagram of an RDNA3 work group processor.

6.4 Shader Engines

For hardware implementation's sake, multiple CUs are grouped together into a Shader Engine or Compute Engine, typically sharing some fixed function units or memory subsystem resources.

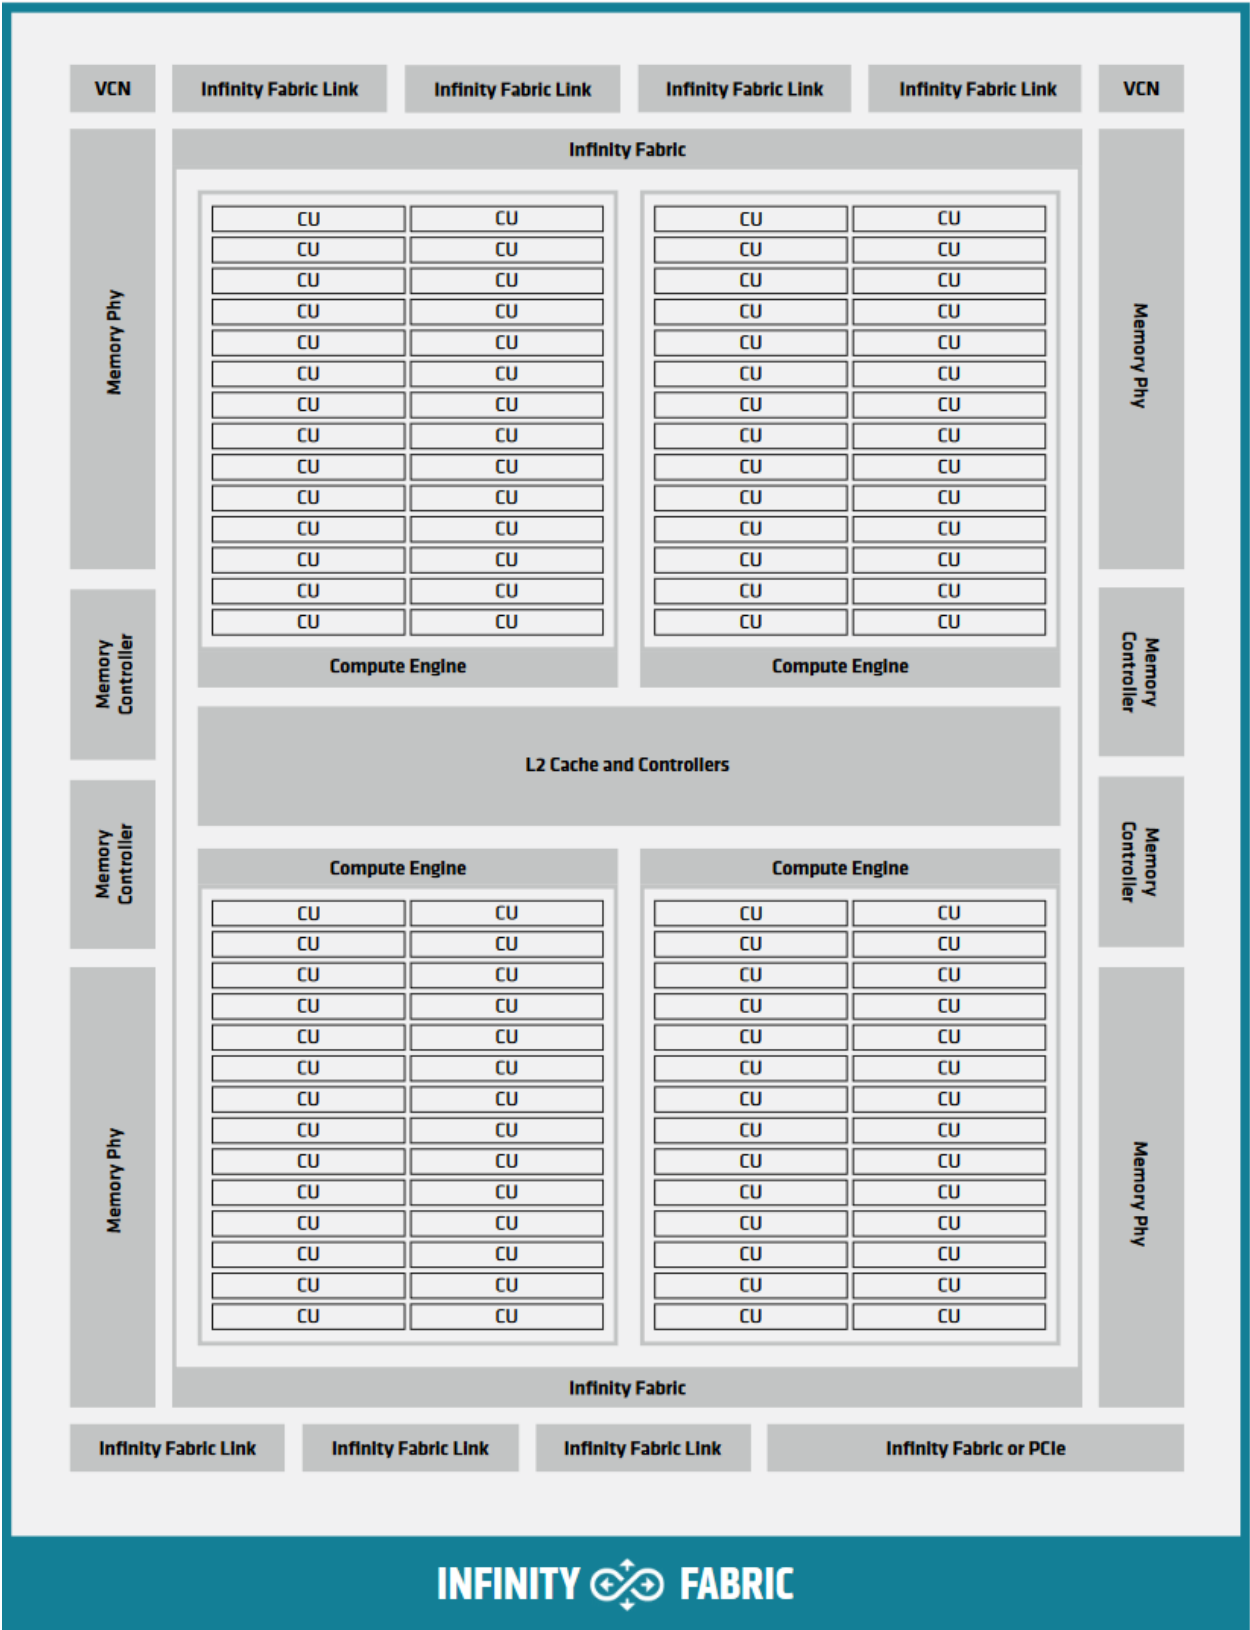


Fig. 6.4: Block Diagram of a CDNA2 Graphics Compute Die.

AMD COMPUTE LANGUAGE RUNTIME (CLR)

CLR contains source codes for AMD's compute languages runtimes: HIP and OpenCL™. CLR is the part of HIP runtime which is supported on the AMD ROCm platform, it provides a header and runtime library built on top of HIP-Clang compiler. For developers and users, CLR implements HIP runtime APIs including streams, events, and memory APIs, which is a object library that is linked with the application. The source codes for all headers and the library implementation are available on GitHub in the [clr repository](#).

7.1 Project Organisation

CLR includes the following source code,

- `hipamd` - contains implementation of HIP language on the AMD platform. It is hosted at [clr/hipamd](#).
- `opencl` - contains implementation of OpenCL™ on AMD platform. It is hosted at [clr/opencl](#).
- `rocclr` - contains ROCm compute runtime used in *HIP* and *OpenCL™*. This is hosted at [clr/rocclr](#).

7.2 How to build/install

7.2.1 Prerequisites

Please refer to Quick Start Guide in [ROCm Docs](#).

Building CLR requires `rocm-hip-libraries` meta package, which provides the pre-requisites for CLR.

7.2.2 Linux

- Clone this repository

```
cd clr && mkdir build && cd build
```

- For HIP

```
cmake .. -DCLR_BUILD_HIP=ON -DHIP_COMMON_DIR=$HIP_COMMON_DIR
```

```
`HIP_COMMON_DIR` points to `HIP <https://github.com/ROCm/HIP>`_.
```

- For OpenCL™

```
cmake .. -DCLR_BUILD_OCL=ON
make
make install
```

Users can also build OCL and HIP at the same time by passing `-DCLR_BUILD_HIP=ON -DCLR_BUILD_OCL=ON` to configure command.

For detail instructions, please refer to [build HIP](#).

7.2.3 Test

`hip-tests` is a separate repository hosted at [hip-tests](#).

To run `hip-tests` please go to the repository and follow the steps.

7.2.4 Release notes

HIP provides release notes in CLR [change log](#), which has records of changes in each release.

HIP PROGRAMMING MANUAL

8.1 Host Memory

8.1.1 Introduction

`hipHostMalloc` allocates pinned host memory which is mapped into the address space of all GPUs in the system, the memory can be accessed directly by the GPU device, and can be read or written with much higher bandwidth than pageable memory obtained with functions such as `malloc()`. There are two use cases for this host memory:

- **Faster HostToDevice and DeviceToHost Data Transfers:** The runtime tracks the `hipHostMalloc` allocations and can avoid some of the setup required for regular unpinned memory. For exact measurements on a specific system, experiment with `--unpinned` and `--pinned` switches for the `hipBusBandwidth` tool.
- **Zero-Copy GPU Access:** GPU can directly access the host memory over the CPU/GPU interconnect, without need to copy the data. This avoids the need for the copy, but during the kernel access each memory access must traverse the interconnect, which can be tens of times slower than accessing the GPU's local device memory. Zero-copy memory can be a good choice when the memory accesses are infrequent (perhaps only once). Zero-copy memory is typically "Coherent" and thus not cached by the GPU but this can be overridden if desired.

8.1.2 Memory allocation flags

There are flags parameter which can specify options how to allocate the memory, for example, `hipHostMallocPortable`, the memory is considered allocated by all contexts, not just the one on which the allocation is made. `hipHostMallocMapped`, will map the allocation into the address space for the current device, and the device pointer can be obtained with the API `hipHostGetDevicePointer()`. `hipHostMallocNumaUser` is the flag to allow host memory allocation to follow Numa policy by user. Please note this flag is currently only applicable on Linux, under development on Windows.

All allocation flags are independent, and can be used in any combination without restriction, for instance, `hipHostMalloc` can be called with both `hipHostMallocPortable` and `hipHostMallocMapped` flags set. Both usage models described above use the same allocation flags, and the difference is in how the surrounding code uses the host memory.

8.1.3 Numa-aware host memory allocation

Numa policy determines how memory is allocated. Target of Numa policy is to select a CPU that is closest to each GPU. Numa distance is the measurement of how far between GPU and CPU devices.

By default, each GPU selects a Numa CPU node that has the least Numa distance between them, that is, host memory will be automatically allocated closest on the memory pool of Numa node of the current GPU device. Using `hipSetDevice` API to a different GPU will still be able to access the host allocation, but can have longer Numa distance. Note, Numa policy is so far implemented on Linux, and under development on Windows.

8.1.4 Coherency Controls

ROCm defines two coherency options for host memory:

- Coherent memory : Supports fine-grain synchronization while the kernel is running. For example, a kernel can perform atomic operations that are visible to the host CPU or to other (peer) GPUs. Synchronization instructions include `threadfence_system` and C++11-style atomic operations. In order to achieve this fine-grained coherence, many AMD GPUs use a limited cache policy, such as leaving these allocations uncached by the GPU, or making them read-only.
- Non-coherent memory : Can be cached by GPU, but cannot support synchronization while the kernel is running. Non-coherent memory can be optionally synchronized only at command (end-of-kernel or copy command) boundaries. This memory is appropriate for high-performance access when fine-grain synchronization is not required.

HIP provides the developer with controls to select which type of memory is used via allocation flags passed to `hipHostMalloc` and the `HIP_HOST_COHERENT` environment variable. By default, the environment variable `HIP_HOST_COHERENT` is set to 0 in HIP. The control logic in the current version of HIP is as follows:

- No flags are passed in: the host memory allocation is coherent, the `HIP_HOST_COHERENT` environment variable is ignored.
- `hipHostMallocCoherent=1`: The host memory allocation will be coherent, the `HIP_HOST_COHERENT` environment variable is ignored.
- `hipHostMallocMapped=1`: The host memory allocation will be coherent, the `HIP_HOST_COHERENT` environment variable is ignored.
- `hipHostMallocNonCoherent=1`, `hipHostMallocCoherent=0`, and `hipHostMallocMapped=0`: The host memory will be non-coherent, the `HIP_HOST_COHERENT` environment variable is ignored.
- `hipHostMallocCoherent=0`, `hipHostMallocNonCoherent=0`, `hipHostMallocMapped=0`, but one of the other `HostMalloc` flags is set:
 - If `HIP_HOST_COHERENT` is defined as 1, the host memory allocation is coherent.
 - If `HIP_HOST_COHERENT` is not defined, or defined as 0, the host memory allocation is non-coherent.
- `hipHostMallocCoherent=1`, `hipHostMallocNonCoherent=1`: Illegal.

8.1.5 Visibility of Zero-Copy Host Memory

Coherent host memory is automatically visible at synchronization points. Non-coherent

HIP API	Synchronization Effect	Fence	Coherent Memory Visibility	Host Visibility	Non-Coherent Host Memory Visibility
<code>hipStreamSync</code>	host waits for all commands in the specified stream to complete	system-scope release	yes		yes
<code>hipDeviceSync</code>	host waits for all commands in all streams on the specified device to complete	system-scope release	yes		yes
<code>hipEventSync</code>	host waits for the specified event to complete	device-scope release	yes		depends - see below
<code>hipStreamWait</code>	stream waits for the specified event to complete	none	yes		no

8.1.6 `hipEventSynchronize`

Developers can control the release scope for `hipEvents`:

- By default, the GPU performs a device-scope acquire and release operation with each recorded event. This will make host and device memory visible to other commands executing on the same device.

A stronger system-level fence can be specified when the event is created with `hipEventCreateWithFlags`:

- `hipEventReleaseToSystem`: Perform a system-scope release operation when the event is recorded. This will make both Coherent and Non-Coherent host memory visible to other agents in the system, but may involve heavy-weight operations such as cache flushing. Coherent memory will typically use lighter-weight in-kernel synchronization mechanisms such as an atomic operation and thus does not need to use `hipEventReleaseToSystem`.
- `hipEventDisableTiming`: Events created with this flag will not record profiling data and provide the best performance if used for synchronization.

8.1.7 Summary and Recommendations

- Coherent host memory is the default and is the easiest to use since the memory is visible to the CPU at typical synchronization points. This memory allows in-kernel synchronization commands such as `threadfence_system` to work transparently.
- HIP/ROCm also supports the ability to cache host memory in the GPU using the “Non-Coherent” host memory allocations. This can provide performance benefit, but care must be taken to use the correct synchronization.

8.1.8 Managed memory allocation

Managed memory, including the `__managed__` keyword, is supported in HIP combined host/device compilation, on Linux, not on Windows (under development).

Managed memory, via unified memory allocation, allows data be shared and accessible to both the CPU and GPU using a single pointer. The allocation will be managed by AMD GPU driver using the Linux HMM (Heterogeneous Memory Management) mechanism, the user can call managed memory API `hipMallocManaged` to allocate a large chunk of HMM memory, execute kernels on device and fetch data between the host and device as needed.

In HIP application, it is recommended to do the capability check before calling the managed memory APIs. For example:

```
int managed_memory = 0;
HIPCHECK(hipDeviceGetAttribute(&managed_memory,
    hipDeviceAttributeManagedMemory, p_gpuDevice));

if (!managed_memory) {
    printf ("info: managed memory access not supported on the device %d\n Skipped\n", p_
    ↪ gpuDevice);
}
else {
    HIPCHECK(hipSetDevice(p_gpuDevice));
    HIPCHECK(hipMallocManaged(&Hmm, N * sizeof(T)));
    . . .
}
```

Please note, the managed memory capability check may not be necessary, but if HMM is not supported, then managed malloc will fall back to using system memory and other managed memory API calls will have undefined behavior.

Note, managed memory management is implemented on Linux, not supported on Windows yet.

8.1.9 HIP Stream Memory Operations

HIP supports Stream Memory Operations to enable direct synchronization between Network Nodes and GPU. Following new APIs are added, `hipStreamWaitValue32` `hipStreamWaitValue64` `hipStreamWriteValue32` `hipStreamWriteValue64`

Note, CPU access to the semaphore's memory requires volatile keyword to disable CPU compiler's optimizations on memory access. For more details, please check the documentation `HIP-API.pdf`.

Please note, HIP stream does not guarantee concurrency on AMD hardware for the case of multiple (at least 6) long-running streams executing concurrently, using `hipStreamSynchronize(nullptr)` for synchronization.

8.2 Direct Dispatch

HIP runtime has Direct Dispatch enabled by default in ROCM 4.4 on Linux. With this feature we move away from our conventional producer-consumer model where the runtime creates a worker thread(consumer) for each HIP Stream, and the host thread(producer) enqueues commands to a command queue(per stream).

For Direct Dispatch, HIP runtime would directly enqueue a packet to the AQL queue (user mode queue on GPU) on the Dispatch API call from the application. That has shown to reduce the latency to launch the first wave on the idle GPU and total time of tiny dispatches synchronized with the host.

In addition, eliminating the threads in runtime has reduced the variance in the dispatch numbers as the thread scheduling delays and atomics/locks synchronization latencies are reduced.

This feature can be disabled via setting the following environment variable, `AMD_DIRECT_DISPATCH=0`

Note, Direct Dispatch is implemented on Linux. It is currently not supported on Windows.

8.3 HIP Runtime Compilation

HIP now supports runtime compilation (HIP RTC), the usage of which will provide the possibility of optimizations and performance improvement compared with other APIs via regular offline static compilation.

HIP RTC APIs accept HIP source files in character string format as input parameters and create handles of programs by compiling the HIP source files without spawning separate processes.

For more details on HIP RTC APIs, refer to [HIP Runtime API Reference](#).

For Linux developers, the link [here](#) shows an example how to program HIP application using runtime compilation mechanism, and a detailed *HIP RTC programming guide* is also available.

8.4 HIP Graph

HIP graph is supported. For more details, refer to the HIP API Guide.

8.5 Device-Side Malloc

HIP-Clang now supports device-side malloc and free. This implementation does not require the use of `hipDeviceSetLimit(hipLimitMallocHeapSize, value)` nor respects any setting. The heap is fully dynamic and can grow until the available free memory on the device is consumed.

8.6 Use of Per-thread default stream

The per-thread default stream is supported in HIP. It is an implicit stream local to both the thread and the current device. This means that the command issued to the per-thread default stream by the thread does not implicitly synchronize with other streams (like explicitly created streams), or default per-thread stream on other threads. The per-thread default stream is a blocking stream and will synchronize with the default null stream if both are used in a program. The per-thread default stream can be enabled via adding a compilation option, `-fgpu-default-stream=per-thread`.

And users can explicitly use `hipStreamPerThread` as per-thread default stream handle as input in API commands. There are test codes as examples in the [link](#).

8.7 Use of Long Double Type

In HIP-Clang, long double type is 80-bit extended precision format for x86_64, which is not supported by AMDGPU. HIP-Clang treats long double type as IEEE double type for AMDGPU. Using long double type in HIP source code will not cause issue as long as data of long double type is not transferred between host and device. However, long double type should not be used as kernel argument type.

8.8 Use of _Float16 Type

If a host function is to be used between clang (or hipcc) and gcc for x86_64, i.e. its definition is compiled by one compiler but the caller is compiled by a different compiler, `_Float16` or aggregates containing `_Float16` should not be used as function argument or return type. This is due to lack of stable ABI for `_Float16` on x86_64. Passing `_Float16` or aggregates containing `_Float16` between clang and gcc could cause undefined behavior.

8.9 FMA and contractions

By default HIP-Clang assumes `-ffp-contract=fast-honor-pragmas`. Users can use `#pragma clang fp contract(on|off|fast)` to control `fp` contraction of a block of code. For x86_64, FMA is off by default since the generic x86_64 target does not support FMA by default. To turn on FMA on x86_64, either use `-mfma` or `-march=native` on CPU's supporting FMA.

When contractions are enabled and the CPU has not enabled FMA instructions, the GPU can produce different numerical results than the CPU for expressions that can be contracted. Tolerance should be used for floating point comparisons.

8.10 Math functions with special rounding modes

Note: Currently, HIP only supports basic math functions with rounding modern (round to nearest). HIP does not support basic math functions with rounding modes `ru` (round up), `rd` (round down), and `rz` (round towards zero).

8.11 Creating Static Libraries

HIP-Clang supports generating two types of static libraries. The first type of static library does not export device functions, and only exports and launches host functions within the same library. The advantage of this type is the ability to link with a non-hipcc compiler such as gcc. The second type exports device functions to be linked by other code objects. However, this requires using hipcc as the linker.

In addition, the first type of library contains host objects with device code embedded as fat binaries. It is generated using the flag `-emit-static-lib`. The second type of library contains relocatable device objects and is generated using `ar`.

Here is an example to create and use static libraries:

- Type 1 using `--emit-static-lib`:

```
hipcc hipOptLibrary.cpp --emit-static-lib -fPIC -o libHipOptLibrary.a
gcc test.cpp -L. -lhipOptLibrary -L/path/to/hip/lib -lamdhip64 -o test.out
```

- Type 2 using system `ar`:

```
hipcc hipDevice.cpp -c -fgpu-rtc -o hipDevice.o  
ar rcsD libHipDevice.a hipDevice.o  
hipcc libHipDevice.a test.cpp -fgpu-rtc -o test.out
```

For more information, please see [HIP samples](#) [host functions](#) and [device_functions](#).

HIP PORTING GUIDE

In addition to providing a portable C++ programming environment for GPUs, HIP is designed to ease the porting of existing CUDA code into the HIP environment. This section describes the available tools and provides practical suggestions on how to port CUDA code and work through common issues.

9.1 Porting a New CUDA Project

9.1.1 General Tips

- Starting the port on a CUDA machine is often the easiest approach, since you can incrementally port pieces of the code to HIP while leaving the rest in CUDA. (Recall that on CUDA machines HIP is just a thin layer over CUDA, so the two code types can interoperate on NVCC platforms.) Also, the HIP port can be compared with the original CUDA code for function and performance.
- Once the CUDA code is ported to HIP and is running on the CUDA machine, compile the HIP code using the HIP compiler on an AMD machine.
- HIP ports can replace CUDA versions: HIP can deliver the same performance as a native CUDA implementation, with the benefit of portability to both NVIDIA and AMD architectures as well as a path to future C++ standard support. You can handle platform-specific features through conditional compilation or by adding them to the open-source HIP infrastructure.
- Use `hipconvertinplace-perl.sh` to hipify all code files in the CUDA source directory.

9.1.2 Scanning existing CUDA code to scope the porting effort

The `hipexamine-perl.sh` tool will scan a source directory to determine which files contain CUDA code and how much of that code can be automatically hipified.

```
> cd examples/rodinia_3.0/cuda/kmeans
> $HIP_DIR/bin/hipexamine-perl.sh.
info: hipify ./kmeans.h =====>
info: hipify ./unistd.h =====>
info: hipify ./kmeans.c =====>
info: hipify ./kmeans_cuda_kernel.cu =====>
    info: converted 40 CUDA->HIP refs( dev:0 mem:0 kern:0 builtin:37 math:0 stream:0_
    ↪event:0 err:0 def:0 tex:3 other:0 ) warn:0 LOC:185
info: hipify ./getopt.h =====>
info: hipify ./kmeans_cuda.cu =====>
    info: converted 49 CUDA->HIP refs( dev:3 mem:32 kern:2 builtin:0 math:0 stream:0_
```

(continues on next page)

(continued from previous page)

```

↪event:0 err:0 def:0 tex:12 other:0 ) warn:0 LOC:311
info: hipify ./rmse.c =====>
info: hipify ./cluster.c =====>
info: hipify ./getopt.c =====>
info: hipify ./kmeans_clustering.c =====>
info: TOTAL-converted 89 CUDA->HIP refs( dev:3 mem:32 kern:2 builtin:37 math:0 stream:0
↪event:0 err:0 def:0 tex:15 other:0 ) warn:0 LOC:3607
kernels (1 total) :    kmeansPoint(1)

```

hipexamine-perl scans each code file (cpp, c, h, hpp, etc.) found in the specified directory:

- Files with no CUDA code (kmeans.h) print one line summary just listing the source file name.
- Files with CUDA code print a summary of what was found - for example the kmeans_cuda_kernel.cu file:

```

info: hipify ./kmeans_cuda_kernel.cu =====>
info: converted 40 CUDA->HIP refs( dev:0 mem:0 kern:0 builtin:37 math:0 stream:0
↪event:0

```

- Interesting information in kmeans_cuda_kernel.cu :
 - How many CUDA calls were converted to HIP (40)
 - Breakdown of the CUDA functionality used (dev:0 mem:0 etc). This file uses many CUDA builtins (37) and texture functions (3).
 - Warning for code that looks like CUDA API but was not converted (0 in this file).
 - Count Lines-of-Code (LOC) - 185 for this file.
- hipexamine-perl also presents a summary at the end of the process for the statistics collected across all files. This has similar format to the per-file reporting, and also includes a list of all kernels which have been called. An example from above:

```

info: TOTAL-converted 89 CUDA->HIP refs( dev:3 mem:32 kern:2 builtin:37 math:0 stream:0
↪event:0 err:0 def:0 tex:15 other:0 ) warn:0 LOC:3607
kernels (1 total) :    kmeansPoint(1)

```

9.1.3 Converting a project “in-place”

```
> hipify-perl --inplace
```

For each input file FILE, this script will:

- If FILE.prehip file does not exist, copy the original code to a new file with extension .prehip. Then hipify the code file.
- If FILE.prehip file exists, hipify FILE.prehip and save to FILE.

This is useful for testing improvements to the hipify toolset.

The `hipconvertinplace-perl.sh` script will perform inplace conversion for all code files in the specified directory. This can be quite handy when dealing with an existing CUDA code base since the script preserves the existing directory structure and filenames - and includes work. After converting in-place, you can review the code to add additional parameters to directory names.

```
> hipconvertinplace-perl.sh MY_SRC_DIR
```

9.1.4 Library Equivalents

Most CUDA libraries have a corresponding ROCm library with similar functionality and APIs. However, ROCm also provides HIP marshallng libraries that greatly simplify the porting process because they more precisely reflect their CUDA counterparts and can be used with either the AMD or NVIDIA platforms (see “Identifying HIP Target Platform” below). There are a few notable exceptions:

- MIOpen does not have a marshallng library interface to ease porting from cuDNN.
- RCCL is a drop-in replacement for NCCL and implements the NCCL APIs.
- hipBLASLt does not have a ROCm library but can still target the NVIDIA platform, as needed.
- EIGEN’s HIP support is part of the library.

CUDA Library	Li-brary	HIP Library	Li-brary	ROCm Library	Li-brary	Comment
cuBLAS		hipBLAS		rocBLAS		Basic Linear Algebra Subroutines
cuBLASLt		hip-BLASLt		N/A		Basic Linear Algebra Subroutines, lightweight and new flexible API
cuFFT		hipFFT		rocFFT		Fast Fourier Transfer Library
cuSPARSE		hipSPARSE		rocSPARSE		Sparse BLAS + SPMV
cuSOLVER		hip-SOLVER		rocSOLVER		Lapack library
AmgX		N/A		rocALU-TION		Sparse iterative solvers and preconditioners with algebraic multigrid
Thrust		N/A		rocThrust		C++ parallel algorithms library
CUB		hipCUB		rocPRIM		Low Level Optimized Parallel Primitives
cuDNN		N/A		MIOpen		Deep learning Solver Library
cuRAND		hipRAND		rocRAND		Random Number Generator Library
EIGEN		EIGEN		N/A		C++ template library for linear algebra: matrices, vectors, numerical solvers,
NCCL		N/A		RCCL		Communications Primitives Library based on the MPI equivalents

9.2 Distinguishing Compiler Modes

9.2.1 Identifying HIP Target Platform

All HIP projects target either AMD or NVIDIA platform. The platform affects which headers are included and which libraries are used for linking.

- `HIP_PLATFORM_AMD` is defined if the HIP platform targets AMD. Note, `HIP_PLATFORM_HCC` was previously defined if the HIP platform targeted AMD, it is deprecated.
- `HIP_PLATFORM_NVIDIA` is defined if the HIP platform targets NVIDIA. Note, `HIP_PLATFORM_NVCC` was previously defined if the HIP platform targeted NVIDIA, it is deprecated.

9.2.2 Identifying the Compiler: hip-clang or NVCC

Often, it's useful to know whether the underlying compiler is HIP-Clang or NVCC. This knowledge can guard platform-specific code or aid in platform-specific performance tuning.

```
#ifdef __HIP_PLATFORM_AMD__  
// Compiled with HIP-Clang  
#endif
```

```
#ifdef __HIP_PLATFORM_NVIDIA__  
// Compiled with nvcc  
// Could be compiling with CUDA language extensions enabled (for example, a ".cu file)  
// Could be in pass-through mode to an underlying host compile OR (for example, a .cpp_  
→file)
```

```
#ifdef __CUDACC__  
// Compiled with nvcc (CUDA language extensions enabled)
```

Compiler directly generates the host code (using the Clang x86 target) and passes the code to another host compiler. Thus, they have no equivalent of the `__CUDACC__` define.

9.2.3 Identifying Current Compilation Pass: Host or Device

NVCC makes two passes over the code: one for host code and one for device code. HIP-Clang will have multiple passes over the code: one for the host code, and one for each architecture on the device code. `__HIP_DEVICE_COMPILE__` is set to a nonzero value when the compiler (HIP-Clang or NVCC) is compiling code for a device inside a `__global__` kernel or for a device function. `__HIP_DEVICE_COMPILE__` can replace `#ifdef` checks on the `__CUDA_ARCH__` define.

```
// #ifdef __CUDA_ARCH__  
#if __HIP_DEVICE_COMPILE__
```

Unlike `__CUDA_ARCH__`, the `__HIP_DEVICE_COMPILE__` value is 1 or undefined, and it doesn't represent the feature capability of the target device.

9.2.4 Compiler Defines: Summary

Define	HIP-Clang	NVCC	Other (GCC, ICC, Clang, etc.)
HIP-related defines:			
<code>__HIP_PLATF</code>	Defined	Undefined	Defined if targeting AMD platform; undefined otherwise
<code>__HIP_PLATF</code>	Undefined	Defined	Defined if targeting NVIDIA platform; undefined otherwise
<code>__HIP_DEVIC</code>	1 if compiling for device; undefined if compiling for host	1 if compiling for device; undefined if compiling for host	Undefined
<code>__HIPCC__</code>	Defined	Defined	Undefined
<code>__HIP_ARCH</code>	0 or 1 depending on feature support (see below)	0 or 1 depending on feature support (see below)	0
NVCC-related defines:			
<code>__CUDACC__</code>	Defined if source code is compiled by NVCC; undefined otherwise	Undefined	
<code>__NVCC__</code>	Defined	Undefined	
Undefined			
<code>__CUDA_ARCH</code>	Undefined	Unsigned representing compute capability (e.g., “130”) if in device code; 0 if in host code	Undefined
hip-clang-related defines:			
<code>__HIP__</code>	Defined	Undefined	Undefined
HIP-Clang common defines:			
<code>__clang__</code>	Defined	Defined	Undefined

9.3 Identifying Architecture Features

9.3.1 HIP_ARCH Defines

Some CUDA code tests `__CUDA_ARCH__` for a specific value to determine whether the machine supports a certain architectural feature. For instance,

```
#if (__CUDA_ARCH__ >= 130)
// doubles are supported
```

This type of code requires special attention, since AMD and CUDA devices have different architectural capabilities. Moreover, you can’t determine the presence of a feature using a simple comparison against an architecture’s version

number. HIP provides a set of defines and device properties to query whether a specific architectural feature is supported.

The `__HIP_ARCH_*` defines can replace comparisons of `__CUDA_ARCH__` values:

```
//#if (__CUDA_ARCH__ >= 130)    // non-portable
if __HIP_ARCH_HAS_DOUBLES__ {  // portable HIP feature query
    // doubles are supported
}
```

For host code, the `__HIP_ARCH_*` defines are set to 0. You should only use the `__HIP_ARCH__` fields in device code.

9.3.2 Device-Architecture Properties

Host code should query the architecture feature flags in the device properties that `hipGetDeviceProperties` returns, rather than testing the “major” and “minor” fields directly:

```
hipGetDeviceProperties(&deviceProp, device);
//if ((deviceProp.major == 1 && deviceProp.minor < 2)) // non-portable
if (deviceProp.arch.hasSharedInt32Atomics) {           // portable HIP feature query
    // has shared int32 atomic operations ...
}
```

9.3.3 Table of Architecture Properties

The table below shows the full set of architectural properties that HIP supports.

Define (use only in device code)	Device Property (run-time query)	Comment
32-bit atomics:		
<code>__HIP_ARCH_HAS_GLOBAL_INT32_ATOMICS__</code>	<code>hasGlobalInt32Atomics</code>	32-bit integer atomics for global memory
<code>__HIP_ARCH_HAS_GLOBAL_FLOAT_ATOMIC_EXCHANGE__</code>	<code>hasGlobalFloatAtomicExchange</code>	32-bit float atomic exchange for global memory
<code>__HIP_ARCH_HAS_SHARED_INT32_ATOMICS__</code>	<code>hasSharedInt32Atomics</code>	32-bit integer atomics for shared memory
<code>__HIP_ARCH_HAS_SHARED_FLOAT_ATOMIC_EXCHANGE__</code>	<code>hasSharedFloatAtomicExchange</code>	32-bit float atomic exchange for shared memory
<code>__HIP_ARCH_HAS_FLOAT_ATOMIC_ADD__</code>	<code>hasFloatAtomicAdd</code>	32-bit float atomic add in global and shared memory
64-bit atomics:		
<code>__HIP_ARCH_HAS_GLOBAL_INT64_ATOMICS__</code>	<code>hasGlobalInt64Atomics</code>	64-bit integer atomics for global memory
<code>__HIP_ARCH_HAS_SHARED_INT64_ATOMICS__</code>	<code>hasSharedInt64Atomics</code>	64-bit integer atomics for shared memory
Doubles:		
<code>__HIP_ARCH_HAS_DOUBLES__</code>	<code>hasDoubles</code>	Double-precision floating point
Warp cross-lane operations:		
<code>__HIP_ARCH_HAS_WARP_VOTE__</code>	<code>hasWarpVote</code>	Warp vote instructions (any, all)
<code>__HIP_ARCH_HAS_WARP_BALLOT__</code>	<code>hasWarpBallot</code>	Warp ballot instructions
<code>__HIP_ARCH_HAS_WARP_SHUFFLE__</code>	<code>hasWarpShuffle</code>	Warp shuffle operations (<code>shfl_*</code>)
<code>__HIP_ARCH_HAS_WARP_FUNNEL_SHIFT__</code>	<code>hasFunnelShift</code>	Funnel shift two input words into one
Sync:		
<code>__HIP_ARCH_HAS_THREAD_FENCE_SYSTEM__</code>	<code>hasThreadFenceSystem</code>	<code>threadfence_system</code>
<code>__HIP_ARCH_HAS_SYNC_THREAD_EXT__</code>	<code>hasSyncThreadsExt</code>	<code>syncthreads_count</code> , <code>syncthreads_and</code> , <code>syncthreads_or</code>
Miscellaneous:		
<code>__HIP_ARCH_HAS_SURFACE_FUNCS__</code>	<code>hasSurfaceFuncs</code>	
<code>__HIP_ARCH_HAS_3DGRID__</code>	<code>has3dGrid</code>	Grids and groups are 3D
<code>__HIP_ARCH_HAS_DYNAMIC_PARALLELISM__</code>	<code>hasDynamicParallelism</code>	

9.4 Finding HIP

Makefiles can use the following syntax to conditionally provide a default `HIP_PATH` if one does not exist:

```
HIP_PATH ?= $(shell hipconfig --path)
```

9.5 Identifying HIP Runtime

HIP can depend on rocclr, or CUDA as runtime

- **AMD platform** On AMD platform, HIP uses ROCm Compute Language Runtime, called ROCclr. ROCclr is a virtual device interface that HIP runtimes interact with different backends which allows runtimes to work on Linux , as well as Windows without much efforts.
- **NVIDIA platform** On NVIDIA platform, HIP is just a thin layer on top of CUDA. On non-AMD platform, HIP runtime determines if CUDA is available and can be used. If available, HIP_PLATFORM is set to nvidia and underneath CUDA path is used.

9.6 hipLaunchKernelGGL

hipLaunchKernelGGL is a macro that can serve as an alternative way to launch kernel, which accepts parameters of launch configurations (grid dims, group dims, stream, dynamic shared size) followed by a variable number of kernel arguments. It can replace <<< >>>, if the user so desires.

9.7 Compiler Options

hipcc is a portable compiler driver that will call NVCC or HIP-Clang (depending on the target system) and attach all required include and library options. It passes options through to the target compiler. Tools that call hipcc must ensure the compiler options are appropriate for the target compiler. The hipconfig script may helpful in identifying the target platform, compiler and runtime. It can also help set options appropriately.

9.7.1 Compiler options supported on AMD platforms

Here are the main compiler options supported on AMD platforms by HIP-Clang.

Option	Description
--amdgpu-t.	[DEPRECATED] This option is being replaced by --offload-arch=<target>. Generate code for the given GPU target. Supported targets are gfx701, gfx801, gfx802, gfx803, gfx900, gfx906, gfx908, gfx1010, gfx1011, gfx1012, gfx1030, gfx1031. This option could appear multiple times on the same command line to generate a fat binary for multiple targets.
--fgpu-rdc	Generate relocatable device code, which allows kernels or device functions calling device functions in different translation units.
-ggdb	Equivalent to -g plus tuning for GDB. This is recommended when using ROCm's GDB to debug GPU code.
--gpu-max-	Generate code to support up to the specified number of threads per block.
-O<n>	Specify the optimization level.
--offload-a.	Specify the AMD GPU target ID .
-save-temp	Save the compiler generated intermediate files.
-v	Show the compilation steps.

9.8 Linking Issues

9.8.1 Linking With hipcc

hipcc adds the necessary libraries for HIP as well as for the accelerator compiler (NVCC or AMD compiler). We recommend linking with hipcc since it automatically links the binary to the necessary HIP runtime libraries. It also has knowledge on how to link and to manage the GPU objects.

9.8.2 `-lm` Option

hipcc adds `-lm` by default to the link command.

9.9 Linking Code With Other Compilers

CUDA code often uses NVCC for accelerator code (defining and launching kernels, typically defined in `.cu` or `.cuh` files). It also uses a standard compiler (g++) for the rest of the application. NVCC is a preprocessor that employs a standard host compiler (gcc) to generate the host code. Code compiled using this tool can employ only the intersection of language features supported by both NVCC and the host compiler. In some cases, you must take care to ensure the data types and alignment of the host compiler are identical to those of the device compiler. Only some host compilers are supported—for example, recent NVCC versions lack Clang host-compiler capability.

HIP-Clang generates both device and host code using the same Clang-based compiler. The code uses the same API as gcc, which allows code generated by different gcc-compatible compilers to be linked together. For example, code compiled using HIP-Clang can link with code compiled using “standard” compilers (such as gcc, ICC and Clang). Take care to ensure all compilers use the same standard C++ header and library formats.

9.9.1 `libc++` and `libstdc++`

hipcc links to `libstdc++` by default. This provides better compatibility between g++ and HIP.

If you pass `--stdlib=libc++` to hipcc, hipcc will use the `libc++` library. Generally, `libc++` provides a broader set of C++ features while `libstdc++` is the standard for more compilers (notably including g++).

When cross-linking C++ code, any C++ functions that use types from the C++ standard library (including `std::string`, `std::vector` and other containers) must use the same standard-library implementation. They include the following:

- Functions or kernels defined in HIP-Clang that are called from a standard compiler
- Functions defined in a standard compiler that are called from HIP-Clang.

Applications with these interfaces should use the default `libstdc++` linking.

Applications which are compiled entirely with hipcc, and which benefit from advanced C++ features not supported in `libstdc++`, and which do not require portability to NVCC, may choose to use `libc++`.

9.9.2 HIP Headers (`hip_runtime.h`, `hip_runtime_api.h`)

The `hip_runtime.h` and `hip_runtime_api.h` files define the types, functions and enumerations needed to compile a HIP program:

- `hip_runtime_api.h`: defines all the HIP runtime APIs (e.g., `hipMalloc`) and the types required to call them. A source file that is only calling HIP APIs but neither defines nor launches any kernels can include `hip_runtime_api.h`. `hip_runtime_api.h` uses no custom Heterogeneous Compute (HC) language features and can be compiled using a standard C++ compiler.
- `hip_runtime.h`: included in `hip_runtime_api.h`. It additionally provides the types and defines required to create and launch kernels. `hip_runtime.h` can be compiled using a standard C++ compiler but will expose a subset of the available functions.

CUDA has slightly different contents for these two files. In some cases you may need to convert hipified code to include the richer `hip_runtime.h` instead of `hip_runtime_api.h`.

9.9.3 Using a Standard C++ Compiler

You can compile `hip_runtime_api.h` using a standard C or C++ compiler (e.g., `gcc` or `ICC`). The HIP include paths and defines (`__HIP_PLATFORM_AMD__` or `__HIP_PLATFORM_NVIDIA__`) must pass to the standard compiler; `hipconfig` then returns the necessary options:

```
> hipconfig --cxx_config
-D__HIP_PLATFORM_AMD__ -I/home/user1/hip/include
```

You can capture the `hipconfig` output and passed it to the standard compiler; below is a sample makefile syntax:

```
CPPFLAGS += $(shell $(HIP_PATH)/bin/hipconfig --cpp_config)
```

NVCC includes some headers by default. However, HIP does not include default headers, and instead all required files must be explicitly included. Specifically, files that call HIP run-time APIs or define HIP kernels must explicitly include the appropriate HIP headers. If the compilation process reports that it cannot find necessary APIs (for example, `error: identifier hipSetDevice is undefined`), ensure that the file includes `hip_runtime.h` (or `hip_runtime_api.h`, if appropriate). The `hipify-perl` script automatically converts `cuda_runtime.h` to `hip_runtime.h`, and it converts `cuda_runtime_api.h` to `hip_runtime_api.h`, but it may miss nested headers or macros.

9.9.3.1 `cuda.h`

The HIP-Clang path provides an empty `cuda.h` file. Some existing CUDA programs include this file but don't require any of the functions.

9.9.4 Choosing HIP File Extensions

Many existing CUDA projects use the `.cu` and `.cuh` file extensions to indicate code that should be run through the NVCC compiler. For quick HIP ports, leaving these file extensions unchanged is often easier, as it minimizes the work required to change file names in the directory and `#include` statements in the files.

For new projects or ports which can be re-factored, we recommend the use of the extension `.hip.cpp` for source files, and `.hip.h` or `.hip.hpp` for header files. This indicates that the code is standard C++ code, but also provides a unique indication for make tools to run `hipcc` when appropriate.

9.10 Workarounds

9.10.1 warpSize

Code should not assume a warp size of 32 or 64. See [Warp Cross-Lane Functions](#) for information on how to write portable wave-aware code.

9.10.2 Kernel launch with group size > 256

Kernel code should use `__attribute__((amdgpu_flat_work_group_size(<min>,<max>)))`.

For example:

```
__global__ void dot(double *a, double *b, const int n) __attribute__((amdgpu_flat_work_
↪group_size(1, 512)))
```

9.11 memcpyToSymbol

HIP support for `hipMemcpyToSymbol` is complete. This feature allows a kernel to define a device-side data symbol which can be accessed on the host side. The symbol can be in `__constant` or device space.

Note that the symbol name needs to be encased in the `HIP_SYMBOL` macro, as shown in the code example below. This also applies to `hipMemcpyFromSymbol`, `hipGetSymbolAddress`, and `hipGetSymbolSize`.

For example:

Device Code:

```
#include<hip/hip_runtime.h>
#include<hip/hip_runtime_api.h>
#include<iostream>

#define HIP_ASSERT(status) \
    assert(status == hipSuccess)

#define LEN 512
#define SIZE 2048

__constant__ int Value[LEN];

__global__ void Get(hipLaunchParm lp, int *Ad)
{
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
    Ad[tid] = Value[tid];
}

int main()
{
    int *A, *B, *Ad;
    A = new int[LEN];
    B = new int[LEN];
    for(unsigned i=0; i<LEN; i++)
```

(continues on next page)

(continued from previous page)

```

{
    A[i] = -1*i;
    B[i] = 0;
}

HIP_ASSERT(hipMalloc((void**)&Ad, SIZE));

HIP_ASSERT(hipMemcpyToSymbol(HIP_SYMBOL(Value), A, SIZE, 0, hipMemcpyHostToDevice));
hipLaunchKernelGGL(Get, dim3(1,1,1), dim3(LEN,1,1), 0, 0, Ad);
HIP_ASSERT(hipMemcpy(B, Ad, SIZE, hipMemcpyDeviceToHost));

for(unsigned i=0;i<LEN;i++)
{
    assert(A[i] == B[i]);
}
std::cout<<"Passed"<<std::endl;
}

```

9.12 CU_POINTER_ATTRIBUTE_MEMORY_TYPE

To get pointer's memory type in HIP/HIP-Clang, developers should use `hipPointerGetAttributes` API. First parameter of the API is `hipPointerAttribute_t` which has 'type' as member variable. 'type' indicates input pointer is allocated on device or host.

For example:

```

double * ptr;
hipMalloc(reinterpret_cast<void**>(&ptr), sizeof(double));
hipPointerAttribute_t attr;
hipPointerGetAttributes(&attr, ptr); /*attr.type will have value as hipMemoryTypeDevice*/

double* ptrHost;
hipHostMalloc(&ptrHost, sizeof(double));
hipPointerAttribute_t attr;
hipPointerGetAttributes(&attr, ptrHost); /*attr.type will have value as
↳hipMemoryTypeHost*/

```

Please note, `hipMemoryType` enum values are different from `cudaMemoryType` enum values.

For example, on AMD platform, `hipMemoryType` is defined in `hip_runtime_api.h`,

```

typedef enum hipMemoryType {
    hipMemoryTypeHost = 0,    ///< Memory is physically located on host
    hipMemoryTypeDevice = 1,  ///< Memory is physically located on device. (see deviceId
↳for specific device)
    hipMemoryTypeArray = 2,   ///< Array memory, physically located on device. (see
↳deviceId for specific device)
    hipMemoryTypeUnified = 3,  ///< Not used currently
    hipMemoryTypeManaged = 4  ///< Managed memory, automaticlly managed by the unified
↳memory system
} hipMemoryType;

```


Looking into CUDA toolkit, it defines `cudaMemoryType` as following,

```
enum cudaMemoryType
{
    cudaMemoryTypeUnregistered = 0, // Unregistered memory.
    cudaMemoryTypeHost = 1, // Host memory.
    cudaMemoryTypeDevice = 2, // Device memory.
    cudaMemoryTypeManaged = 3, // Managed memory
}
```

In this case, memory type translation for `hipPointerGetAttributes` needs to be handled properly on NVIDIA platform to get the correct memory type in CUDA, which is done in the file `nvidia_hip_runtime_api.h`.

So in any HIP applications which use HIP APIs involving memory types, developers should use `#ifdef` in order to assign the correct enum values depending on NVIDIA or AMD platform.

As an example, please see the code from the [link](#).

With the `#ifdef` condition, HIP APIs work as expected on both AMD and NVIDIA platforms.

Note, `cudaMemoryTypeUnregistered` is currently not supported in `hipMemoryType` enum, due to HIP functionality backward compatibility.

9.13 threadfence_system

`threadfence_system` makes all device memory writes, all writes to mapped host memory, and all writes to peer memory visible to CPU and other GPU devices. Some implementations can provide this behavior by flushing the GPU L2 cache. HIP/HIP-Clang does not provide this functionality. As a workaround, users can set the environment variable `HSA_DISABLE_CACHE=1` to disable the GPU L2 cache. This will affect all accesses and for all kernels and so may have a performance impact.

9.13.1 Textures and Cache Control

Compute programs sometimes use textures either to access dedicated texture caches or to use the texture-sampling hardware for interpolation and clamping. The former approach uses simple point samplers with linear interpolation, essentially only reading a single point. The latter approach uses the sampler hardware to interpolate and combine multiple samples. AMD hardware, as well as recent competing hardware, has a unified texture/L1 cache, so it no longer has a dedicated texture cache. But the NVCC path often caches global loads in the L2 cache, and some programs may benefit from explicit control of the L1 cache contents. We recommend the `__ldg` instruction for this purpose.

AMD compilers currently load all data into both the L1 and L2 caches, so `__ldg` is treated as a no-op.

We recommend the following for functional portability:

- For programs that use textures only to benefit from improved caching, use the `__ldg` instruction
- Programs that use texture object and reference APIs, work well on HIP

9.14 More Tips

9.14.1 HIP Logging

On an AMD platform, set the AMD_LOG_LEVEL environment variable to log HIP application execution information.

The value of the setting controls different logging level,

```
enum LogLevel {
LOG_NONE = 0,
LOG_ERROR = 1,
LOG_WARNING = 2,
LOG_INFO = 3,
LOG_DEBUG = 4
};
```

Logging mask is used to print types of functionalities during the execution of HIP application. It can be set as one of the following values,

```
enum LogMask {
LOG_API      = 1,      //!< (0x1)      API call
LOG_CMD      = 2,      //!< (0x2)      Kernel and Copy Commands and Barriers
LOG_WAIT     = 4,      //!< (0x4)      Synchronization and waiting for commands to_
↳ finish
LOG_AQL      = 8,      //!< (0x8)      Decode and display AQL packets
LOG_QUEUE    = 16,     //!< (0x10)     Queue commands and queue contents
LOG_SIG      = 32,     //!< (0x20)     Signal creation, allocation, pool
LOG_LOCK     = 64,     //!< (0x40)     Locks and thread-safety code.
LOG_KERN     = 128,    //!< (0x80)     Kernel creations and arguments, etc.
LOG_COPY     = 256,    //!< (0x100)    Copy debug
LOG_COPY2    = 512,    //!< (0x200)    Detailed copy debug
LOG_RESOURCE = 1024,   //!< (0x400)    Resource allocation, performance-impacting_
↳ events.
LOG_INIT     = 2048,   //!< (0x800)    Initialization and shutdown
LOG_MISC     = 4096,   //!< (0x1000)   Misc debug, not yet classified
LOG_AQL2     = 8192,   //!< (0x2000)   Show raw bytes of AQL packet
LOG_CODE     = 16384,  //!< (0x4000)   Show code creation debug
LOG_CMD2     = 32768,  //!< (0x8000)   More detailed command info, including barrier_
↳ commands
LOG_LOCATION = 65536,  //!< (0x10000)  Log message location
LOG_MEM      = 131072, //!< (0x20000)  Memory allocation
LOG_MEM_POOL = 262144, //!< (0x40000)  Memory pool allocation, including memory in_
↳ graphs
LOG_ALWAYS   = -1      //!< (0xFFFFFFFF) Log always even mask flag is zero
};
```

9.14.2 Debugging hipcc

To see the detailed commands that hipcc issues, set the environment variable HIPCC_VERBOSE to 1. Doing so will print to `stderr` the HIP-clang (or NVCC) commands that hipcc generates.

```
export HIPCC_VERBOSE=1
make
...
hipcc-cmd: /opt/rocm/bin/hipcc --offload-arch=native -x hip backprop_cuda.cu
```

9.14.3 Editor Highlighting

See the `utils/vim` or `utils/gedit` directories to add handy highlighting to hip files.

PORTING CUDA DRIVER API

10.1 Introduction to the CUDA Driver and Runtime APIs

CUDA provides a separate CUDA Driver and Runtime APIs. The two APIs have significant overlap in functionality:

- Both APIs support events, streams, memory management, memory copy, and error handling.
- Both APIs deliver similar performance.
- Driver APIs calls begin with the prefix `cu` while Runtime APIs begin with the prefix `cuda`. For example, the Driver API contains `cuEventCreate` while the Runtime API contains `cudaEventCreate`, with similar functionality.
- The Driver API defines a different but largely overlapping error code space than the Runtime API, and uses a different coding convention. For example, Driver API defines `CUDA_ERROR_INVALID_VALUE` while the Runtime API defines `cudaErrorInvalidValue`

The Driver API offers two additional pieces of functionality not provided by the Runtime API: `cuModule` and `cuCtx` APIs.

10.1.1 `cuModule` API

The Module section of the Driver API provides additional control over how and when accelerator code objects are loaded. For example, the driver API allows code objects to be loaded from files or memory pointers. Symbols for kernels or global data can be extracted from the loaded code objects. In contrast, the Runtime API automatically loads and (if necessary) compiles all of the kernels from an executable binary when run. In this mode, NVCC must be used to compile kernel code so the automatic loading can function correctly.

Both Driver and Runtime APIs define a function for launching kernels (called `cuLaunchKernel` or `cudaLaunchKernel`). The kernel arguments and the execution configuration (grid dimensions, group dimensions, dynamic shared memory, and stream) are passed as arguments to the launch function. The Runtime additionally provides the `<<< >>>` syntax for launching kernels, which resembles a special function call and is easier to use than explicit launch API (in particular with respect to handling of kernel arguments). However, this syntax is not standard C++ and is available only when NVCC is used to compile the host code.

The Module features are useful in an environment which generates the code objects directly, such as a new accelerator language front-end. Here, NVCC is not used. Instead, the environment may have a different kernel language or different compilation flow. Other environments have many kernels and do not want them to be all loaded automatically. The Module functions can be used to load the generated code objects and launch kernels. As we will see below, HIP defines a Module API which provides similar explicit control over code object management.

10.1.2 cuCtx API

The Driver API defines “Context” and “Devices” as separate entities. Contexts contain a single device, and a device can theoretically have multiple contexts. Each context contains a set of streams and events specific to the context. Historically contexts also defined a unique address space for the GPU, though this may no longer be the case in Unified Memory platforms (since the CPU and all the devices in the same process share a single unified address space). The Context APIs also provide a mechanism to switch between devices, which allowed a single CPU thread to send commands to different GPUs. HIP as well as a recent versions of CUDA Runtime provide other mechanisms to accomplish this feat - for example using streams or `cudaSetDevice`.

The CUDA Runtime API unifies the Context API with the Device API. This simplifies the APIs and has little loss of functionality since each Context can contain a single device, and the benefits of multiple contexts has been replaced with other interfaces. HIP provides a context API to facilitate easy porting from existing Driver codes. In HIP, the Ctx functions largely provide an alternate syntax for changing the active device.

Most new applications will prefer to use `hipSetDevice` or the stream APIs, therefore HIP has marked `hipCtx` APIs as **deprecated**. Support for these APIs may not be available in future releases. For more details on deprecated APIs please refer [HIP deprecated APIs](#).

10.2 HIP Module and Ctx APIs

Rather than present two separate APIs, HIP extends the HIP API with new APIs for Modules and Ctx control.

10.2.1 hipModule API

Like the CUDA Driver API, the Module API provides additional control over how code is loaded, including options to load code from files or from in-memory pointers. NVCC and HIP-Clang target different architectures and use different code object formats: NVCC is `cubin` or `ptx` files, while the HIP-Clang path is the `hsaco` format. The external compilers which generate these code objects are responsible for generating and loading the correct code object for each platform. Notably, there is not a fat binary format that can contain code for both NVCC and HIP-Clang platforms. The following table summarizes the formats used on each platform:

Format	APIs	NVCC	HIP-CLANG
Code Object	<code>hipModuleLoad</code> , <code>hipModuleLoadData</code>	<code>.cubin</code> or <code>PTX</code> text	<code>.hsaco</code>
Fat Binary	<code>hipModuleLoadFatBin</code>	<code>.fatbin</code>	<code>.hip_fatbin</code>

`hipcc` uses HIP-Clang or NVCC to compile host codes. Both of these may embed code objects into the final executable, and these code objects will be automatically loaded when the application starts. The `hipModule` API can be used to load additional code objects, and in this way provides an extended capability to the automatically loaded code objects. HIP-Clang allows both of these capabilities to be used together, if desired. Of course it is possible to create a program with no kernels and thus no automatic loading.

10.2.2 hipCtx API

HIP provides a Ctx API as a thin layer over the existing Device functions. This Ctx API can be used to set the current context, or to query properties of the device associated with the context. The current context is implicitly used by other APIs such as `hipStreamCreate`.

10.2.3 hipify translation of CUDA Driver API

The HIPIFY tools convert CUDA Driver APIs for streams, events, modules, devices, memory management, context, profiler to the equivalent HIP driver calls. For example, `cuEventCreate` will be translated to `hipEventCreate`. HIPIFY tools also convert error codes from the Driver namespace and coding convention to the equivalent HIP error code. Thus, HIP unifies the APIs for these common functions.

The memory copy API requires additional explanation. The CUDA driver includes the memory direction in the name of the API (`cuMemcpyH2D`) while the CUDA driver API provides a single memory copy API with a parameter that specifies the direction and additionally supports a “default” direction where the runtime determines the direction automatically. HIP provides APIs with both styles: for example, `hipMemcpyH2D` as well as `hipMemcpy`. The first flavor may be faster in some cases since they avoid host overhead to detect the different memory directions.

HIP defines a single error space, and uses camel-case for all errors (i.e. `hipErrorInvalidValue`).

10.2.3.1 Address Spaces

HIP-Clang defines a process-wide address space where the CPU and all devices allocate addresses from a single unified pool. Thus addresses may be shared between contexts, and unlike the original CUDA definition a new context does not create a new address space for the device.

10.2.3.2 Using `hipModuleLaunchKernel`

`hipModuleLaunchKernel` is `cuLaunchKernel` in HIP world. It takes the same arguments as `cuLaunchKernel`.

10.2.3.3 Additional Information

- HIP-Clang creates a primary context when the HIP API is called. So in a pure driver API code, HIP-Clang will create a primary context while HIP/NVCC will have empty context stack. HIP-Clang will push primary context to context stack when it is empty. This can have subtle differences on applications which mix the runtime and driver APIs.

10.2.4 hip-clang Implementation Notes

10.2.4.1 `.hip_fatbin`

hip-clang links device code from different translation units together. For each device target, a code object is generated. Code objects for different device targets are bundled by `clang-offload-bundler` as one fatbinary, which is embedded as a global symbol `__hip_fatbin` in the `.hip_fatbin` section of the ELF file of the executable or shared object.

10.2.4.2 Initialization and Termination Functions

hip-clang generates initialization and termination functions for each translation unit for host code compilation. The initialization functions call `__hipRegisterFatBinary` to register the fatbinary embedded in the ELF file. They also call `__hipRegisterFunction` and `__hipRegisterVar` to register kernel functions and device side global variables. The termination functions call `__hipUnregisterFatBinary`. hip-clang emits a global variable `__hip_gpubin_handle` of `void**` type with `linkonce` linkage and initial value 0 for each host translation unit. Each initialization function checks `__hip_gpubin_handle` and register the fatbinary only if `__hip_gpubin_handle` is 0 and saves the return value of `__hip_gpubin_handle` to `__hip_gpubin_handle`. This is to guarantee that the fatbinary is only registered once. Similar check is done in the termination functions.

10.2.4.3 Kernel Launching

hip-clang supports kernel launching by CUDA `<<<>>>` syntax, `hipLaunchKernelGGL`. The latter one is macro which expand to CUDA `<<<>>>` syntax.

When the executable or shared library is loaded by the dynamic linker, the initialization functions are called. In the initialization functions, when `__hipRegisterFatBinary` is called, the code objects containing all kernels are loaded; when `__hipRegisterFunction` is called, the stub functions are associated with the corresponding kernels in code objects.

hip-clang implements two sets of kernel launching APIs.

By default, in the host code, for the `<<<>>>` statement, hip-clang first emits call of `hipConfigureCall` to set up the threads and grids, then emits call of the stub function with the given arguments. In the stub function, `hipSetupArgument` is called for each kernel argument, then `hipLaunchByPtr` is called with a function pointer to the stub function. In `hipLaunchByPtr`, the real kernel associated with the stub function is launched.

10.2.5 NVCC Implementation Notes

10.2.5.1 Interoperation between HIP and CUDA Driver

CUDA applications may want to mix CUDA driver code with HIP code (see example below). This table shows the type equivalence to enable this interaction.

HIP Type	CU Driver Type	CUDA Runtime Type
<code>hipModule_t</code>	<code>CUmodule</code>	
<code>hipFunction_t</code>	<code>CUfunction</code>	
<code>hipCtx_t</code>	<code>CUcontext</code>	
<code>hipDevice_t</code>	<code>CUdevice</code>	
<code>hipStream_t</code>	<code>CUstream</code>	<code>cudaStream_t</code>
<code>hipEvent_t</code>	<code>CUevent</code>	<code>cudaEvent_t</code>
<code>hipArray</code>	<code>CUarray</code>	<code>cudaArray</code>

10.2.5.2 Compilation Options

The `hipModule_t` interface does not support `cuModuleLoadDataEx` function, which is used to control PTX compilation options. HIP-Clang does not use PTX and does not support these compilation options. In fact, HIP-Clang code objects always contain fully compiled ISA and do not require additional compilation as a part of the load step. The corresponding HIP function `hipModuleLoadDataEx` behaves as `hipModuleLoadData` on HIP-Clang path (compilation options are not used) and as `cuModuleLoadDataEx` on NVCC path. For example (CUDA):

```
CUmodule module;
void *imagePtr = ...; // Somehow populate data pointer with code object

const int numOptions = 1;
CUjit_option options[numOptions];
void * optionValues[numOptions];

options[0] = CU_JIT_MAX_REGISTERS;
unsigned maxRegs = 15;
optionValues[0] = (void*)&maxRegs;

cuModuleLoadDataEx(module, imagePtr, numOptions, options, optionValues);

CUfunction k;
cuModuleGetFunction(&k, module, "myKernel");
```

HIP:

```
hipModule_t module;
void *imagePtr = ...; // Somehow populate data pointer with code object

const int numOptions = 1;
hipJitOption options[numOptions];
void * optionValues[numOptions];

options[0] = hipJitOptionMaxRegisters;
unsigned maxRegs = 15;
optionValues[0] = (void*)&maxRegs;

// hipModuleLoadData(module, imagePtr) will be called on HIP-Clang path, JIT options_
↳ will not be used, and
// cupModuleLoadDataEx(module, imagePtr, numOptions, options, optionValues) will be_
↳ called on NVCC path
hipModuleLoadDataEx(module, imagePtr, numOptions, options, optionValues);

hipFunction_t k;
hipModuleGetFunction(&k, module, "myKernel");
```

The below sample shows how to use `hipModuleGetFunction`.

```
#include<hip_runtime.h>
#include<hip_runtime_api.h>
#include<iostream>
#include<fstream>
#include<vector>
```

(continues on next page)

(continued from previous page)

```

#define LEN 64
#define SIZE LEN<<2

#ifdef __HIP_PLATFORM_AMD__
#define fileName "vcpy_isa.co"
#endif

#ifdef __HIP_PLATFORM_NVIDIA__
#define fileName "vcpy_isa.ptx"
#endif

#define kernel_name "hello_world"

int main(){
    float *A, *B;
    hipDeviceptr_t Ad, Bd;
    A = new float[LEN];
    B = new float[LEN];

    for(uint32_t i=0;i<LEN;i++){
        A[i] = i*1.0f;
        B[i] = 0.0f;
        std::cout<<A[i] << " " <<B[i]<<std::endl;
    }

#ifdef __HIP_PLATFORM_NVIDIA__
    hipInit(0);
    hipDevice_t device;
    hipCtx_t context;
    hipDeviceGet(&device, 0);
    hipCtxCreate(&context, 0, device);
#endif

    hipMalloc((void**)&Ad, SIZE);
    hipMalloc((void**)&Bd, SIZE);

    hipMemcpyHtoD(Ad, A, SIZE);
    hipMemcpyHtoD(Bd, B, SIZE);
    hipModule_t Module;
    hipFunction_t Function;
    hipModuleLoad(&Module, fileName);
    hipModuleGetFunction(&Function, Module, kernel_name);

    std::vector<void*>argBuffer(2);
    memcpy(&argBuffer[0], &Ad, sizeof(void*));
    memcpy(&argBuffer[1], &Bd, sizeof(void*));

    size_t size = argBuffer.size()*sizeof(void*);

    void *config[] = {
        HIP_LAUNCH_PARAM_BUFFER_POINTER, &argBuffer[0],

```

(continues on next page)

(continued from previous page)

```

    HIP_LAUNCH_PARAM_BUFFER_SIZE, &size,
    HIP_LAUNCH_PARAM_END
};

hipModuleLaunchKernel(Function, 1, 1, 1, LEN, 1, 1, 0, 0, NULL, (void**)&config);

hipMemcpyDtoH(B, Bd, SIZE);
for(uint32_t i=0;i<LEN;i++){
    std::cout<<A[i]<<" - "<<B[i]<<std::endl;
}

#ifdef __HIP_PLATFORM_NVIDIA__
    hipCtxDetach(context);
#endif

    return 0;
}

```

10.3 HIP Module and Texture Driver API

HIP supports texture driver APIs however texture reference should be declared in host scope. Following code explains the use of texture reference for __HIP_PLATFORM_AMD__ platform.

```

// Code to generate code object

#include "hip/hip_runtime.h"
extern texture<float, 2, hipReadModeElementType> tex;

__global__ void tex2dKernel(hipLaunchParm lp, float* outputData,
                           int width,
                           int height)
{
    int x = blockIdx.x*blockDim.x + threadIdx.x;
    int y = blockIdx.y*blockDim.y + threadIdx.y;
    outputData[y*width + x] = tex2D(tex, x, y);
}

```

```

// Host code:

texture<float, 2, hipReadModeElementType> tex;

void myFunc ()
{
    // ...

    textureReference* texref;
    hipModuleGetTexRef(&texref, Module1, "tex");
    hipTexRefSetAddressMode(texref, 0, hipAddressModeWrap);
    hipTexRefSetAddressMode(texref, 1, hipAddressModeWrap);
    hipTexRefSetFilterMode(texref, hipFilterModePoint);
}

```

(continues on next page)

(continued from previous page)

```
hipTexRefSetFlags(texref, 0);
hipTexRefSetFormat(texref, HIP_AD_FORMAT_FLOAT, 1);
hipTexRefSetArray(texref, array, HIP_TRSA_OVERRIDE_FORMAT);

// ...
}
```

PROGRAMMING FOR HIP RUNTIME COMPILER (RTC)

HIP lets you compile kernels at runtime with the `hiprtc*` APIs. Kernels can be stored as a text string and can be passed to HIPRTC APIs alongside options to guide the compilation.

NOTE:

- This library can be used on systems without HIP installed nor AMD GPU driver installed at all (offline compilation). Therefore, it does not depend on any HIP runtime library.
- But it does depend on COMGr. You may try to statically link COMGr into HIPRTC to avoid any ambiguity.
- Developers can decide to bundle this library with their application.

11.1 Compilation APIs

To use HIPRTC functionality, HIPRTC header needs to be included first. `#include <hip/hiprtc.h>`

Kernels can be stored in a string:

```
static constexpr auto kernel_source {
R"(
    extern "C"
    __global__ void vector_add(float* output, float* input1, float* input2, size_t size)
    ↪{
        int i = threadIdx.x;
        if (i < size) {
            output[i] = input1[i] + input2[i];
        }
    }
)"};
```

Now to compile this kernel, it needs to be associated with `hiprtcProgram` type, which is done by declaring `hiprtcProgram prog`; and associating the string of kernel with this program:

```
hiprtcCreateProgram(&prog,                // HIPRTC program handle
                  kernel_source,          // HIP kernel source string
                  "vector_add.cpp",       // Name of the HIP program, can be null or an
    ↪empty string
                  0,                      // Number of headers
                  NULL,                   // Header sources
                  NULL);                  // Name of header files
```

hiprtcCreateProgram API also allows you to add headers which can be included in your RTC program. For on-line compilation, the compiler pre-defines HIP device API functions, HIP specific types and macros for device compilation, but does not include standard C/C++ headers by default. Users can only include header files provided to hiprtcCreateProgram.

After associating the kernel string with hiprtcProgram, you can now compile this program using:

```
hiprtcCompileProgram(prog,      // hiprtcProgram
                    0,          // Number of options
                    options);   // Clang Options [Supported Clang Options](clang_options.
↪md)
```

hiprtcCompileProgram returns a status value which can be converted to string via hiprtcGetErrorString. If compilation is successful, hiprtcCompileProgram will return HIPRTC_SUCCESS.

If the compilation fails, you can look up the logs via:

```
size_t logSize;
hiprtcGetProgramLogSize(prog, &logSize);

if (logSize) {
    string log(logSize, '\\0');
    hiprtcGetProgramLog(prog, &log[0]);
    // Corrective action with logs
}
```

If the compilation is successful, you can load the compiled binary in a local variable.

```
size_t codeSize;
hiprtcGetCodeSize(prog, &codeSize);

vector<char> kernel_binary(codeSize);
hiprtcGetCode(prog, kernel_binary.data());
```

After loading the binary, hiprtcProgram can be destroyed. hiprtcDestroyProgram(&prog);

The binary present in kernel_binary can now be loaded via hipModuleLoadData API.

```
hipModule_t module;
hipFunction_t kernel;

hipModuleLoadData(&module, kernel_binary.data());
hipModuleGetFunction(&kernel, module, "vector_add");
```

And now this kernel can be launched via hipModule APIs.

The full example is below:

```
#include <hip/hip_runtime.h>
#include <hip/hiprtc.h>

#include <iostream>
#include <string>
#include <vector>

#define CHECK_RET_CODE(call, ret_code)
```

(continues on next page)

(continued from previous page)

```

    ↪      \
    ↪      {
    ↪          \
    ↪          if ((call) != ret_code) {
    ↪              \
    ↪              std::cout << "Failed in call: " << #call << std::endl;
    ↪              \
    ↪              std::abort();
    ↪              \
    ↪          }
    ↪      \
    ↪  }

#define HIP_CHECK(call) CHECK_RET_CODE(call, hipSuccess)
#define HIPRTC_CHECK(call) CHECK_RET_CODE(call, HIPRTC_SUCCESS)

// source code for hiprtc
static constexpr auto kernel_source{
    R"(
    extern "C"
    __global__ void vector_add(float* output, float* input1, float* input2, size_t size)
    ↪{
        int i = threadIdx.x;
        if (i < size) {
            output[i] = input1[i] + input2[i];
        }
    }
    )"};

int main() {
    hiprtcProgram prog;
    auto rtc_ret_code = hiprtcCreateProgram(&prog,           // HIPRTC program handle
                                           kernel_source,   // kernel source string
                                           "vector_add.cpp", // Name of the file
                                           0,               // Number of headers
                                           NULL,            // Header sources
                                           NULL);           // Name of header file

    if (rtc_ret_code != HIPRTC_SUCCESS) {
        std::cout << "Failed to create program" << std::endl;
        std::abort();
    }

    hipDeviceProp_t props;
    int device = 0;
    HIP_CHECK(hipGetDeviceProperties(&props, device));
    std::string sarg = std::string("--gpu-architecture=") +
        props.gcnArchName; // device for which binary is to be generated

    const char* options[] = {sarg.c_str()};

    rtc_ret_code = hiprtcCompileProgram(prog,           // hiprtcProgram
                                         0,             // Number of options

```

(continues on next page)

(continued from previous page)

```

                                options); // Clang Options
if (rtc_ret_code != HIPRTC_SUCCESS) {
    std::cout << "Failed to create program" << std::endl;
    std::abort();
}

size_t logSize;
HIPRTC_CHECK(hiprtcGetProgramLogSize(prog, &logSize));

if (logSize) {
    std::string log(logSize, '\\0');
    HIPRTC_CHECK(hiprtcGetProgramLog(prog, &log[0]));
    std::cout << "Compilation failed with: " << log << std::endl;
    std::abort();
}

size_t codeSize;
HIPRTC_CHECK(hiprtcGetCodeSize(prog, &codeSize));

std::vector<char> kernel_binary(codeSize);
HIPRTC_CHECK(hiprtcGetCode(prog, kernel_binary.data()));

HIPRTC_CHECK(hiprtcDestroyProgram(&prog));

hipModule_t module;
hipFunction_t kernel;

HIP_CHECK(hipModuleLoadData(&module, kernel_binary.data()));
HIP_CHECK(hipModuleGetFunction(&kernel, module, "vector_add"));

constexpr size_t ele_size = 256; // total number of items to add
std::vector<float> hinput, output;
hinput.reserve(ele_size);
output.reserve(ele_size);
for (size_t i = 0; i < ele_size; i++) {
    hinput.push_back(static_cast<float>(i + 1));
    output.push_back(0.0f);
}

float *dinput1, *dinput2, *doutput;
HIP_CHECK(hipMalloc(&dinput1, sizeof(float) * ele_size));
HIP_CHECK(hipMalloc(&dinput2, sizeof(float) * ele_size));
HIP_CHECK(hipMalloc(&doutput, sizeof(float) * ele_size));

HIP_CHECK(hipMemcpy(dinput1, hinput.data(), sizeof(float) * ele_size,
↳hipMemcpyHostToDevice));
HIP_CHECK(hipMemcpy(dinput2, hinput.data(), sizeof(float) * ele_size,
↳hipMemcpyHostToDevice));

struct {
    float* output;
    float* input1;

```

(continues on next page)

(continued from previous page)

```

    float* input2;
    size_t size;
} args{doutput, dinput1, dinput2, ele_size};

auto size = sizeof(args);
void* config[] = {HIP_LAUNCH_PARAM_BUFFER_POINTER, &args, HIP_LAUNCH_PARAM_BUFFER_SIZE,
↪ &size,
                HIP_LAUNCH_PARAM_END};

HIP_CHECK(hipModuleLaunchKernel(kernel, 1, 1, 1, ele_size, 1, 1, 0, nullptr, nullptr, ↪
↪ config));

HIP_CHECK(hipMemcpy(output.data(), doutput, sizeof(float) * ele_size, ↪
↪ hipMemcpyDeviceToHost));

for (size_t i = 0; i < ele_size; i++) {
    if ((hinput[i] + hinput[i]) != output[i]) {
        std::cout << "Failed in validation: " << (hinput[i] + hinput[i]) << " - " << ↪
↪ output[i]
        << std::endl;
        std::abort();
    }
}
std::cout << "Passed" << std::endl;

HIP_CHECK(hipFree(dinput1));
HIP_CHECK(hipFree(dinput2));
HIP_CHECK(hipFree(doutput));
}

```

11.2 HIPRTC specific options

HIPRTC provides a few HIPRTC specific flags

- `--gpu-architecture` : This flag can guide the code object generation for a specific gpu arch. Example: `--gpu-architecture=gfx906:sramecc+:xnack-`, its equivalent to `--offload-arch`.
 - This option is compulsory if compilation is done on a system without AMD GPUs supported by HIP runtime.
 - Otherwise, HIPRTC will load the hip runtime and gather the current device and its architecture info and use it as option.
- `-fgpu-rdc` : This flag when provided during the `hiprtcCompileProgram` generates the bitcode (HIPRTC doesn't convert this bitcode into ISA and binary). This bitcode can later be fetched using `hiprtcGetBitcode` and `hiprtcGetBitcodeSize` APIs.

11.2.1 Bitcode

In the usual scenario, the kernel associated with `hiprtcProgram` is compiled into the binary which can be loaded and run. However, if `-fpu-rdc` option is provided in the compile options, HIPRTC calls `comgr` and generates only the LLVM bitcode. It doesn't convert this bitcode to ISA and generate the final binary.

```
std::string sarg = std::string("-fgpu-rdc");
const char* options[] = {
    sarg.c_str() };
hiprtcCompileProgram(prog, // hiprtcProgram
                    1,     // Number of options
                    options);
```

If the compilation is successful, one can load the bitcode in a local variable using the bitcode APIs provided by HIPRTC.

```
size_t bitCodeSize;
hiprtcGetBitCodeSize(prog, &bitCodeSize);

vector<char> kernel_bitcode(bitCodeSize);
hiprtcGetBitCode(prog, kernel_bitcode.data());
```

11.2.2 CU Mode vs WGP mode

AMD GPUs consist of an array of workgroup processors, each built with 2 compute units (CUs) capable of executing SIMD32. All the CUs inside a workgroup processor use local data share (LDS).

gfx10+ support execution of wavefront in CU mode and work-group processor mode (WGP). Please refer to section 2.3 of [RDNA3 ISA reference](#).

gfx9 and below only supports CU mode.

In WGP mode, 4 warps of a block can simultaneously be executed on the workgroup processor, where as in CU mode only 2 warps of a block can simultaneously execute on a CU. In theory, WGP mode might help with occupancy and increase the performance of certain HIP programs (if not bound to inter warp communication), but might incur performance penalty on other HIP programs which rely on atomics and inter warp communication. This also has effect of how the LDS is split between warps, please refer to [RDNA3 ISA reference](#) for more information.

HIPRTC assumes **WGP mode by default** for gfx10+. This can be overridden by passing `-mcumode` to HIPRTC compile options in `hiprtcCompileProgram`.

11.3 Linker APIs

The bitcode generated using the HIPRTC Bitcode APIs can be loaded using `hipModule` APIs and also can be linked with other generated bitcodes with appropriate linker flags using the HIPRTC linker APIs. This also provides more flexibility and optimizations to the applications who want to generate the binary dynamically according to their needs. The input bitcodes can be generated only for a specific architecture or it can be a bundled bitcode which is generated for multiple architectures.

11.3.1 Example

Firstly, HIPRTC link instance or a pending linker invocation must be created using `hiprtcLinkCreate`, with the appropriate linker options provided.

```
hiprtcLinkCreate( num_options,      // number of options
                  options,         // Array of options
                  option_vals,     // Array of option values cast to void*
                  &rtc_link_state ); // HIPRTC link state created upon success
```

Following which, the bitcode data can be added to this link instance via `hiprtcLinkAddData` (if the data is present as a string) or `hiprtcLinkAddFile` (if the data is present as a file) with the appropriate input type according to the data or the bitcode used.

```
hiprtcLinkAddData(rtc_link_state, // HIPRTC link state
                  input_type,     // type of the input data or bitcode
                  bit_code_ptr,   // input data which is null terminated
                  bit_code_size,  // size of the input data
                  "a",            // optional name for this input
                  0,              // size of the options
                  0,              // Array of options applied to this input
                  0);             // Array of option values cast to void*
```

```
hiprtcLinkAddFile(rtc_link_state, // HIPRTC link state
                  input_type,     // type of the input data or bitcode
                  bc_file_path.c_str(), // path to the input file where bitcode is
↪present
                  0,              // size of the options
                  0,              // Array of options applied to this input
                  0);             // Array of option values cast to void*
```

Once the bitcodes for multiple architectures are added to the link instance, the linking of the device code must be completed using `hiprtcLinkComplete` which generates the final binary.

```
hiprtcLinkComplete(rtc_link_state, // HIPRTC link state
                   &binary,        // upon success, points to the output binary
                   &binarySize);   // size of the binary is stored (optional)
```

If the `hiprtcLinkComplete` returns successfully, the generated binary can be loaded and run using the `hipModule*` APIs.

```
hipModuleLoadData(&module, binary);
```

11.3.1.1 Note

- The compiled binary must be loaded before HIPRTC link instance is destroyed using the `hiprtcLinkDestroy` API.

```
hiprtcLinkDestroy(rtc_link_state);
```

- The correct sequence of calls is : `hiprtcLinkCreate`, `hiprtcLinkAddData` or `hiprtcLinkAddFile`, `hiprtcLinkComplete`, `hiprtcModuleLoadData`, `hiprtcLinkDestroy`.

11.3.2 Input Types

HIPRTC provides `hiprtcJITInputType` enumeration type which defines the input types accepted by the Linker APIs. Here are the enum values of `hiprtcJITInputType`. However only the input types `HIPRTC_JIT_INPUT_LLVM_BITCODE`, `HIPRTC_JIT_INPUT_LLVM_BUNDLED_BITCODE` and `HIPRTC_JIT_INPUT_LLVM_ARCHIVES_OF_BUNDLED_BITCODE` are supported currently.

`HIPRTC_JIT_INPUT_LLVM_BITCODE` can be used to load both LLVM bitcode or LLVM IR assembly code. However, `HIPRTC_JIT_INPUT_LLVM_BUNDLED_BITCODE` and `HIPRTC_JIT_INPUT_LLVM_ARCHIVES_OF_BUNDLED_BITCODE` are only for bundled bitcode and archive of bundled bitcode.

```
HIPRTC_JIT_INPUT_CUBIN = 0,
HIPRTC_JIT_INPUT_PTX,
HIPRTC_JIT_INPUT_FATBINARY,
HIPRTC_JIT_INPUT_OBJECT,
HIPRTC_JIT_INPUT_LIBRARY,
HIPRTC_JIT_INPUT_NVVM,
HIPRTC_JIT_NUM_LEGACY_INPUT_TYPES,
HIPRTC_JIT_INPUT_LLVM_BITCODE = 100,
HIPRTC_JIT_INPUT_LLVM_BUNDLED_BITCODE = 101,
HIPRTC_JIT_INPUT_LLVM_ARCHIVES_OF_BUNDLED_BITCODE = 102,
HIPRTC_JIT_NUM_INPUT_TYPES = (HIPRTC_JIT_NUM_LEGACY_INPUT_TYPES + 3)
```

11.3.3 Backward Compatibility of LLVM Bitcode/IR

For HIP applications utilizing HIPRTC to compile LLVM bitcode/IR, compatibility is assured only when the ROCm or HIP SDK version used for generating the LLVM bitcode/IR matches the version used during the runtime compilation. When an application requires the ingestion of bitcode/IR not derived from the currently installed AMD compiler, it must run with HIPRTC and `comgr` dynamic libraries that are compatible with the version of the bitcode/IR.

`comgr`, a shared library, incorporates the LLVM/Clang compiler that HIPRTC relies on. To identify the bitcode/IR version that `comgr` is compatible with, one can execute “`clang -v`” using the clang binary from the same ROCm or HIP SDK package. For instance, if compiling bitcode/IR version 14, the HIPRTC and `comgr` libraries released by AMD around mid 2022 would be the best choice, assuming the LLVM/Clang version included in the package is also version 14.

To ensure smooth operation and compatibility, an application may choose to ship the specific versions of HIPRTC and `comgr` dynamic libraries, or it may opt to clearly specify the version requirements and dependencies. This approach guarantees that the application can correctly compile the specified version of bitcode/IR.

11.3.4 Link Options

- `HIPRTC_JIT_IR_TO_ISA_OPT_EXT` - AMD Only. Options to be passed on to link step of compiler by `hiprtcLinkCreate`.
- `HIPRTC_JIT_IR_TO_ISA_OPT_COUNT_EXT` - AMD Only. Count of options passed on to link step of compiler.

Example:

```
const char* isaopts[] = {"-mllvm", "-inline-threshold=1", "-mllvm", "-inlinehint-
↪threshold=1"};
std::vector<hiprtcJIT_option> jit_options = {HIPRTC_JIT_IR_TO_ISA_OPT_EXT,
HIPRTC_JIT_IR_TO_ISA_OPT_COUNT_EXT};
size_t isaoptssize = 4;
```

(continues on next page)

(continued from previous page)

```
const void* lopts[] = {(void*)isaopts, (void*)(isaoptssize)};
hiprtcLinkState linkstate;
hiprtcLinkCreate(2, jit_options.data(), (void**)lopts, &linkstate);
```

11.4 Error Handling

HIPRTC defines the `hiprtcResult` enumeration type and a function `hiprtcGetErrorString` for API call error handling. `hiprtcResult` enum defines the API result codes. HIPRTC APIs return `hiprtcResult` to indicate the call result. `hiprtcGetErrorString` function returns a string describing the given `hiprtcResult` code, e.g., `HIPRTC_SUCCESS` to “`HIPRTC_SUCCESS`”. For unrecognized enumeration values, it returns “Invalid HIPRTC error code”.

`hiprtcResult` enum supported values and the `hiprtcGetErrorString` usage are mentioned below.

```
HIPRTC_SUCCESS = 0,
HIPRTC_ERROR_OUT_OF_MEMORY = 1,
HIPRTC_ERROR_PROGRAM_CREATION_FAILURE = 2,
HIPRTC_ERROR_INVALID_INPUT = 3,
HIPRTC_ERROR_INVALID_PROGRAM = 4,
HIPRTC_ERROR_INVALID_OPTION = 5,
HIPRTC_ERROR_COMPILATION = 6,
HIPRTC_ERROR_LINKING = 7,
HIPRTC_ERROR_BUILTIN_OPERATION_FAILURE = 8,
HIPRTC_ERROR_NO_NAME_EXPRESSIONS_AFTER_COMPILATION = 9,
HIPRTC_ERROR_NO_LOWERED_NAMES_BEFORE_COMPILATION = 10,
HIPRTC_ERROR_NAME_EXPRESSION_NOT_VALID = 11,
HIPRTC_ERROR_INTERNAL_ERROR = 12
```

```
hiprtcResult result;
result = hiprtcCompileProgram(prog, 1, opts);
if (result != HIPRTC_SUCCESS) {
    std::cout << "hiprtcCompileProgram fails with error " << hiprtcGetErrorString(result);
}
```

11.5 HIPRTC General APIs

HIPRTC provides the following API for querying the version.

`hiprtcVersion(int* major, int* minor)` - This sets the output parameters `major` and `minor` with the HIP Runtime compilation major version and minor version number respectively.

Currently, it returns hardcoded value. This should be implemented to return HIP runtime major and minor version in the future releases.

11.6 Lowered Names (Mangled Names)

HIPRTC mangles the `__global__` function names and names of `__device__` and `__constant__` variables. If the generated binary is being loaded using the HIP Runtime API, the kernel function or `__device__`/`__constant__` variable must be looked up by name, but this is very hard when the name has been mangled. To overcome this, HIPRTC provides API functions that map `__global__` function or `__device__`/`__constant__` variable names in the source to the mangled names present in the generated binary.

The two APIs `hiprtcAddNameExpression` and `hiprtcGetLoweredName` provide this functionality. First, a ‘name expression’ string denoting the address for the `__global__` function or `__device__`/`__constant__` variable is provided to `hiprtcAddNameExpression`. Then, the program is compiled with `hiprtcCompileProgram`. During compilation, HIPRTC will parse the name expression string as a C++ constant expression at the end of the user program. Finally, the function `hiprtcGetLoweredName` is called with the original name expression and it returns a pointer to the lowered name. The lowered name can be used to refer to the kernel or variable in the HIP Runtime API.

11.6.1 Note

- The identical name expression string must be provided on a subsequent call to `hiprtcGetLoweredName` to extract the lowered name.
- The correct sequence of calls is : `hiprtcAddNameExpression`, `hiprtcCompileProgram`, `hiprtcGetLoweredName`, `hiprtcDestroyProgram`.
- The lowered names must be fetched using `hiprtcGetLoweredName` only after the HIPRTC program has been compiled, and before it has been destroyed.

11.6.2 Example

kernel containing various definitions `__global__` functions/function templates and `__device__`/`__constant__` variables can be stored in a string.

```
static constexpr const char gpu_program[] {
R"(
__device__ int V1; // set from host code
static __global__ void f1(int *result) { *result = V1 + 10; }
namespace N1 {
namespace N2 {
__constant__ int V2; // set from host code
__global__ void f2(int *result) { *result = V2 + 20; }
}
}
template<typename T>
__global__ void f3(int *result) { *result = sizeof(T); }
)"};
```

`hiprtcAddNameExpression` is called with various name expressions referring to the address of `__global__` functions and `__device__`/`__constant__` variables.

```
kernel_name_vec.push_back("&f1");
kernel_name_vec.push_back("N1:N2::f2");
kernel_name_vec.push_back("f3<int>");
for (auto& x : kernel_name_vec) hiprtcAddNameExpression(prog, x.c_str());
variable_name_vec.push_back("&V1");
```

(continues on next page)

(continued from previous page)

```
variable_name_vec.push_back("&N1::N2::V2");
for (auto&& x : variable_name_vec) hiprtcAddNameExpression(prog, x.c_str());
```

After which, the program is compiled using `hiprtcCompileProgram` and the generated binary is loaded using `hipModuleLoadData`. And the mangled names can be fetched using `hiprtcGetLoweredName`.

```
for (decltype(variable_name_vec.size()) i = 0; i != variable_name_vec.size(); ++i) {
    const char* name;
    hiprtcGetLoweredName(prog, variable_name_vec[i].c_str(), &name);
}
```

```
for (decltype(kernel_name_vec.size()) i = 0; i != kernel_name_vec.size(); ++i) {
    const char* name;
    hiprtcGetLoweredName(prog, kernel_name_vec[i].c_str(), &name);
}
```

The mangled name of the variables are used to look up the variable in the module and update its value.

```
hipDeviceptr_t variable_addr;
size_t bytes{};
hipModuleGetGlobal(&variable_addr, &bytes, module, name);
hipMemcpyHtoD(variable_addr, &initial_value, sizeof(initial_value));
```

Finally, the mangled name of the kernel is used to launch it using the `hipModule` APIs.

```
hipFunction_t kernel;
hipModuleGetFunction(&kernel, module, name);
hipModuleLaunchKernel(kernel, 1, 1, 1, 1, 1, 1, 0, nullptr, nullptr, config);
```

Please have a look at `hiprtcGetLoweredName.cpp` for the detailed example.

11.7 Versioning

HIPRTC follows the below versioning.

- Linux
 - HIPRTC follows the same versioning as HIP runtime library.
 - The so name field for the shared library is set to MAJOR version. For example, for HIP 5.3 the so name is set to 5 (`hiprtc.so.5`).
- Windows
 - HIPRTC dll is named as `hiprtcXXYY.dll` where XX is MAJOR version and YY is MINOR version. For example, for HIP 5.3 the name is `hiprtc0503.dll`.

11.8 HIP header support

- Added HIPRTC support for all the hip common header files such as `library_types.h`, `hip_math_constants.h`, `hip_complex.h`, `math_functions.h`, `surface_types.h` etc. from 6.1. HIPRTC users need not include any HIP macros or constants explicitly in their header files. All of these should get included via HIPRTC builtins when the app links to HIPRTC library.

11.9 Deprecation notice

- Currently HIPRTC APIs are separated from HIP APIs and HIPRTC is available as a separate library `libhiprtc.so/libhiprtc.dll`. But on Linux, HIPRTC symbols are also present in `libamdhip64.so` in order to support the existing applications. Gradually, these symbols will be removed from HIP library and applications using HIPRTC will be required to explicitly link to HIPRTC library. However, on Windows `hiprtc.dll` must be used as the `amdhip64.dll` doesn't contain the HIPRTC symbols.
- Data types such as `uint32_t`, `uint64_t`, `int32_t`, `int64_t` defined in `std` namespace in HIPRTC are deprecated earlier and are being removed from ROCm release 6.1 since these can conflict with the standard C++ data types. These data types are now prefixed with `__hip__`, e.g. `__hip_uint32_t`. Applications previously using `std::uint32_t` or similar types can use `__hip_` prefixed types to avoid conflicts with standard `std` namespace or application can have their own definitions for these types. Also, `type_traits` templates previously defined in `std` namespace are moved to `__hip_internal` namespace as implementation details.

PERFORMANCE GUIDELINES

The AMD HIP Performance Guidelines are a set of best practices designed to help developers optimize the performance of AMD GPUs. They cover established parallelization and optimization techniques, coding metaphors, and idioms that can greatly simplify programming for HIP-capable GPU architectures.

By following four main cornerstones, we can exploit the performance optimization potential of HIP.

- parallel execution
- memory usage optimization
- optimization for maximum throughput
- minimizing memory thrashing

In the following chapters, we will show you their benefits and how to use them effectively.

12.1 Parallel execution

For optimal use, the application should reveal and efficiently imply as much parallelism as possible to keep all system components active.

12.1.1 Application level

The application should optimize parallel execution across the host and devices using asynchronous calls and streams. Workloads should be assigned based on efficiency: serial to the host, parallel to the devices.

For parallel workloads, when threads need to synchronize to share data, if they belong to the same block, they should use `__syncthreads()` (see: *Synchronization functions*) within the same kernel invocation. If they belong to different blocks, they must use global memory with two separate kernel invocations. The latter should be minimized as it adds overhead.

12.1.2 Device level

Device-level optimization primarily involves maximizing parallel execution across the multiprocessors of the device. This can be achieved by executing multiple kernels concurrently on a device. The management of these kernels is facilitated by streams, which allow for the overlapping of computation and data transfers, enhancing performance. The aim is to keep all multiprocessors busy by executing enough kernels concurrently. However, launching too many kernels can lead to resource contention, so a balance must be found for optimal performance. This approach helps in achieving maximum utilization of the resources of the device.

12.1.3 Multiprocessor level

Multiprocessor-level optimization involves maximizing parallel execution within each multiprocessor on a device. Each multiprocessor can execute a number of threads concurrently, and the total number of threads that can run in parallel is determined by the number of concurrent threads each multiprocessor can handle.

The key to multiprocessor-level optimization is to efficiently utilize the various functional units within a multiprocessor. This can be achieved by ensuring a sufficient number of resident warps, as at every instruction issue time, a warp scheduler selects an instruction that is ready to execute. This instruction can be another independent instruction of the same warp, exploiting *Optimization for maximum instruction throughput*, or more commonly an instruction of another warp, exploiting thread-level parallelism.

In comparison, device-level optimization focuses on the device as a whole, aiming to keep all multiprocessors busy by executing enough kernels concurrently. Both levels of optimization are crucial for achieving maximum performance. They work together to ensure efficient utilization of the resources of the GPU, from the individual multiprocessors to the device as a whole.

12.2 Memory optimization

The first step in maximizing memory throughput is to minimize low-bandwidth data transfers. This involves reducing data transfers between the host and the device, as these have lower bandwidth than transfers between global memory and the device.

Additionally, data transfers between global memory and the device should be minimized by maximizing the use of on-chip memory: shared memory and caches. Shared memory acts as a user-managed cache, where the application explicitly allocates and accesses it. A common programming pattern is to stage data from device memory into shared memory. This involves each thread of a block loading data from device memory to shared memory, synchronizing with all other threads of the block, processing the data in shared memory, synchronizing again if necessary, and writing the results back to device global memory.

For some applications, a traditional hardware-managed cache is more appropriate to exploit data locality. On devices of certain compute capabilities, the same on-chip memory is used for both L1 and shared memory, and the amount dedicated to each is configurable for each kernel call.

Finally, the throughput of memory accesses by a kernel can vary significantly depending on the access pattern for each type of memory. Therefore, the next step in maximizing memory throughput is to organize memory accesses as optimally as possible. This is especially important for global memory accesses, as global memory bandwidth is low compared to available on-chip bandwidths and arithmetic instruction throughput. Thus, non-optimal global memory accesses generally have a high impact on performance.

12.2.1 Data Transfer

Applications should aim to minimize data transfers between the host and the device. This can be achieved by moving more computations from the host to the device, even if it means running kernels that do not fully utilize the parallelism for device. Intermediate data structures can be created, used, and discarded in device memory without being mapped or copied to host memory.

Batching small transfers into a single large transfer can improve performance due to the overhead associated with each transfer. On systems with a front-side bus, using page-locked host memory can enhance data transfer performance.

When using mapped page-locked memory, there is no need to allocate device memory or explicitly copy data between device and host memory. Data transfers occur implicitly each time the kernel accesses the mapped memory. For optimal performance, these memory accesses should be coalesced, similar to global memory accesses.

On integrated systems where device and host memory are physically the same, any copy operation between host and device memory is unnecessary, and mapped page-locked memory should be used instead. Applications can check if a device is integrated by querying the integrated device property.

12.2.2 Device Memory Access

Memory access instructions may be repeated due to the spread of memory addresses across warp threads. The impact on throughput varies with memory type and is generally reduced when addresses are more scattered, especially in global memory.

Device memory is accessed via 32-, 64-, or 128-byte transactions that must be naturally aligned. Maximizing memory throughput involves coalescing memory accesses of threads within a warp into minimal transactions, following optimal access patterns, using properly sized and aligned data types, and padding data when necessary.

Global memory instructions support reading or writing data of specific sizes (1, 2, 4, 8, or 16 bytes) that are naturally aligned. If the size and alignment requirements are not met, it leads to multiple instructions, reducing performance. Therefore, using data types that meet these requirements, ensuring alignment for structures, and maintaining alignment for all values or arrays is crucial for correct results and optimal performance.

Threads often access 2D arrays at an address calculated as `BaseAddress + xIndex + width * yIndex`. For efficient memory access, the array and thread block widths should be multiples of the warp size. If the array width is not a multiple of the warp size, it is usually more efficient to allocate it with a width rounded up to the nearest multiple and pad the rows accordingly.

Local memory is used for certain automatic variables, such as arrays with non-constant indices, large structures or arrays, and any variable when the kernel uses more registers than available. Local memory resides in device memory, leading to high latency and low bandwidth similar to global memory accesses. However, it is organized for consecutive 32-bit words to be accessed by consecutive thread IDs, allowing full coalescing when all threads in a warp access the same relative address.

Shared memory, located on-chip, provides higher bandwidth and lower latency than local or global memory. It is divided into banks that can be simultaneously accessed, boosting bandwidth. However, bank conflicts, where two addresses fall in the same bank, lead to serialized access and decreased throughput. Therefore, understanding how memory addresses map to banks and scheduling requests to minimize conflicts is crucial for optimal performance.

Constant memory is in device memory and cached in the constant cache. Requests are split based on different memory addresses, affecting throughput, and are serviced at the throughput of the constant cache for cache hits, or the throughput of the device memory otherwise.

Texture and surface memory are stored in device memory and cached in texture cache. This setup optimizes 2D spatial locality, leading to better performance for threads reading close 2D addresses. Reading device memory through texture or surface fetching can be advantageous, offering higher bandwidth for local texture fetches or surface reads, offloading addressing calculations, allowing data broadcasting, and optional conversion of 8-bit and 16-bit integer input data to 32-bit floating-point values on-the-fly.

12.3 Optimization for maximum instruction throughput

To maximize instruction throughput:

- minimize low throughput arithmetic instructions
- minimize divergent warps inflicted by control flow instructions
- minimize the number of instruction as possible
- maximize instruction parallelism

12.3.1 Arithmetic instructions

The type and complexity of arithmetic operations can significantly impact the performance of your application. We are highlighting some hints how to maximize it.

Using efficient operations: Some arithmetic operations are more costly than others. For example, multiplication is typically faster than division, and integer operations are usually faster than floating-point operations, especially with double-precision.

Minimizing low-throughput instructions: This might involve trading precision for speed when it does not affect the final result. For instance, consider using single-precision arithmetic instead of double-precision.

Leverage intrinsic functions: Intrinsic functions are pre-defined functions available in HIP that can often be executed faster than equivalent arithmetic operations (subject to some input or accuracy restrictions). They can help optimize performance by replacing more complex arithmetic operations.

Avoiding divergent warps: Divergent warps occur when threads within the same warp follow different execution paths. This can happen due to conditional statements that lead to different arithmetic operations being performed by different threads. Divergent warps can significantly reduce instruction throughput, so try to structure your code to minimize divergence.

Optimizing memory access: The efficiency of memory access can impact the speed of arithmetic operations. Coalesced memory access, where threads in a warp access consecutive memory locations, can improve memory throughput and thus the speed of arithmetic operations.

Maximizing instruction parallelism: Some GPU architectures could issue parallel independent instructions simultaneously, for example integer and floating point, or two operations with independent inputs and outputs. Mostly this is a work for compiler, but expressing parallelism in the code explicitly can improve instructions throughput.

12.3.2 Control flow instructions

Flow control instructions (`if`, `else`, `for`, `do`, `while`, `break`, `continue`, `switch`) can impact instruction throughput by causing threads within a warp to diverge and follow different execution paths. To optimize performance, control conditions should be written to minimize divergent warps. For example, when the control condition depends on (`threadIdx / warpSize`), no warp diverges. The compiler may optimize loops or short `if` or `switch` blocks using branch predication, preventing warp divergence. With branch predication, instructions associated with a false predicate are scheduled but not executed, avoiding unnecessary operations.

12.3.3 Synchronization

Synchronization ensures that all threads within a block have completed their computations and memory accesses before moving forward, which is critical when threads are dependent on the results of other threads. However, synchronization can also lead to performance overhead, as it requires threads to wait, potentially leading to idle GPU resources.

`__syncthreads()` is used to synchronize all threads in a block, ensuring that all threads have reached the same point in the code and that shared memory is visible to all threads after the point of synchronization.

An alternative way to synchronize is using streams. Different streams can execute commands out of order with respect to one another or concurrently. This allows for more fine-grained control over the execution order of commands, which can be beneficial in certain scenarios.

12.4 Minimizing memory thrashing

Applications frequently allocating and freeing memory may experience slower allocation calls over time. This is expected as memory is released back to the operating system. To optimize performance in such scenarios, consider some recommendations:

- avoid allocating all available memory with `hipMalloc` / `hipHostMalloc`, as this immediately reserves memory and can block other applications from using it. This could strain the operating system schedulers or even prevent other applications from running on the same GPU.
- aim to allocate memory in suitably sized blocks early in the lifecycle of the application and deallocate only when the application no longer needs it. Minimize the number of `hipMalloc` and `hipFree` calls in your application, particularly in areas critical to performance.
- if an application is unable to allocate sufficient device memory, consider resorting to other memory types such as `hipHostMalloc` or `hipMallocManaged`. While these may not offer the same performance, they can allow the application to continue running.
- For supported platforms, `hipMallocManaged` allows for oversubscription. With the right memory advise policies, it can maintain most, if not all, of the performance of `hipMalloc`. `hipMallocManaged` does not require an allocation to be resident until it is needed or prefetched, easing the load on the operating system schedulers and facilitating multi-tenant scenarios.

DEBUGGING WITH HIP

AMD debugging tools include *ltrace* and *ROCgdb*. External tools are available and can be found online. For example, if you're using Windows, you can use *Microsoft Visual Studio* and *WinGDB*.

You can trace and debug your code using the following tools and techniques.

13.1 Tracing

You can use tracing to quickly observe the flow of an application before reviewing the detailed information provided by a command-line debugger. Tracing can be used to identify issues ranging from accidental API calls to calls made on a critical path.

ltrace is a standard Linux tool that provides a message to `stderr` on every dynamic library call. You can use *ltrace* to visualize the runtime behavior of the entire ROCm software stack.

Here's a simple command-line example that uses *ltrace* to trace HIP APIs and output:

```
$ ltrace -C -e "hip*" ./hipGetChanDesc
hipGetChanDesc->hipCreateChannelDesc(0x7ffdc4b66860, 32, 0, 0) = 0x7ffdc4b66860
hipGetChanDesc->hipMallocArray(0x7ffdc4b66840, 0x7ffdc4b66860, 8, 8) = 0
hipGetChanDesc->hipGetChannelDesc(0x7ffdc4b66848, 0xa63990, 5, 1) = 0
hipGetChanDesc->hipFreeArray(0xa63990, 0, 0x7f8c7fe13778, 0x7ffdc4b66848) = 0
PASSED!
+++ exited (status 0) +++
```

Here's another example that uses *ltrace* to trace hsa APIs and output:

```
$ ltrace -C -e "hsa*" ./hipGetChanDesc
libamdhip64.so.4->hsa_init(0, 0x7fff325a69d0, 0x9c80e0, 0 <unfinished ...>
libhsa-runtime64.so.1->hsaKmtOpenKFD(0x7fff325a6590, 0x9c38c0, 0, 1) = 0
libhsa-runtime64.so.1->hsaKmtGetVersion(0x7fff325a6608, 0, 0, 0) = 0
libhsa-runtime64.so.1->hsaKmtReleaseSystemProperties(3, 0x80084b01, 0, 0) = 0
libhsa-runtime64.so.1->hsaKmtAcquireSystemProperties(0x7fff325a6610, 0, 0, 1) = 0
libhsa-runtime64.so.1->hsaKmtGetNodeProperties(0, 0x7fff325a66a0, 0, 0) = 0
libhsa-runtime64.so.1->hsaKmtGetNodeMemoryProperties(0, 1, 0x9c42b0, 0x936012) = 0
...
<... hsaKmtCreateEvent resumed> ) = 0
libhsa-runtime64.so.1->hsaKmtAllocMemory(0, 4096, 64, 0x7fff325a6690) = 0
libhsa-runtime64.so.1->hsaKmtMapMemoryToGPUNodes(0x7f1202749000, 4096, 0x7fff325a6690, 0, 0) = 0
libhsa-runtime64.so.1->hsaKmtCreateEvent(0x7fff325a6700, 0, 0, 0x7fff325a66f0) = 0
```

(continues on next page)

(continued from previous page)

```

libhsa-runtime64.so.1->hsaKmtAllocMemory(1, 0x100000000, 576, 0x7fff325a67d8) = 0
libhsa-runtime64.so.1->hsaKmtAllocMemory(0, 8192, 64, 0x7fff325a6790) = 0
libhsa-runtime64.so.1->hsaKmtMapMemoryToGPUNodes(0x7f120273c000, 8192, 0x7fff325a6790, ↵
↵0) = 0
libhsa-runtime64.so.1->hsaKmtAllocMemory(0, 4096, 4160, 0x7fff325a6450) = 0
libhsa-runtime64.so.1->hsaKmtMapMemoryToGPUNodes(0x7f120273a000, 4096, 0x7fff325a6450, ↵
↵0) = 0
libhsa-runtime64.so.1->hsaKmtSetTrapHandler(1, 0x7f120273a000, 4096, 0x7f120273c000) = 0
<... hsa_init resumed> ) = 0
libamdhip64.so.4->hsa_system_get_major_extension_table(513, 1, 24, 0x7f1202597930) = 0
libamdhip64.so.4->hsa_iterate_agents(0x7f120171f050, 0, 0x7fff325a67f8, 0 <unfinished ...
↵>
libamdhip64.so.4->hsa_agent_get_info(0x94f110, 17, 0x7fff325a67e8, 0) = 0
libamdhip64.so.4->hsa_amd_agent_iterate_memory_pools(0x94f110, 0x7f1201722816, ↵
↵0x7fff325a67f0, 0x7f1201722816 <unfinished ...>
libamdhip64.so.4->hsa_amd_memory_pool_get_info(0x9c7fb0, 0, 0x7fff325a6744, ↵
↵0x7fff325a67f0) = 0
libamdhip64.so.4->hsa_amd_memory_pool_get_info(0x9c7fb0, 1, 0x7fff325a6748, ↵
↵0x7f1200d82df4) = 0
...
<... hsa_amd_agent_iterate_memory_pools resumed> ) = 0
libamdhip64.so.4->hsa_agent_get_info(0x9dbf30, 17, 0x7fff325a67e8, 0) = 0
<... hsa_iterate_agents resumed> ) = 0
libamdhip64.so.4->hsa_agent_get_info(0x9dbf30, 0, 0x7fff325a6850, 3) = 0
libamdhip64.so.4->hsa_agent_get_info(0x9dbf30, 0xa000, 0x9e7cd8, 0) = 0
libamdhip64.so.4->hsa_agent_iterate_isas(0x9dbf30, 0x7f1201720411, 0x7fff325a6760, ↵
↵0x7f1201720411) = 0
libamdhip64.so.4->hsa_isa_get_info_alt(0x94e7c8, 0, 0x7fff325a6728, 1) = 0
libamdhip64.so.4->hsa_isa_get_info_alt(0x94e7c8, 1, 0x9e7f90, 0) = 0
libamdhip64.so.4->hsa_agent_get_info(0x9dbf30, 4, 0x9e7ce8, 0) = 0
...
<... hsa_amd_memory_pool_allocate resumed> ) = 0
libamdhip64.so.4->hsa_ext_image_create(0x9dbf30, 0xa1c4c8, 0x7f10f2800000, 3 <unfinished ↵
↵...>
libhsa-runtime64.so.1->hsaKmtAllocMemory(0, 4096, 64, 0x7fff325a6740) = 0
libhsa-runtime64.so.1->hsaKmtQueryPointerInfo(0x7f1202736000, 0x7fff325a65e0, 0, 0) = 0
libhsa-runtime64.so.1->hsaKmtMapMemoryToGPUNodes(0x7f1202736000, 4096, 0x7fff325a66e8, ↵
↵0) = 0
<... hsa_ext_image_create resumed> ) = 0
libamdhip64.so.4->hsa_ext_image_destroy(0x9dbf30, 0x7f1202736000, 0x9dbf30, 0
↵<unfinished ...>
libhsa-runtime64.so.1->hsaKmtUnmapMemoryToGPU(0x7f1202736000, 0x7f1202736000, 4096, ↵
↵0x9c8050) = 0
libhsa-runtime64.so.1->hsaKmtFreeMemory(0x7f1202736000, 4096, 0, 0) = 0
<... hsa_ext_image_destroy resumed> ) = 0
libamdhip64.so.4->hsa_amd_memory_pool_free(0x7f10f2800000, 0x7f10f2800000, 256, ↵
↵0x9e76f0) = 0
PASSED!

```


13.2 Debugging

You can use ROCgdb for debugging and profiling.

ROCgdb is the ROCm source-level debugger for Linux and is based on GNU Project debugger (GDB). the GNU source-level debugger, equivalent of CUDA-GDB, can be used with debugger frontends, such as Eclipse, Visual Studio Code, or GDB dashboard. For details, see (<https://github.com/ROCm/ROCgdb>).

Below is a sample how to use ROCgdb run and debug HIP application, ROCgdb is installed with ROCM package in the folder /opt/rocm/bin.

```
$ export PATH=$PATH:/opt/rocm/bin
$ rocgdb ./hipTexObjPitch
GNU gdb (rocm-dkms-no-npi-hipclang-6549) 10.1
Copyright (C) 2020 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
...
For bug reporting instructions, please see:
<https://github.com/ROCm/ROCgdb/issues>.
Find the GDB manual and other documentation resources online at:
    <http://www.gnu.org/software/gdb/documentation/>.
...
Reading symbols from ./hipTexObjPitch...
(gdb) break main
Breakpoint 1 at 0x4013d1: file /home/test/hip/tests/src/texture/hipTexObjPitch.cpp, line 98.
(gdb) run
Starting program: /home/test/hip/build/directed_tests/texture/hipTexObjPitch
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, main ()
    at /home/test/hip/tests/src/texture/hipTexObjPitch.cpp:98
98     texture2Dtest<float>();
(gdb) c
```

13.2.1 Debugging HIP applications

The following Linux example shows how to get useful information from the debugger while running a simple memory copy test, which caused a segmentation fault issue.

```
test: simpleTest2<?> numElements=4194304 sizeElements=4194304 bytes
Segmentation fault (core dumped)

(gdb) run
Starting program: /home/test/hipamd/build/directed_tests/runtimeApi/memory/hipMemcpy_simple
simple
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, main (argc=1, argv=0x7fffffffdea8)
    at /home/test/hip/tests/src/runtimeApi/memory/hipMemcpy_simple.cpp:147
```

(continues on next page)

(continued from previous page)

```

147     int main(int argc, char* argv[]) {
(gdb) c
Continuing.
[New Thread 0x7ffff64c4700 (LWP 146066)]

Thread 1 "hipMemcpy_simple" received signal SIGSEGV, Segmentation fault.
0x000000000020f78e in simpleTest2<float> (numElements=4194304, usePinnedHost=true)
    at /home/test/hip/tests/src/runtimeApi/memory/hipMemcpy_simple.cpp:104
104         A_h1[i] = 3.14f + 1000 * i;
(gdb) bt
#0 0x000000000020f78e in simpleTest2<float> (numElements=4194304, usePinnedHost=true)
    at /home/test/hip/tests/src/runtimeApi/memory/hipMemcpy_simple.cpp:104
#1 0x000000000020e96c in main (argc=<optimized out>, argv=<optimized out>)
    at /home/test/hip/tests/src/runtimeApi/memory/hipMemcpy_simple.cpp:163
(gdb) info thread
Id      Target Id                                     Frame
* 1      Thread 0x7ffff64c5880 (LWP 146060) "hipMemcpy_simple" 0x000000000020f78e in
↳ simpleTest2<float> (numElements=4194304, usePinnedHost=true)
    at /home/test/hip/tests/src/runtimeApi/memory/hipMemcpy_simple.cpp:104
2      Thread 0x7ffff64c4700 (LWP 146066) "hipMemcpy_simple" 0x00007ffff6b0850b in ioctl
    () from /lib/x86_64-linux-gnu/libc.so.6
(gdb) thread 2
[Switching to thread 2 (Thread 0x7ffff64c4700 (LWP 146066))]
#0 0x00007ffff6b0850b in ioctl () from /lib/x86_64-linux-gnu/libc.so.6
(gdb) bt
#0 0x00007ffff6b0850b in ioctl () from /lib/x86_64-linux-gnu/libc.so.6
#1 0x00007ffff6604568 in ?? () from /opt/rocm/lib/libhsa-runtime64.so.1
#2 0x00007ffff65fe73a in ?? () from /opt/rocm/lib/libhsa-runtime64.so.1
#3 0x00007ffff659e4d6 in ?? () from /opt/rocm/lib/libhsa-runtime64.so.1
#4 0x00007ffff65807de in ?? () from /opt/rocm/lib/libhsa-runtime64.so.1
#5 0x00007ffff65932a2 in ?? () from /opt/rocm/lib/libhsa-runtime64.so.1
#6 0x00007ffff654f547 in ?? () from /opt/rocm/lib/libhsa-runtime64.so.1
#7 0x00007ffff7f76609 in start_thread () from /lib/x86_64-linux-gnu/libpthread.so.0
#8 0x00007ffff6b13293 in clone () from /lib/x86_64-linux-gnu/libc.so.6
(gdb) thread 1
[Switching to thread 1 (Thread 0x7ffff64c5880 (LWP 146060))]
#0 0x000000000020f78e in simpleTest2<float> (numElements=4194304, usePinnedHost=true)
    at /home/test/hip/tests/src/runtimeApi/memory/hipMemcpy_simple.cpp:104
104         A_h1[i] = 3.14f + 1000 * i;
(gdb) bt
#0 0x000000000020f78e in simpleTest2<float> (numElements=4194304, usePinnedHost=true)
    at /home/test/hip/tests/src/runtimeApi/memory/hipMemcpy_simple.cpp:104
#1 0x000000000020e96c in main (argc=<optimized out>, argv=<optimized out>)
    at /home/test/hip/tests/src/runtimeApi/memory/hipMemcpy_simple.cpp:163
(gdb)
...

```

Debugging HIP applications using Windows tools can be more informative than on Linux. Windows tools provides more visibility into debug codes, which makes it easier to inspect variables, watch multiple details, and examine call stacks.

13.3 Useful environment variables

HIP provides environment variables that allow HIP, hip-clang, or HSA drivers to prevent certain features and optimizations. These are not intended for production, but can be useful to diagnose synchronization problems in the application (or driver).

Some of the more widely used environment variables are described in this section. These are supported on the Linux ROCm path and Windows.

13.3.1 Kernel enqueue serialization

You can control kernel command serialization from the host:

AMD_SERIALIZE_KERNEL, for serializing kernel enqueue

AMD_SERIALIZE_KERNEL = 1, Wait for completion before enqueue AMD_SERIALIZE_KERNEL = 2, Wait for completion after enqueue AMD_SERIALIZE_KERNEL = 3, Both

Or

AMD_SERIALIZE_COPY, for serializing copies

AMD_SERIALIZE_COPY = 1, Wait for completion before enqueue AMD_SERIALIZE_COPY = 2, Wait for completion after enqueue AMD_SERIALIZE_COPY = 3, Both

So HIP runtime can wait for GPU idle before/after any GPU command depending on the environment setting.

13.3.2 Making device visible

For systems with multiple devices, you can choose to make only certain device(s) visible to HIP using HIP_VISIBLE_DEVICES (or CUDA_VISIBLE_DEVICES on an NVIDIA platform). Once enabled, HIP can only view devices that have indices present in the sequence. For example:

```
$ HIP_VISIBLE_DEVICES=0,1
```

Or in the application:

```
if (totalDeviceNum > 2) {
    setenv("HIP_VISIBLE_DEVICES", "0,1,2", 1);
    assert(getDeviceNumber(false) == 3);
    ...
}
```

13.3.3 Dump code object

To analyze compiler-related issues, you can use the dump code object: GPU_DUMP_CODE_OBJECT.

13.3.4 HSA-related environment variables (Linux)

HSA provides environment variables that help analyze issues in drivers or hardware.

- To isolate issues with hardware copy engines, you can use `HSA_ENABLE_SDMA`.

`HSA_ENABLE_SDMA=0` causes host-to-device and device-to-host copies to use compute shader blit kernels, rather than the dedicated DMA copy engines. Compute shader copies have low latency (typically < 5 us) and can achieve approximately 80% of the bandwidth of the DMA copy engine.

- To diagnose interrupt storm issues in the driver, you can use `HSA_ENABLE_INTERRUPT`.

`HSA_ENABLE_INTERRUPT=0` causes completion signals to be detected with memory-based polling, rather than interrupts.

13.3.5 HIP environment variable summary

Here are some of the more commonly used environment variables:

Environment variable	De-fault value	Usage
AMD_LOG_LEV Enable HIP log on different Level	0	0: Disable log. 1: Enable log on error level 2: Enable log on warning and below levels 0x3: Enable log on information and below levels 0x4: Decode and display AQL packets
AMD_LOG_MA Enable HIP log on different Level	0x7F	0x1: Log API calls 0x02: Kernel and Copy Commands and Barriers 0x4: Synchronization and waiting for commands to finish 0x8: Enable log on information and below levels 0x20: Queue commands and queue contents 0x40: Signal creation, allocation, pool 0x80: Locks and thread-safety code 0x100: Copy debug 0x200: Detailed copy debug 0x400: Resource allocation, performance-impacting events 0x800: Initialization and shutdown 0x1000: Misc debug, not yet classified 0x2000: Show raw bytes of AQL packet 0x4000: Show code creation debug 0x8000: More detailed command info, including barrier commands 0x10000: Log message location 0xFFFFFFFF: Log always even mask flag is zero
HIP_LAUNCH_E Used for serialization on kernel execution.	0	0: Disable. Kernel executes normally. 1: Enable. Serializes kernel enqueue, behaves the same as AMD_SERIALIZE_KERNEL.
HIP_VISIBLE_D (or CUDA_VISIBLE_Only devices whose index is present in the sequence are visible to HIP		0,1,2: Depending on the number of devices on the system
GPU_DUMP_CO Dump code object	0	0: Disable 1: Enable
AMD_SERIALIZ Serialize kernel enqueue	0	1: Wait for completion before enqueue 2: Wait for completion after enqueue 3: Both
AMD_SERIALIZ Serialize copies	0	1: Wait for completion before enqueue 2: Wait for completion after enqueue 3: Both
HIP_HOST_COH Coherent memory in hipHost-Malloc	0	0: memory is not coherent between host and GPU 1: memory is coherent with host
AMD_DIRECT_I Enable direct kernel dispatch (Currently for Linux; under development for Windows)	1	0: Disable 1: Enable
GPU_MAX_HW The maximum number of hardware queues allocated per device	4	The variable controls how many independent hardware queues HIP runtime can create per process, per device. If an application allocates more HIP streams than this number, then HIP runtime reuses the same hardware queues for the new streams in a round-robin manner. Note that this maximum number does not apply to hardware queues that are created for CU-masked HIP streams, or cooperative queues for HIP Cooperative Groups (single queue per device).

13.4 General debugging tips

- `gdb --args` can be used to pass the executable and arguments to `gdb`.
- You can set environment variables (`set env`) from within GDB on Linux:

```
(gdb) set env AMD_SERIALIZE_KERNEL 3
```

Note: This `gdb` command does not use an equal (=) sign.

- The GDB backtrace shows a path in the runtime. This is because a fault is caught by the runtime, but it is generated by an asynchronous command running on the GPU.
- To determine the true location of a fault, you can force the kernels to run synchronously by setting the environment variables `AMD_SERIALIZE_KERNEL=3` and `AMD_SERIALIZE_COPY=3`. This forces HIP runtime to wait for the kernel to finish running before returning. If the fault occurs when a kernel is running, you can see the code that launched the kernel inside the backtrace. The thread that's causing the issue is typically the one inside `libhsa-runtime64.so`.
- VM faults inside kernels can be caused by:
 - Incorrect code (e.g., a for loop that extends past array boundaries)
 - Memory issues, such as invalid kernel arguments (null pointers, unregistered host pointers, bad pointers)
 - Synchronization issues
 - Compiler issues (incorrect code generation from the compiler)
 - Runtime issues

LOGGING HIP ACTIVITY

HIP provides a logging mechanism that allows you to trace HIP API and runtime codes when running a HIP application. In addition to being useful to our users/developers, the HIP development team uses these logs to improve the HIP runtime.

By adjusting the logging settings and logging mask, you can get different types of information for different functionalities, such as HIP APIs, executed kernels, queue commands, and queue contents. Refer to the following sections for examples.

Tip: Logging works for the release and debug versions of HIP. If you want to save logging output in a file, define the file when running the application via command line. For example:

```
user@user-test:~/hip/bin$ ./hipinfo > ~/hip_log.txt
```

14.1 Logging level

HIP logging is disabled by default. You can enable it via the `AMD_LOG_LEVEL` environment variable. The value of this variable controls your logging level. Levels are defined as follows:

```
enum LogLevel {  
    LOG_NONE    = 0,  
    LOG_ERROR   = 1,  
    LOG_WARNING = 2,  
    LOG_INFO    = 3,  
    LOG_DEBUG   = 4  
};
```

Tip: You can call a logging function with different logging levels. All information under the value set for `AMD_LOG_LEVEL` is printed.

14.2 Logging mask

The logging mask is designed to print functionality types when you're running a HIP application. Once you set `AMD_LOG_LEVEL`, the logging mask is set as the default value (`0x7FFFFFFF`). You can change this to any of the valid values:

```
enum LogMask {
    LOG_API      = 0x00000001, //!< API call
    LOG_CMD      = 0x00000002, //!< Kernel and Copy Commands and Barriers
    LOG_WAIT     = 0x00000004, //!< Synchronization and waiting for commands to finish
    LOG_AQL      = 0x00000008, //!< Decode and display AQL packets
    LOG_QUEUE    = 0x00000010, //!< Queue commands and queue contents
    LOG_SIG      = 0x00000020, //!< Signal creation, allocation, pool
    LOG_LOCK     = 0x00000040, //!< Locks and thread-safety code.
    LOG_KERN     = 0x00000080, //!< kernel creations and arguments, etc.
    LOG_COPY     = 0x00000100, //!< Copy debug
    LOG_COPY2    = 0x00000200, //!< Detailed copy debug
    LOG_RESOURCE = 0x00000400, //!< Resource allocation, performance-impacting events.
    LOG_INIT     = 0x00000800, //!< Initialization and shutdown
    LOG_MISC     = 0x00001000, //!< misc debug, not yet classified
    LOG_AQL2     = 0x00002000, //!< Show raw bytes of AQL packet
    LOG_CODE     = 0x00004000, //!< Show code creation debug
    LOG_CMD2     = 0x00008000, //!< More detailed command info, including barrier commands
    LOG_LOCATION = 0x00010000, //!< Log message location
    LOG_MEM      = 0x00020000, //!< Memory allocation
    LOG_MEM_POOL = 0x00040000, //!< Memory pool allocation, including memory in graphs
    LOG_ALWAYS   = 0xFFFFFFFF, //!< Log always even mask flag is zero
};
```

You can also define the logging mask via the `AMD_LOG_MASK` environment variable.

14.3 Logging command

You can use the following code to print HIP logging information:

```
#define CIPrint(level, mask, format, ...) \
do { \
    if (AMD_LOG_LEVEL >= level) { \
        if (AMD_LOG_MASK & mask || mask == amd::LOG_ALWAYS) { \
            if (AMD_LOG_MASK & amd::LOG_LOCATION) { \
                amd::log_printf(level, __FILENAME__, __LINE__, format, ##__VA_ARGS__); \
            } else { \
                amd::log_printf(level, "", 0, format, ##__VA_ARGS__); \
            } \
        } \
    } \
} while (false)
```

Using HIP code, call the `CIPrint()` function with the desired input variables. For example:

```
CIPrint(amd::LOG_INFO, amd::LOG_INIT, "Initializing HSA stack.");
```


14.4 Logging examples

On **Linux**, you can enable HIP logging and retrieve logging information when you run `hipinfo`.

```

user@user-test:~/hip/bin$ export AMD_LOG_LEVEL=4
user@user-test:~/hip/bin$ ./hipinfo

:3:rocdevice.cpp           :453 : 23647210092: Initializing HSA stack.
:3:comgrctx.cpp            :33  : 23647639336: Loading COMGR library.
:3:rocdevice.cpp           :203 : 23647687108: Numa select cpu.
↪agent[0]=0x13407c0(fine=0x13409a0,coarse=0x1340ad0) for gpu agent=0x1346150
:4:runtime.cpp             :82  : 23647698669: init
:3:hip_device_runtime.cpp  :473 : 23647698869: 5617 : [7fad295dd840] hipGetDeviceCount:
↪Returned hipSuccess
:3:hip_device_runtime.cpp  :502 : 23647698990: 5617 : [7fad295dd840] hipSetDevice ( 0 )
:3:hip_device_runtime.cpp  :507 : 23647699042: 5617 : [7fad295dd840] hipSetDevice:
↪Returned hipSuccess

-----
device#                      0
:3:hip_device.cpp           :150 : 23647699276: 5617 : [7fad295dd840]
↪hipGetDeviceProperties ( 0x7ffdb7db730, 0 )
:3:hip_device.cpp           :237 : 23647699335: 5617 : [7fad295dd840]
↪hipGetDeviceProperties: Returned hipSuccess
Name:                        Device 7341
pciBusID:                    3
pciDeviceID:                 0
pciDomainID:                 0
multiProcessorCount:         11
maxThreadsPerMultiProcessor: 2560
isMultiGpuBoard:             0
clockRate:                   1900 Mhz
memoryClockRate:             875 Mhz
memoryBusWidth:              0
clockInstructionRate:        1000 Mhz
totalGlobalMem:              7.98 GB
maxSharedMemoryPerMultiProcessor: 64.00 KB
totalConstMem:               8573157376
sharedMemPerBlock:           64.00 KB
canMapHostMemory:            1
regsPerBlock:                0
warpSize:                    32
l2CacheSize:                 0
computeMode:                 0
maxThreadsPerBlock:          1024
maxThreadsDim.x:             1024
maxThreadsDim.y:             1024
maxThreadsDim.z:             1024
maxGridSize.x:               2147483647
maxGridSize.y:               2147483647
maxGridSize.z:               2147483647
major:                       10
minor:                        12
concurrentKernels:           1

```

(continues on next page)

(continued from previous page)

```

cooperativeLaunch:          0
cooperativeMultiDeviceLaunch: 0
arch.hasGlobalInt32Atomics: 1
...
gcnArch:                    1012
isIntegrated:               0
maxTexture1D:               65536
maxTexture2D.width:         16384
maxTexture2D.height:        16384
maxTexture3D.width:         2048
maxTexture3D.height:        2048
maxTexture3D.depth:         2048
isLargeBar:                 0
:3:hip_device_runtime.cpp   :471 : 23647701557: 5617 : [7fad295dd840] hipGetDeviceCount_
↪ ( 0x7ffdb7db714 )
:3:hip_device_runtime.cpp   :473 : 23647701608: 5617 : [7fad295dd840] hipGetDeviceCount:
↪ Returned hipSuccess
:3:hip_peer.cpp             :76  : 23647701731: 5617 : [7fad295dd840]
↪ hipDeviceCanAccessPeer ( 0x7ffdb7db728, 0, 0 )
:3:hip_peer.cpp             :60  : 23647701784: 5617 : [7fad295dd840] canAccessPeer:
↪ Returned hipSuccess
:3:hip_peer.cpp             :77  : 23647701831: 5617 : [7fad295dd840]
↪ hipDeviceCanAccessPeer: Returned hipSuccess
peers:
:3:hip_peer.cpp             :76  : 23647701921: 5617 : [7fad295dd840]
↪ hipDeviceCanAccessPeer ( 0x7ffdb7db728, 0, 0 )
:3:hip_peer.cpp             :60  : 23647701965: 5617 : [7fad295dd840] canAccessPeer:
↪ Returned hipSuccess
:3:hip_peer.cpp             :77  : 23647701998: 5617 : [7fad295dd840]
↪ hipDeviceCanAccessPeer: Returned hipSuccess
non-peers:                  device#0
:3:hip_memory.cpp           :345 : 23647702191: 5617 : [7fad295dd840] hipMemGetInfo (
↪ 0x7ffdb7db718, 0x7ffdb7db720 )
:3:hip_memory.cpp           :360 : 23647702243: 5617 : [7fad295dd840] hipMemGetInfo:
↪ Returned hipSuccess
memInfo.total:              7.98 GB
memInfo.free:               7.98 GB (100%)

```

On **Windows**, you can set `AMD_LOG_LEVEL` via environment variable from the advanced system settings or the command prompt (when run as administrator). The following example shows debug log information when calling the backend runtime.

```

C:\hip\bin>set AMD_LOG_LEVEL=4
C:\hip\bin>hipinfo
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\vdi\device\comgrctx.cpp:33 :
↪ 605413686305 us: 29864: [tid:0x9298] Loading COMGR library.
:4:C:\constructicon\builds\gfx\two\22.40\drivers\compute\vdi\platform\runtime.cpp:83 :
↪ 605413869411 us: 29864: [tid:0x9298] init
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_context.cpp:47 :
↪ : 605413869502 us: 29864: [tid:0x9298] Direct Dispatch: 0
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_device_runtime.

```

(continues on next page)

(continued from previous page)

```

→cpp:543 : 605413870553 us: 29864: [tid:0x9298] hipGetDeviceCount: Returned hipSuccess :
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_device_runtime.
→cpp:556 : 605413870631 us: 29864: [tid:0x9298] ←[32m hipSetDevice ( 0 ) ←[0m
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_device_runtime.
→cpp:561 : 605413870848 us: 29864: [tid:0x9298] hipSetDevice: Returned hipSuccess :
-----
device#                                0
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_device.cpp:346 :_
→605413871623 us: 29864: [tid:0x9298] ←[32m hipGetDeviceProperties ( 0000008AEBEFF8C8, _
→0 ) ←[0m
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_device.cpp:348 :_
→605413871695 us: 29864: [tid:0x9298] hipGetDeviceProperties: Returned hipSuccess :
Name:                                AMD Radeon(TM) Graphics
pciBusID:                            3
pciDeviceID:                         0
pciDomainID:                         0
multiProcessorCount:                 7
maxThreadsPerMultiProcessor:         2560
isMultiGpuBoard:                     0
clockRate:                           1600 Mhz
memoryClockRate:                     1333 Mhz
memoryBusWidth:                       0
totalGlobalMem:                      12.06 GB
totalConstMem:                       2147483647
sharedMemPerBlock:                    64.00 KB
...
gcnArchName:                          gfx90c:xnack-
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_device_runtime.
→cpp:541 : 605413924779 us: 29864: [tid:0x9298] ←[32m hipGetDeviceCount ( _
→0000008AEBEFF8A4 ) ←[0m
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_device_runtime.
→cpp:543 : 605413925075 us: 29864: [tid:0x9298] hipGetDeviceCount: Returned hipSuccess :
peers:                               :3:C:\constructicon\builds\gfx\two\22.40\drivers\
→compute\hipamd\src\hip_peer.cpp:176 : 605413928643 us: 29864: [tid:0x9298] ←[32m _
→hipDeviceCanAccessPeer ( 0000008AEBEFF890, 0, 0 ) ←[0m
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_peer.cpp:177 :_
→605413928743 us: 29864: [tid:0x9298] hipDeviceCanAccessPeer: Returned hipSuccess :
non-peers:                           :3:C:\constructicon\builds\gfx\two\22.40\drivers\
→compute\hipamd\src\hip_peer.cpp:176 : 605413930830 us: 29864: [tid:0x9298] ←[32m _
→hipDeviceCanAccessPeer ( 0000008AEBEFF890, 0, 0 ) ←[0m
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_peer.cpp:177 :_
→605413930882 us: 29864: [tid:0x9298] hipDeviceCanAccessPeer: Returned hipSuccess :
device#0
...
:4:C:\constructicon\builds\gfx\two\22.40\drivers\compute\vdi\device\pal\palmemory.
→cpp:430 : 605414517802 us: 29864: [tid:0x9298] Free:-      8000 bytes, VM[ 3007c8000, _
→3007d0000]
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\vdi\device\devprogram.cpp:2979:_
→605414517893 us: 29864: [tid:0x9298] For Init/Fini: Kernel Name: __amd_rocclr_
→copyBufferToImage
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\vdi\device\devprogram.cpp:2979:_
→605414518259 us: 29864: [tid:0x9298] For Init/Fini: Kernel Name: __amd_rocclr_

```

(continues on next page)

(continued from previous page)

```
→copyBuffer
...
:4:C:\constructicon\builds\gfx\two\22.40\drivers\compute\vdi\device\pal\palmemory.
→cpp:206 : 605414523422 us: 29864: [tid:0x9298] Alloc: 100000 bytes,
→ptr[000000003008D0000-000000003009D0000], obj[000000003007D0000-000000003047D0000]
:4:C:\constructicon\builds\gfx\two\22.40\drivers\compute\vdi\device\pal\palmemory.
→cpp:206 : 605414523767 us: 29864: [tid:0x9298] Alloc: 100000 bytes,
→ptr[000000003009D0000-00000000300AD0000], obj[000000003007D0000-000000003047D0000]
:3:C:\constructicon\builds\gfx\two\22.40\drivers\compute\hipamd\src\hip_memory.cpp:681 :
→605414524092 us: 29864: [tid:0x9298] hipMemGetInfo: Returned hipSuccess :
memInfo.total:          12.06 GB
memInfo.free:           11.93 GB (99%)
```

FREQUENTLY ASKED QUESTIONS

15.1 What APIs and features does HIP support?

HIP provides the following:

- Devices (`hipSetDevice()`, `hipGetDeviceProperties()`, etc.)
- Memory management (`hipMalloc()`, `hipMemcpy()`, `hipFree()`, etc.)
- Streams (`hipStreamCreate()`, `hipStreamSynchronize()`, `hipStreamWaitEvent()`, etc.)
- Events (`hipEventRecord()`, `hipEventElapsedTime()`, etc.)
- Kernel launching (`hipLaunchKernel`/`hipLaunchKernelGGL` is the preferred way of launching kernels. `hipLaunchKernelGGL` is a standard C/C++ macro that can serve as an alternative way to launch kernels, replacing the CUDA triple-chevron (`<<< >>>`) syntax).
- HIP Module API to control when and how code is loaded.
- CUDA-style kernel coordinate functions (`threadIdx`, `blockIdx`, `blockDim`, `gridDim`)
- Cross-lane instructions including `shfl`, `ballot`, `any`, `all`
- Most device-side math built-ins
- Error reporting (`hipGetLastError()`, `hipGetErrorString()`)

The HIP API documentation describes each API and its limitations, if any, compared with the equivalent CUDA API.

15.2 What is not supported?

15.2.1 Runtime/Driver API features

At a high-level, the following features are not supported:

- Textures (partial support available)
- Dynamic parallelism (CUDA 5.0)
- Graphics interoperability with OpenGL or Direct3D
- CUDA IPC Functions (Under Development)
- CUDA array, `mipmappedArray` and pitched memory
- Queue priority controls

See the [API Support Table](#) for more detailed information.

15.2.2 Kernel language features

- C++ style device-side dynamic memory allocations (free, new, delete) (CUDA 4.0)
- Virtual functions, indirect functions and try/catch (CUDA 4.0)
- `__prof_trigger`
- PTX assembly (CUDA 4.0). HIP-Clang supports inline GCN assembly.
- Several kernel features are under development. See the [C++ Language Extensions](#) for more information.

15.3 Is HIP a drop-in replacement for CUDA?

No. HIP provides porting tools which do most of the work to convert CUDA code into portable C++ code that uses the HIP APIs. Most developers will port their code from CUDA to HIP and then maintain the HIP version. HIP code provides the same performance as native CUDA code, plus the benefits of running on AMD platforms.

15.4 What specific version of CUDA does HIP support?

HIP APIs and features do not map to a specific CUDA version. HIP provides a strong subset of the functionality provided in CUDA, and the hipify tools can scan code to identify any unsupported CUDA functions - this is useful for identifying the specific features required by a given application.

However, we can provide a rough summary of the features included in each CUDA SDK and the support level in HIP. Each bullet below lists the major new language features in each CUDA release and then indicate which are supported/not supported in HIP:

- CUDA 4.0 and earlier :
 - HIP supports CUDA 4.0 except for the limitations described above.
- CUDA 5.0 :
 - Dynamic Parallelism (not supported)
 - `cuIpc` functions (under development).
- CUDA 6.0 :
 - Managed memory (under development)
- CUDA 6.5 :
 - `__shfl` intrinsic (supported)
- CUDA 7.0 :
 - Per-thread default streams (supported)
 - C++11 (Hip-Clang supports all of C++11, all of C++14 and some C++17 features)
- CUDA 7.5 :
 - `float16` (supported)
- CUDA 8.0 :
 - Page Migration including `cudaMemAdvise`, `cudaMemPrefetch`, other `cudaMem*` APIs(not supported)
- CUDA 9.0 :

- Cooperative Launch, Surface Object Management, Version Management

15.5 What libraries does HIP support?

HIP includes growing support for the four key math libraries using hipBLAS, hipFFT, hipRAND and hipSPARSE, as well as MIOpen for machine intelligence applications. These offer pointer-based memory interfaces (as opposed to opaque buffers) and can be easily interfaced with other HIP applications. The hip interfaces support both ROCm and CUDA paths, with familiar library interfaces.

- `hipBLAS`, which utilizes `rocBlas`.
- `hipFFT`
- `hipsPARSE`
- `hipRAND`
- `MIOpen`

Additionally, some of the cuBLAS routines are automatically converted to hipblas equivalents by the HIPIFY tools. These APIs use cuBLAS or hcBLAS depending on the platform and replace the need to use conditional compilation.

15.6 How does HIP compare with OpenCL?

Both AMD and NVIDIA support OpenCL 1.2 on their devices so that developers can write portable code. HIP offers several benefits over OpenCL:

- Developers can code in C++ as well as mix host and device C++ code in their source files. HIP C++ code can use templates, lambdas, classes and so on.
- The HIP API is less verbose than OpenCL and is familiar to CUDA developers.
- Because both CUDA and HIP are C++ languages, porting from CUDA to HIP is significantly easier than porting from CUDA to OpenCL.
- HIP uses the best available development tools on each platform: on NVIDIA GPUs, HIP code compiles using NVCC and can employ the Nsight profiler and debugger (unlike OpenCL on NVIDIA GPUs).
- HIP provides pointers and host-side pointer arithmetic.
- HIP provides device-level control over memory allocation and placement.
- HIP offers an offline compilation model.

15.7 How does porting CUDA to HIP compare to porting CUDA to OpenCL?

Both HIP and CUDA are dialects of C++, and thus porting between them is relatively straightforward. Both dialects support templates, classes, lambdas, and other C++ constructs. As one example, the hipify-perl tool was originally a Perl script that used simple text conversions from CUDA to HIP. HIP and CUDA provide similar math library calls as well. In summary, the HIP philosophy was to make the HIP language close enough to CUDA that the porting effort is relatively simple. This reduces the potential for error, and also makes it easy to automate the translation. HIP goal is to quickly get the ported program running on both platforms with little manual intervention, so that the programmer can focus on performance optimizations.

There have been several tools that have attempted to convert CUDA into OpenCL, such as CU2CL. OpenCL is a C99-based kernel language (rather than C++) and also does not support single-source compilation. As a result, the OpenCL syntax is different from CUDA, and the porting tools have to perform some heroic transformations to bridge this gap. The tools also struggle with more complex CUDA applications, in particular, those that use templates, classes, or other C++ features inside the kernel.

15.8 What hardware does HIP support?

- For AMD platforms, see the [ROCm documentation](#) for the list of supported platforms.
- For NVIDIA platforms, HIP requires unified memory and should run on any device supporting CUDA SDK 6.0 or newer. We have tested the NVIDIA Titan and Tesla K40.

15.9 Do HIPIFY tools automatically convert all source code?

Typically, HIPIFY tools can automatically convert almost all run-time code. Most device code needs no additional conversion since HIP and CUDA have similar names for math and built-in functions. The hipify-clang tool will automatically modify the kernel signature as needed (automating a step that used to be done manually). Additional porting may be required to deal with architecture feature queries or with CUDA capabilities that HIP doesn't support. In general, developers should always expect to perform some platform-specific tuning and optimization.

15.10 What is NVCC?

NVCC is NVIDIA's compiler driver for compiling "CUDA C++" code into PTX or device code for NVIDIA GPUs. It's a closed-source binary compiler that is provided by the CUDA SDK.

15.11 What is HIP-Clang?

HIP-Clang is a Clang/LLVM based compiler to compile HIP programs which can run on AMD platform.

15.12 Why use HIP rather than supporting CUDA directly?

While HIP is a strong subset of the CUDA, it is a subset. The HIP layer allows that subset to be clearly defined and documented. Developers who code to the HIP API can be assured their code will remain portable across NVIDIA and AMD platforms. In addition, HIP defines portable mechanisms to query architectural features and supports a larger 64-bit WaveSize which expands the return type for cross-lane functions like ballot and shuffle from 32-bit integers to 64-bit integers.

15.13 Can I develop HIP code on an NVIDIA CUDA platform?

Yes. HIP's CUDA path only exposes the APIs and functionality that work on both NVCC and AMDGPU back-ends. "Extra" APIs, parameters, and features which exist in CUDA but not in HIP-Clang will typically result in compile-time or run-time errors. Developers need to use the HIP API for most accelerator code and bracket any CUDA-specific code with preprocessor conditionals. Developers concerned about portability should, of course, run on both platforms, and should expect to tune for performance. In some cases, CUDA has a richer set of modes for some APIs, and some C++ capabilities such as virtual functions - see the HIP @API documentation for more details.

15.14 Can I develop HIP code on an AMD HIP-Clang platform?

Yes. HIP's HIP-Clang path only exposes the APIs and functions that work on AMD runtime back ends. "Extra" APIs, parameters and features that appear in HIP-Clang but not CUDA will typically cause compile- or run-time errors. Developers must use the HIP API for most accelerator code and bracket any HIP-Clang specific code with preprocessor conditionals. Those concerned about portability should, of course, test their code on both platforms and should tune it for performance. Typically, HIP-Clang supports a more modern set of C++11/C++14/C++17 features, so HIP developers who want portability should be careful when using advanced C++ features on the HIP-Clang path.

15.15 How to use HIP-Clang to build HIP programs?

The environment variable can be used to set compiler path:

- `HIP_CLANG_PATH`: path to hip-clang. When set, this variable let hipcc to use hip-clang for compilation/linking.

There is an alternative environment variable to set compiler path:

- `HIP_ROCCLR_HOME`: path to root directory of the HIP-ROCclr runtime. When set, this variable let hipcc use hip-clang from the ROCclr distribution. NOTE: If `HIP_ROCCLR_HOME` is set, there is no need to set `HIP_CLANG_PATH` since hipcc will deduce them from `HIP_ROCCLR_HOME`.

15.16 What is AMD clr?

AMD Compute Language Runtime (CLR) is a repository for the AMD platform, which contains source codes for AMD's compute languages runtimes as follows,

- `hipamd` - contains implementation of HIP language for AMD GPU.
- `rocclr` - contains virtual device interfaces that compute runtimes interact with backends, such as ROCr on Linux and PAL on Windows.
- `opencl` - contains implementation of OpenCL™ on the AMD platform.

15.17 What is hipother?

A new repository ‘hipother’ is added in the ROCm 6.1 release, which is branched out from HIP. hipother supports the HIP back-end implementation on some non-AMD platforms, like NVIDIA.

15.18 Can I get HIP open source repository for Windows?

No, there is no HIP repository open publicly on Windows.

15.19 Can a HIP binary run on both AMD and NVIDIA platforms?

HIP is a source-portable language that can be compiled to run on either AMD or NVIDIA platform. HIP tools don’t create a “fat binary” that can run on either platform, however.

15.20 On HIP-Clang, can I link HIP code with host code compiled with another compiler such as gcc, icc, or clang?

Yes. HIP generates the object code which conforms to the GCC ABI, and also links with libstdc++. This means you can compile host code with the compiler of your choice and link the generated object code with GPU code compiled with HIP. Larger projects often contain a mixture of accelerator code (initially written in CUDA with NVCC) and host code (compiled with gcc, icc, or clang). These projects can convert the accelerator code to HIP, compile that code with hipcc, and link with object code from their preferred compiler.

15.21 Can HIP API support C style application? What is the difference between C and C++?

HIP is C++ runtime API that supports C style applications as well.

Some C style applications (and interfaces to other languages (FORTRAN, Python)) would call certain HIP APIs but not use kernel programming. They can be compiled with a C compiler and run correctly, however, small details must be considered in the code. For example, initialization, as shown in the simple application below, uses HIP structs dim3 with the file name “test.hip.cpp”

```
#include "hip/hip_runtime_api.h"
#include "stdio.h"

int main(int argc, char** argv) {
    dim3 grid1;
    printf("dim3 grid1; x=%d, y=%d, z=%d\n", grid1.x, grid1.y, grid1.z);
    dim3 grid2 = {1,1,1};
    printf("dim3 grid2 = {1,1,1}; x=%d, y=%d, z=%d\n", grid2.x, grid2.y, grid2.z);
    return 0;
}
```

When using a C++ compiler,

```
$ gcc -x c++ $(hipconfig --cpp_config) test3.hip.cpp -o test
$ ./test
dim3 grid1; x=1, y=1, z=1
dim3 grid2 = {1,1,1}; x=1, y=1, z=1
```

In which “dim3 grid1;” will yield a dim3 grid with all dimensional members x,y,z initialized to 1, as the default constructor behaves that way. Further, if written:

```
dim3 grid(2); // yields {2,1,1}
dim3 grid(2,3); yields {2,3,1}
```

In comparison, when using the C compiler,

```
$ gcc -x c $(hipconfig --cpp_config) test.hip.cpp -o test
$ ./test
dim3 grid1; x=646881376, y=21975, z=1517277280
dim3 grid2 = {1,1,1}; x=1, y=1, z=1
```

In which “dim3 grid;” does not imply any initialization, no constructor is called, and dimensional values x,y,z of grid are undefined. NOTE: To get the C++ default behavior, C programmers must additionally specify the right-hand side as shown below,

```
dim3 grid = {1,1,1}; // initialized as in C++
```

15.22 Can I install both CUDA SDK and HIP-Clang on the same machine?

Yes. You can use HIP_PLATFORM to choose which path hipcc targets. This configuration can be useful when using HIP to develop an application which is portable to both AMD and NVIDIA.

15.23 HIP detected my platform (HIP-Clang vs NVCC) incorrectly * what should I do?

HIP will set the platform to AMD and use HIP-Clang as compiler if it sees that the AMD graphics driver is installed and has detected an AMD GPU. Sometimes this isn't what you want * you can force HIP to recognize the platform by setting the following,

```
export HIP_PLATFORM=amd
```

HIP then set and use correct AMD compiler and runtime, HIP_COMPILER=clang HIP_RUNTIME=roclr

To choose NVIDIA platform, you can set,

```
export HIP_PLATFORM=nvidia
```

In this case, HIP will set and use the following,

```
HIP_COMPILER=cuda
HIP_RUNTIME=nvcc
```

One symptom of this problem is the message “error: ‘unknown error’(11) at `square.hipref.cpp:56`. This can occur if you have a CUDA installation on an AMD platform, and HIP incorrectly detects the platform as NVCC. HIP may be able to compile the application using the NVCC tool-chain but will generate this error at runtime since the platform does not have a CUDA device.

15.24 On CUDA, can I mix CUDA code with HIP code?

Yes. Most HIP data structures (`hipStream_t`, `hipEvent_t`) are typedefs to CUDA equivalents and can be intermixed. Both CUDA and HIP use integer device ids. One notable exception is that `hipError_t` is a new type, and cannot be used where a `cudaError_t` is expected. In these cases, refactor the code to remove the expectation. Alternatively, `hip_runtime_api.h` defines functions which convert between the error code spaces:

```
hipErrorToCudaError hipCudaErrorTohipError hipCudaResultTohipError
```

If platform portability is important, use `#ifdef __HIP_PLATFORM_NVIDIA__` to guard the CUDA-specific code.

15.25 How do I trace HIP application flow?

See *Logging HIP activity* for more information.

15.26 What are the maximum limits of kernel launch parameters?

Product of `block.x`, `block.y`, and `block.z` should be less than 1024. Please note, HIP does not support kernel launch with total work items defined in dimension with size `gridDim.x * blockDim.x`, `gridDim.y * blockDim.y` and `gridDim.z * blockDim.z` are always less than 2^{32} .

15.27 Are `__shfl*_sync` functions supported on HIP platform?

`__shfl*_sync` is not supported on HIP but for NVCC path CUDA 9.0 and above all shuffle calls get redirected to its sync version.

15.28 How to create a guard for code that is specific to the host or the GPU?

The compiler defines the `__HIP_DEVICE_COMPILE__` macro only when compiling the code for the GPU. It could be used to guard code that is specific to the host or the GPU.

15.29 Why `_OpenMP` is undefined when compiling with `-fopenmp`?

When compiling an OpenMP source file with `hipcc -fopenmp`, the compiler may generate error if there is a reference to the `_OPENMP` macro. This is due to a limitation in `hipcc` that treats any source file type (for example `.cpp`) as an HIP translation unit leading to some conflicts with the OpenMP language switch. If the OpenMP source file doesn't contain any HIP language constructs you could work around this issue by adding the `-x c++` switch to force the compiler to treat the file as regular C++. Another approach would be to guard the OpenMP code with `#ifdef _OPENMP` so that the code block is disabled when compiling for the GPU. The `__HIP_DEVICE_COMPILE__` macro defined by the HIP compiler when compiling GPU code could also be used for guarding code paths specific to the host or the GPU.

15.30 Does the HIP-Clang compiler support extern shared declarations?

Previously, it was essential to declare dynamic shared memory using the `HIP_DYNAMIC_SHARED` macro for accuracy, as using static shared memory in the same kernel could result in overlapping memory ranges and data-races.

Now, the HIP-Clang compiler provides support for extern shared declarations, and the `HIP_DYNAMIC_SHARED` option is no longer required. You may use the standard extern definition: `extern shared type var[];`

15.31 I have multiple HIP enabled devices and I am getting an error code `hipErrorSharedObjectInitFailed` with the message “Error: shared object initialization failed”?

This error message is seen due to the fact that you do not have valid code object for all of your devices.

If you have compiled the application yourself, make sure you have given the correct device name(s) and its features via: `--offload-arch`. If you are not mentioning the `--offload-arch`, make sure that `hipcc` is using the correct offload arch by verifying the `hipcc` output generated by setting the environment variable `HIPCC_VERBOSE=1`.

If you have a precompiled application/library (like `roclblas`, `TensorFlow` etc) which gives you such error, there are one of two possibilities.

- The application/library does not ship code object bundles for **all** of your device(s): in this case you need to recompile the application/library yourself with correct `--offload-arch`.
- The application/library does not ship code object bundles for **some** of your device(s), for example you have a system with an APU + GPU and the library does not ship code objects for your APU. For this you can set the environment variable `HIP_VISIBLE_DEVICES` or `CUDA_VISIBLE_DEVICES` on NVIDIA platform, to only enable GPUs for which code object is available. This will limit the GPUs visible to your application and allow it to run.

Note: In previous releases, the error code is `hipErrorNoBinaryForGpu` with message “Unable to find code object for all current devices”. The error code handling behavior is changed. HIP runtime shows the error code `hipErrorSharedObjectInitFailed` with message “Error: shared object initialization failed” on unsupported GPU.

15.32 How to use per-thread default stream in HIP?

The per-thread default stream is an implicit stream local to both the thread and the current device. It does not do any implicit synchronization with other streams (like explicitly created streams), or default per-thread stream on other threads.

The per-thread default stream is a blocking stream and will synchronize with the default null stream if both are used in a program.

In ROCm, a compilation option should be added in order to compile the translation unit with per-thread default stream enabled. `-fgpu-default-stream=per-thread`. Once source is compiled with per-thread default stream enabled, all APIs will be executed on per thread default stream, hence there will not be any implicit synchronization with other streams.

Besides, per-thread default stream be enabled per translation unit, users can compile some files with feature enabled and some with feature disabled. Feature enabled translation unit will have default stream as per thread and there will not be any implicit synchronization done but other modules will have legacy default stream which will do implicit synchronization.

15.33 How to use complex multiplication and division operations?

In HIP, `hipFloatComplex` and `hipDoubleComplex` are defined as complex data types,

```
typedef float2 hipFloatComplex;
typedef double2 hipDoubleComplex;
```

Any application uses complex multiplication and division operations, need to replace `*` and `/` operators with the following,

- `hipCmulf()` and `hipCdivf()` for `hipFloatComplex`
- `hipCmul()` and `hipCdiv()` for `hipDoubleComplex`

Note: These complex operations are equivalent to corresponding types/functions on the NVIDIA platform.

15.34 Can I develop applications with HIP APIs on Windows the same on Linux?

Yes, HIP APIs are available to use on both Linux and Windows. Due to different working mechanisms on operating systems like Windows vs Linux, HIP APIs call corresponding lower level backend runtime libraries and kernel drivers for the OS, in order to control the executions on GPU hardware accordingly. There might be a few differences on the related backend software and driver support, which might affect usage of HIP APIs. See OS support details in HIP API document.

15.35 Does HIP support LUID?

Starting ROCm 6.0, HIP runtime supports Locally Unique Identifier (LUID). This feature enables the local physical device(s) to interoperate with other devices. For example, DirectX 12.

HIP runtime sets device LUID properties so the driver can query LUID to identify each device for interoperability.

Note: HIP supports LUID only on Windows OS.

15.36 How can I know the version of HIP?

HIP version definition has been updated since ROCm 4.2 release as the following:

```
HIP_VERSION=HIP_VERSION_MAJOR * 100000000 + HIP_VERSION_MINOR * 100000 + HIP_VERSION_PATCH
```

HIP version can be queried from HIP API call,

```
hipRuntimeGetVersion(&runtimeVersion);
```

The version returned will always be greater than the versions in previous ROCm releases.

Note: The version definition of HIP runtime is different from CUDA. On AMD platform, the function returns HIP runtime version, while on NVIDIA platform, it returns CUDA runtime version. And there is no mapping/correlation between HIP version and CUDA version.

HIP RUNTIME API REFERENCE

16.1 Related Pages

16.2 Modules

16.3 Data Structures

16.3.1 Data Structures

16.3.2 Data Structure Index

16.3.3 Data Fields

16.3.3.1 All

16.3.3.1.1 Data Fields

16.3.3.1.2 Data Fields

16.3.3.1.3 Data Fields

16.3.3.1.4 Data Fields

16.3.3.1.5 Data Fields

16.3.3.1.6 Data Fields

16.3.3.1.7 Data Fields

16.3.3.1.8 Data Fields

16.3.3.1.9 Data Fields

16.3.3.1.10 Data Fields

16.3.3.1.11 Data Fields

16.3.3.1.12 Data Fields

16.3.3.1.13 Data Fields

C++ LANGUAGE EXTENSIONS

HIP provides a C++ syntax that is suitable for compiling most code that commonly appears in compute kernels (classes, namespaces, operator overloading, and templates). HIP also defines other language features that are designed to target accelerators, such as:

- A kernel-launch syntax that uses standard C++ (this resembles a function call and is portable to all HIP targets)
- Short-vector headers that can serve on a host or device
- Math functions that resemble those in `math.h`, which is included with standard C++ compilers
- Built-in functions for accessing specific GPU hardware capabilities

Note: This chapter describes the built-in variables and functions that are accessible from the HIP kernel. It's intended for users who are familiar with CUDA kernel syntax and want to learn how HIP differs from CUDA.

Features are labeled with one of the following keywords:

- **Supported:** HIP supports the feature with a CUDA-equivalent function
- **Not supported:** HIP does not support the feature
- **Under development:** The feature is under development and not yet available

17.1 Function-type qualifiers

17.1.1 `__device__`

Supported `__device__` functions are:

- Run on the device
- Called from the device only

You can combine `__device__` with the host keyword (`__host__`).

17.1.2 `__global__`

Supported `__global__` functions are:

- Run on the device
- Called (launched) from the host

HIP `__global__` functions must have a `void` return type. The first parameter in a HIP `__global__` function must have the type `hipLaunchParm`. Refer to [Kernel launch example](#) to see usage.

HIP doesn't support dynamic-parallelism, which means that you can't call `__global__` functions from the device.

17.1.3 `__host__`

Supported `__host__` functions are:

- Run on the host
- Called from the host

You can combine `__host__` with `__device__`; in this case, the function compiles for the host and the device. Note that these functions can't use the HIP grid coordinate functions (e.g., `threadIdx.x`). If you need to use HIP grid coordinate functions, you can pass the necessary coordinate information as an argument.

You can't combine `__host__` with `__global__`.

HIP parses the `__noinline__` and `__forceinline__` keywords and converts them into the appropriate Clang attributes.

17.2 Calling `__global__` functions

`__global__` functions are often referred to as *kernels*. When you call a global function, you're *launching a kernel*. When launching a kernel, you must specify an execution configuration that includes the grid and block dimensions. The execution configuration can also include other information for the launch, such as the amount of additional shared memory to allocate and the stream where you want to execute the kernel.

HIP introduces a standard C++ calling convention (`hipLaunchKernelGGL`) to pass the run configuration to the kernel. However, you can also use the CUDA `<<< >>>` syntax.

When using `hipLaunchKernelGGL`, your first five parameters must be:

- `symbol kernelName`: The name of the kernel you want to launch. To support template kernels that contain `" , "`, use the `HIP_KERNEL_NAME` macro (HIPIFY tools insert this automatically).
- `dim3 gridDim`: 3D-grid dimensions that specify the number of blocks to launch.
- `dim3 blockDim`: 3D-block dimensions that specify the number of threads in each block.
- `size_t dynamicShared`: The amount of additional shared memory that you want to allocate when launching the kernel (see [__shared__](#)).
- `hipStream_t`: The stream where you want to run the kernel. A value of `0` corresponds to the NULL stream (see [Synchronization functions](#)).

You can include your kernel arguments after these parameters.

```
// Example hipLaunchKernelGGL pseudocode:
__global__ MyKernel(hipLaunchParm lp, float *A, float *B, float *C, size_t N)
{
    ...
}

MyKernel<<<dim3(gridDim), dim3(groupDim), 0, 0>>> (a,b,c,n);

// Alternatively, you can launch the kernel using:
// hipLaunchKernelGGL(MyKernel, dim3(gridDim), dim3(groupDim), 0/*dynamicShared*/, 0/
//   ↳*stream), a, b, c, n);
```

You can use HIPIFY tools to convert CUDA launch syntax to `hipLaunchKernelGGL`. This includes the conversion of optional `<<< >>>` arguments into the five required `hipLaunchKernelGGL` parameters.

Note: HIP doesn't support dimension sizes of $gridDim * blockDim \geq 2^{32}$ when launching a kernel.

17.3 Kernel launch example

```
// Example showing device function, __device__ __host__
// <- compile for both device and host
float PlusOne(float x)
{
    return x + 1.0;
}

__global__
void
MyKernel (hipLaunchParm lp, /*lp parm for execution configuration */
          const float *a, const float *b, float *c, unsigned N)
{
    unsigned gid = threadIdx.x; // <- coordinate index function
    if (gid < N) {
        c[gid] = a[gid] + PlusOne(b[gid]);
    }
}

void callMyKernel()
{
    float *a, *b, *c; // initialization not shown...
    unsigned N = 1000000;
    const unsigned blockSize = 256;

    MyKernel<<<dim3(gridDim), dim3(groupDim), 0, 0>>> (a,b,c,n);
    // Alternatively, kernel can be launched by
    // hipLaunchKernelGGL(MyKernel, dim3(N/blockSize), dim3(blockSize), 0, 0, a,b,c,N);
}
```

17.4 Variable type qualifiers

17.4.1 `__constant__`

The host writes constant memory before launching the kernel. This memory is read-only from the GPU while the kernel is running. The functions for accessing constant memory are:

- `hipGetSymbolAddress()`
- `hipGetSymbolSize()`
- `hipMemcpyToSymbol()`
- `hipMemcpyToSymbolAsync()`
- `hipMemcpyFromSymbol()`
- `hipMemcpyFromSymbolAsync()`

17.4.2 `__shared__`

To allow the host to dynamically allocate shared memory, you can specify `extern __shared__` as a launch parameter.

Note: Prior to the HIP-Clang compiler, dynamic shared memory had to be declared using the `HIP_DYNAMIC_SHARED` macro in order to ensure accuracy. This is because using static shared memory in the same kernel could've resulted in overlapping memory ranges and data-races. The HIP-Clang compiler provides support for `extern __shared__` declarations, so `HIP_DYNAMIC_SHARED` is no longer required.

17.4.3 `__managed__`

Managed memory, including the `__managed__` keyword, is supported in HIP combined host/device compilation.

17.4.4 `__restrict__`

`__restrict__` tells the compiler that the associated memory pointer not to alias with any other pointer in the kernel or function. This can help the compiler generate better code. In most use cases, every pointer argument should use this keyword in order to achieve the benefit.

17.5 Built-in variables

17.5.1 Coordinate built-ins

The kernel uses coordinate built-ins (`thread*`, `block*`, `grid*`) to determine the coordinate index and bounds for the active work item.

Built-ins are defined in `amd_hip_runtime.h`, rather than being implicitly defined by the compiler.

Coordinate variable definitions for built-ins are the same for HIP and CUDA. For example: `threadIdx.x`, `blockIdx.y`, and `gridDim.y`. The products `gridDim.x * blockDim.x`, `gridDim.y * blockDim.y`, and `gridDim.z * blockDim.z` are always less than 2^{32} .

Coordinate built-ins are implemented as structures for improved performance. When used with `printf`, they must be explicitly cast to integer types.

17.5.2 warpSize

The `warpSize` variable type is `int`. It contains the warp size (in threads) for the target device. `warpSize` should only be used in device functions that develop portable wave-aware code.

Note: NVIDIA devices return 32 for this variable; AMD devices return 64 for gfx9 and 32 for gfx10 and above.

17.6 Vector types

The following vector types are defined in `hip_runtime.h`. They are not automatically provided by the compiler.

17.6.1 Short vector types

Short vector types derive from basic integer and floating-point types. These structures are defined in `hip_vector_types.h`. The first, second, third, and fourth components of the vector are defined by the `x`, `y`, `z`, and `w` fields, respectively. All short vector types support a constructor function of the form `make_<type_name>()`. For example, `float4 make_float4(float x, float y, float z, float w)` creates a vector with type `float4` and value `(x,y,z,w)`.

HIP supports the following short vector formats:

- Signed Integers:
 - `char1`, `char2`, `char3`, `char4`
 - `short1`, `short2`, `short3`, `short4`
 - `int1`, `int2`, `int3`, `int4`
 - `long1`, `long2`, `long3`, `long4`
 - `longlong1`, `longlong2`, `longlong3`, `longlong4`
- Unsigned Integers:
 - `uchar1`, `uchar2`, `uchar3`, `uchar4`
 - `ushort1`, `ushort2`, `ushort3`, `ushort4`
 - `uint1`, `uint2`, `uint3`, `uint4`
 - `ulong1`, `ulong2`, `ulong3`, `ulong4`
 - `ulonglong1`, `ulonglong2`, `ulonglong3`, `ulonglong4`
- Floating Points:
 - `float1`, `float2`, `float3`, `float4`
 - `double1`, `double2`, `double3`, `double4`

17.6.2 dim3

`dim3` is a three-dimensional integer vector type that is commonly used to specify grid and group dimensions.

The `dim3` constructor accepts between zero and three arguments. By default, it initializes unspecified dimensions to 1.

```
typedef struct dim3 {  
    uint32_t x;  
    uint32_t y;  
    uint32_t z;  
  
    dim3(uint32_t _x=1, uint32_t _y=1, uint32_t _z=1) : x(_x), y(_y), z(_z) {};  
};
```

17.7 Memory fence instructions

HIP supports `__threadfence()` and `__threadfence_block()`. If you're using `threadfence_system()` in the HIP-Clang path, you can use the following workaround:

1. Build HIP with the `HIP_COHERENT_HOST_ALLOC` environment variable enabled.
2. Modify kernels that use `__threadfence_system()` as follows:
 - Ensure the kernel operates only on fine-grained system memory, which should be allocated with `hipHostMalloc()`.
 - Remove `memcpy` for all allocated fine-grained system memory regions.

17.8 Synchronization functions

The `__syncthreads()` built-in function is supported in HIP. The `__syncthreads_count(int)`, `__syncthreads_and(int)`, and `__syncthreads_or(int)` functions are under development.

17.9 Math functions

HIP-Clang supports a set of math operations that are callable from the device. HIP supports most of the device functions supported by CUDA. These are described in the following sections.

17.9.1 Single precision mathematical functions

Following is the list of supported single precision mathematical functions.

Table 17.1: Single precision mathematical functions

Function	Supported on Host	Supported on Device
<code>float abs(float x)</code> Returns the absolute value of x	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>float acosf(float x)</code> Returns the arc cosine of x .	✓	✓
<code>float acoshf(float x)</code> Returns the nonnegative arc hyperbolic cosine of x .	✓	✓
<code>float asinf(float x)</code> Returns the arc sine of x .	✓	✓
<code>float asinhf(float x)</code> Returns the arc hyperbolic sine of x .	✓	✓
<code>float atanf(float x)</code> Returns the arc tangent of x .	✓	✓
<code>float atan2f(float x, float y)</code> Returns the arc tangent of the ratio of x and y .	✓	✓
<code>float atanhf(float x)</code> Returns the arc hyperbolic tangent of x .	✓	✓
<code>float cbrtf(float x)</code> Returns the cube root of x .	✓	✓
<code>float ceilf(float x)</code> Returns ceiling of x .	✓	✓
<code>float copysignf(float x, float y)</code> Create value with given magnitude, copying sign of second value.	✓	✓
<code>float cosf(float x)</code> Returns the cosine of x .	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>float coshf(float x)</code> Returns the hyperbolic cosine of x .	✓	✓
<code>float cospif(float x)</code> Returns the cosine of $\pi \cdot x$.	✓	✓
<code>float cyl_bessel_i0f(float x)</code> Returns the value of the regular modified cylindrical Bessel function of order 0 for x .		
<code>float cyl_bessel_i1f(float x)</code> Returns the value of the regular modified cylindrical Bessel function of order 1 for x .		
<code>float erff(float x)</code> Returns the error function of x .	✓	✓
<code>float erfcf(float x)</code> Returns the complementary error function of x .	✓	✓
<code>float erfcinvf(float x)</code> Returns the inverse complementary function of x .	✓	✓
<code>float erfcxf(float x)</code> Returns the scaled complementary error function of x .	✓	✓
<code>float erfinvf(float x)</code> Returns the inverse error function of x .	✓	✓
<code>float expf(float x)</code> Returns e^x .	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>float exp10f(float x)</code> Returns 10^x .	✓	✓
<code>float exp2f(float x)</code> Returns 2^x .	✓	✓
<code>float expm1f(float x)</code> Returns $\ln(x + 1)$	✓	✓
<code>float fabsf(float x)</code> Returns the absolute value of x	✓	✓
<code>float fdimf(float x, float y)</code> Returns the positive difference between x and y .	✓	✓
<code>float fdivdef(float x, float y)</code> Divide two floating point values.	✓	✓
<code>float floorf(float x)</code> Returns the largest integer less than or equal to x .	✓	✓
<code>float fmaf(float x, float y, float z)</code> Returns $x \cdot y + z$ as a single operation.	✓	✓
<code>float fmaxf(float x, float y)</code> Determine the maximum numeric value of x and y .	✓	✓
<code>float fminf(float x, float y)</code> Determine the minimum numeric value of x and y .	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>float fmodf(float x, float y)</code> Returns the floating-point remainder of x/y .	✓	✓
<code>float modff(float x, float* iptr)</code> Break down x into fractional and integral parts.	✓	
<code>float frexpf(float x, int* nptr)</code> Extract mantissa and exponent of x .	✓	
<code>float hypotf(float x, float y)</code> Returns the square root of the sum of squares of x and y .	✓	✓
<code>int ilogbf(float x)</code> Returns the unbiased integer exponent of x .	✓	✓
<code>bool isfinite(float x)</code> Determine whether x is finite.	✓	✓
<code>bool isinf(float x)</code> Determine whether x is infinite.	✓	✓
<code>bool isnan(float x)</code> Determine whether x is a NAN.	✓	✓
<code>float j0f(float x)</code> Returns the value of the Bessel function of the first kind of order 0 for x .	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>float j1f(float x)</code> Returns the value of the Bessel function of the first kind of order 1 for x .	✓	✓
<code>float jnf(int n, float x)</code> Returns the value of the Bessel function of the first kind of order n for x .	✓	✓
<code>float ldexpf(float x, int exp)</code> Returns the natural logarithm of the absolute value of the gamma function of x .	✓	✓
<code>float lgammaf(float x)</code> Returns the natural logarithm of the absolute value of the gamma function of x .	✓	
<code>long int lrintf(float x)</code> Round x to nearest integer value.	✓	✓
<code>long long int llrintf(float x)</code> Round x to nearest integer value.	✓	✓
<code>long int lroundf(float x)</code> Round to nearest integer value.	✓	✓
<code>long long int llroundf(float x)</code> Round to nearest integer value.	✓	✓
<code>float log10f(float x)</code> Returns the base 10 logarithm of x .	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>float log1pf(float x)</code> Returns the natural logarithm of $x + 1$.	✓	✓
<code>float log2f(float x)</code> Returns the base 2 logarithm of x .	✓	✓
<code>float logf(float x)</code> Returns the natural logarithm of x .	✓	✓
<code>float logbf(float x)</code> Returns the floating point representation of the exponent of x .	✓	✓
<code>float nanf(const char* tagp)</code> Returns “Not a Number” value.		✓
<code>float nearbyintf(float x)</code> Round x to the nearest integer.	✓	✓
<code>float nextafterf(float x, float y)</code> Returns next representable single-precision floating-point value after argument.	✓	
<code>float norm3df(float x, float y, float z)</code> Returns the square root of the sum of squares of x , y and z .	✓	✓
<code>float norm4df(float x, float y, float z, float w)</code> Returns the square root of the sum of squares of x , y , z and w .	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>float normcdf(float y)</code> Returns the standard normal cumulative distribution function.	✓	✓
<code>float normcdfinv(float y)</code> Returns the inverse of the standard normal cumulative distribution function.	✓	✓
<code>float normf(int dim, const float *a)</code> Returns the square root of the sum of squares of any number of coordinates.	✓	✓
<code>float powf(float x, float y)</code> Returns x^y .	✓	✓
<code>float powif(float base, int iexp)</code> Returns the value of first argument to the power of second argument.	✓	✓
<code>float remainderf(float x, float y)</code> Returns single-precision floating-point remainder.	✓	✓
<code>float remquof(float x, float y, int* quo)</code> Returns single-precision floating-point remainder and part of quotient.	✓	✓
<code>float roundf(float x)</code> Round to nearest integer value in floating-point.	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>float rcbrtf(float x)</code> Returns the reciprocal cube root function.	✓	✓
<code>float rhypotf(float x, float y)</code> Returns one over the square root of the sum of squares of two arguments.	✓	✓
<code>float rintf(float x)</code> Round input to nearest integer value in floating-point.	✓	✓
<code>float rnorm3df(float x, float y, float z)</code> Returns one over the square root of the sum of squares of three coordinates of the argument.	✓	✓
<code>float rnorm4df(float x, float y, float z, float w)</code> Returns one over the square root of the sum of squares of four coordinates of the argument.	✓	✓
<code>float rnormf(int dim, const float *a)</code> Returns the reciprocal of square root of the sum of squares of any number of coordinates.	✓	✓
<code>float scalblnf(float x, long int n)</code> Scale x by 2^n .	✓	✓
<code>float scalbnf(float x, int n)</code> Scale x by 2^n .	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>bool signbit(float x)</code> Return the sign bit of x .	✓	✓
<code>float sinf(float x)</code> Returns the sine of x .	✓	✓
<code>float sinhf(float x)</code> Returns the hyperbolic sine of x .	✓	✓
<code>float sinpif(float x)</code> Returns the hyperbolic sine of $\pi \cdot x$.	✓	✓
<code>void sincosf(float x, float *sptr, float *cptr)</code> Returns the sine and cosine of x .	✓	✓
<code>void sincospif(float x, float *sptr, float *cptr)</code> Returns the sine and cosine of $\pi \cdot x$.	✓	✓
<code>float sqrtf(float x)</code> Returns the square root of x .	✓	✓
<code>float rsqrtf(float x)</code> Returns the reciprocal of the square root of x .		✓
<code>float tanf(float x)</code> Returns the tangent of x .	✓	✓
<code>float tanhf(float x)</code> Returns the hyperbolic tangent of x .	✓	✓
<code>float tgammaf(float x)</code> Returns the gamma function of x .	✓	✓

continues on next page

Table 17.1 – continued from previous page

<code>float truncf(float x)</code> Truncate x to the integral part.	✓	✓
<code>float y0f(float x)</code> Returns the value of the Bessel function of the second kind of order 0 for x .	✓	✓
<code>float y1f(float x)</code> Returns the value of the Bessel function of the second kind of order 1 for x .	✓	✓
<code>float ynf(int n, float x)</code> Returns the value of the Bessel function of the second kind of order n for x .	✓	✓

17.9.2 Double precision mathematical functions

Following is the list of supported double precision mathematical functions.

Table 17.2: Double precision mathematical functions

Function	Supported on Host	Supported on Device
<code>double abs(double x)</code> Returns the absolute value of x	✓	✓
<code>double acos(double x)</code> Returns the arc cosine of x .	✓	✓
<code>double acosh(double x)</code> Returns the nonnegative arc hyperbolic cosine of x .	✓	✓
<code>double asin(double x)</code> Returns the arc sine of x .	✓	✓

continues on next page

Table 17.2 – continued from previous page

<code>double asinh(double x)</code> Returns the arc hyperbolic sine of x .	✓	✓
<code>double atan(double x)</code> Returns the arc tangent of x .	✓	✓
<code>double atan2(double x, double y)</code> Returns the arc tangent of the ratio of x and y .	✓	✓
<code>double atanh(double x)</code> Returns the arc hyperbolic tangent of x .	✓	✓
<code>double cbrt(double x)</code> Returns the cube root of x .	✓	✓
<code>double ceil(double x)</code> Returns ceiling of x .	✓	✓
<code>double copysign(double x, double y)</code> Create value with given magnitude, copying sign of second value.	✓	✓
<code>double cos(double x)</code> Returns the cosine of x .	✓	✓
<code>double cosh(double x)</code> Returns the hyperbolic cosine of x .	✓	✓
<code>double cospi(double x)</code> Returns the cosine of $\pi \cdot x$.	✓	✓

continues on next page

Table 17.2 – continued from previous page

double cyl_bessel_i0(double x) Returns the value of the regular modified cylindrical Bessel function of order 0 for x .		
double cyl_bessel_i1(double x) Returns the value of the regular modified cylindrical Bessel function of order 1 for x .		
double erf(double x) Returns the error function of x .	✓	✓
double erfc(double x) Returns the complementary error function of x .	✓	✓
double erfcinv(double x) Returns the inverse complementary function of x .	✓	✓
double erfcx(double x) Returns the scaled complementary error function of x .	✓	✓
double erfinv(double x) Returns the inverse error function of x .	✓	✓
double exp(double x) Returns e^x .	✓	✓
double exp10(double x) Returns 10^x .	✓	✓
double exp2(double x) Returns 2^x .	✓	✓

continues on next page

Table 17.2 – continued from previous page

<code>double expm1(double x)</code> Returns $\ln(x - 1)$	✓	✓
<code>double fabs(double x)</code> Returns the absolute value of x	✓	✓
<code>double fdim(double x, double y)</code> Returns the positive difference between x and y .	✓	✓
<code>double floor(double x)</code> Returns the largest integer less than or equal to x .	✓	✓
<code>double fma(double x, double y, double z)</code> Returns $x \cdot y + z$ as a single operation.	✓	✓
<code>double fmax(double x, double y)</code> Determine the maximum numeric value of x and y .	✓	✓
<code>double fmin(double x, double y)</code> Determine the minimum numeric value of x and y .	✓	✓
<code>double fmod(double x, double y)</code> Returns the floating-point remainder of x/y .	✓	✓
<code>double modf(double x, double* iptr)</code> Break down x into fractional and integral parts.	✓	

continues on next page

Table 17.2 – continued from previous page

<code>double frexp(double x, int* nptr)</code> Extract mantissa and exponent of x .	✓	
<code>double hypot(double x, double y)</code> Returns the square root of the sum of squares of x and y .	✓	✓
<code>int ilogb(double x)</code> Returns the unbiased integer exponent of x .	✓	✓
<code>bool isfinite(double x)</code> Determine whether x is finite.	✓	✓
<code>bool isin(double x)</code> Determine whether x is infinite.	✓	✓
<code>bool isnan(double x)</code> Determine whether x is a NAN.	✓	✓
<code>double j0(double x)</code> Returns the value of the Bessel function of the first kind of order 0 for x .	✓	✓
<code>double j1(double x)</code> Returns the value of the Bessel function of the first kind of order 1 for x .	✓	✓
<code>double jn(int n, double x)</code> Returns the value of the Bessel function of the first kind of order n for x .	✓	✓

continues on next page

Table 17.2 – continued from previous page

<code>double ldexp(double x, int exp)</code> Returns the natural logarithm of the absolute value of the gamma function of x .	✓	✓
<code>double lgamma(double x)</code> Returns the natural logarithm of the absolute value of the gamma function of x .	✓	
<code>long int lrint(double x)</code> Round x to nearest integer value.	✓	✓
<code>long long int llrint(double x)</code> Round x to nearest integer value.	✓	✓
<code>long int lround(double x)</code> Round to nearest integer value.	✓	✓
<code>long long int llround(double x)</code> Round to nearest integer value.	✓	✓
<code>double log10(double x)</code> Returns the base 10 logarithm of x .	✓	✓
<code>double log1p(double x)</code> Returns the natural logarithm of $x + 1$.	✓	✓
<code>double log2(double x)</code> Returns the base 2 logarithm of x .	✓	✓
<code>double log(double x)</code> Returns the natural logarithm of x .	✓	✓

continues on next page

Table 17.2 – continued from previous page

<code>double logb(double x)</code> Returns the floating point representation of the exponent of x .	✓	✓
<code>double nan(const char* tagp)</code> Returns “Not a Number” value.		✓
<code>double nearbyint(double x)</code> Round x to the nearest integer.	✓	✓
<code>double nextafter(double x, double y)</code> Returns next representable double-precision floating-point value after argument.	✓	✓
<code>double norm3d(double x, double y, double z)</code> Returns the square root of the sum of squares of x , y and z .	✓	✓
<code>double norm4d(double x, double y, double z, double w)</code> Returns the square root of the sum of squares of x , y , z and w .	✓	✓
<code>double normcdf(double y)</code> Returns the standard normal cumulative distribution function.	✓	✓
<code>double normcdfinv(double y)</code> Returns the inverse of the standard normal cumulative distribution function.	✓	✓

continues on next page

Table 17.2 – continued from previous page

<code>double norm(int dim, const double *a)</code> Returns the square root of the sum of squares of any number of coordinates.	✓	✓
<code>double pow(double x, double y)</code> Returns x^y .	✓	✓
<code>double powi(double base, int iexp)</code> Returns the value of first argument to the power of second argument.	✓	✓
<code>double remainder(double x, double y)</code> Returns double-precision floating-point remainder.	✓	✓
<code>double remquo(double x, double y, int* quo)</code> Returns double-precision floating-point remainder and part of quotient.	✓	
<code>double round(double x)</code> Round to nearest integer value in floating-point.	✓	✓
<code>double rcbrt(double x)</code> Returns the reciprocal cube root function.	✓	✓
<code>double rhypot(double x, double y)</code> Returns one over the square root of the sum of squares of two arguments.	✓	✓

continues on next page

Table 17.2 – continued from previous page

<code>double rint(double x)</code> Round input to nearest integer value in floating-point.	✓	✓
<code>double rnorm3d(double x, double y, double z)</code> Returns one over the square root of the sum of squares of three coordinates of the argument.	✓	✓
<code>double rnorm4d(double x, double y, double z, double w)</code> Returns one over the square root of the sum of squares of four coordinates of the argument.	✓	✓
<code>double rnorm(int dim, const double *a)</code> Returns the reciprocal of square root of the sum of squares of any number of coordinates.	✓	✓
<code>double scalbln(double x, long int n)</code> Scale x by 2^n .	✓	✓
<code>double scalbn(double x, int n)</code> Scale x by 2^n .	✓	✓
<code>bool signbit(double x)</code> Return the sign bit of x .	✓	✓
<code>double sin(double x)</code> Returns the sine of x .	✓	✓
<code>double sinh(double x)</code> Returns the hyperbolic sine of x .	✓	✓

continues on next page

Table 17.2 – continued from previous page

<code>double sinpi(double x)</code> Returns the hyperbolic sine of $\pi \cdot x$.	✓	✓
<code>void sincos(double x, double *sptr, double *cptr)</code> Returns the sine and cosine of x .	✓	✓
<code>void sincospi(double x, double *sptr, double *cptr)</code> Returns the sine and cosine of $\pi \cdot x$.	✓	✓
<code>double sqrt(double x)</code> Returns the square root of x .	✓	✓
<code>double rsqrt(double x)</code> Returns the reciprocal of the square root of x .		✓
<code>double tan(double x)</code> Returns the tangent of x .	✓	✓
<code>double tanh(double x)</code> Returns the hyperbolic tangent of x .	✓	✓
<code>double tgamma(double x)</code> Returns the gamma function of x .	✓	✓
<code>double trunc(double x)</code> Truncate x to the integral part.	✓	✓
<code>double y0(double x)</code> Returns the value of the Bessel function of the second kind of order 0 for x .	✓	✓

continues on next page

Table 17.2 – continued from previous page

<code>double y1(double x)</code> Returns the value of the Bessel function of the second kind of order 1 for x .	✓	✓
<code>double yn(int n, double x)</code> Returns the value of the Bessel function of the second kind of order n for x .	✓	✓

17.9.3 Integer intrinsics

Following is the list of supported integer intrinsics. Note that intrinsics are supported on device only.

Table 17.3: Integer intrinsics mathematical functions

Function
<pre>unsigned int __brev(unsigned int x)</pre> Reverse the bit order of a 32 bit unsigned integer.
<pre>unsigned long long int __brevll(unsigned long long int x)</pre> Reverse the bit order of a 64 bit unsigned integer.
<pre>unsigned int __byte_perm(unsigned int x, unsigned int y, unsigned int z)</pre> Return selected bytes from two 32-bit unsigned integers.
<pre>unsigned int __clz(int x)</pre> Return the number of consecutive high-order zero bits in 32 bit integer.
<pre>unsigned int __clzll(long long int x)</pre> Return the number of consecutive high-order zero bits in 64 bit integer.
<pre>unsigned int __ffs(int x)</pre> Find the position of least significant bit set to 1 in a 32 bit integer.
<pre>unsigned int __ffsll(long long int x)</pre> Find the position of least significant bit set to 1 in a 64 bit signed integer.
<pre>unsigned int __fns32(unsigned long long mask, unsigned int base, int offset)</pre> Find the position of the n-th set to 1 bit in a 32-bit integer.
<pre>unsigned int __fns64(unsigned long long int mask, unsigned int base, int offset)</pre> Find the position of the n-th set to 1 bit in a 64-bit integer.
<pre>unsigned int __funnelshift_l(unsigned int lo, unsigned int hi, unsigned int shift)</pre> Concatenate <i>hi</i> and <i>lo</i>, shift left by shift & 31 bits, return the most significant 32 bits.
<pre>unsigned int __funnelshift_lc(unsigned int lo, unsigned int hi, unsigned int shift)</pre> Concatenate <i>hi</i> and <i>lo</i>, shift left by min(shift, 32) bits, return the most significant 32 bits.
<pre>unsigned int __funnelshift_r(unsigned int lo, unsigned int hi, unsigned int shift)</pre> Concatenate <i>hi</i> and <i>lo</i>, shift right by shift & 31 bits, return the least significant 32 bits.

The HIP-Clang implementation of `__ffs()` and `__ffsll()` contains code to add a constant +1 to produce the `ffs` result format. For the cases where this overhead is not acceptable and programmer is willing to specialize for the platform, HIP-Clang provides `__lastbit_u32_u32(unsigned int input)` and `__lastbit_u32_u64(unsigned long long int input)`. The index returned by `__lastbit_` instructions starts at -1, while for `ffs` the index starts at 0.

17.9.4 Floating-point Intrinsics

Following is the list of supported floating-point intrinsics. Note that intrinsics are supported on device only.

Note: Only the nearest even rounding mode supported on AMD GPUs by defaults. The `_rz`, `_ru` and `_rd` suffixed intrinsic functions are existing in HIP AMD backend, if the `OCML_BASIC_ROUNDED_OPERATIONS` macro is defined.

Table 17.4: Single precision intrinsics mathematical functions

Function
<code>float __cosf(float x)</code> Returns the fast approximate cosine of x.
<code>float __exp10f(float x)</code> Returns the fast approximate for 10^x.
<code>float __expf(float x)</code> Returns the fast approximate for e^x.
<code>float __fadd_rn(float x, float y)</code> Add two floating-point values in round-to-nearest-even mode.
<code>float __fdiv_rn(float x, float y)</code> Divide two floating point values in round-to-nearest-even mode.
<code>float __fmaf_rn(float x, float y, float z)</code> Returns $x \times y + z$ as a single operation in round-to-nearest-even mode.
<code>float __fmul_rn(float x, float y)</code> Multiply two floating-point values in round-to-nearest-even mode.
<code>float __frcp_rn(float x, float y)</code> Returns $1 / x$ in round-to-nearest-even mode.
<code>float __frsqrt_rn(float x)</code> Returns $1 / \sqrt{x}$ in round-to-nearest-even mode.
<code>float __fsqrt_rn(float x)</code> Returns $\sqrt{x}$ in round-to-nearest-even mode.
<code>float __fsub_rn(float x, float y)</code> Subtract two floating-point values in round-to-nearest-even mode.
17.9. Math functions <code>float __log10f(float x)</code> Returns the fast approximate for base 10 logarithm of x.

Table 17.5: Double precision intrinsics mathematical functions

Function
<code>double __dadd_rn(double x, double y)</code> Add two floating-point values in round-to-nearest-even mode.
<code>double __ddiv_rn(double x, double y)</code> Divide two floating-point values in round-to-nearest-even mode.
<code>double __dmul_rn(double x, double y)</code> Multiply two floating-point values in round-to-nearest-even mode.
<code>double __drcp_rn(double x, double y)</code> Returns $1 / x$ in round-to-nearest-even mode.
<code>double __dsqrt_rn(double x)</code> Returns x in round-to-nearest-even mode.
<code>double __dsub_rn(double x, double y)</code> Subtract two floating-point values in round-to-nearest-even mode.
<code>double __fma_rn(double x, double y, double z)</code> Returns $x \times y + z$ as a single operation in round-to-nearest-even mode.

17.10 Texture functions

The supported texture functions are listed in `texture_fetch_functions.h` and `texture_indirect_functions.h` header files in the [HIP-AMD backend repository](#).

Texture functions are not supported on some devices. To determine if texture functions are supported on your device, use Macro `__HIP_NO_IMAGE_SUPPORT == 1`. You can query the attribute `hipDeviceAttributeImageSupport` to check if texture functions are supported in the host runtime code.

17.11 Surface functions

Surface functions are not supported.

17.12 Timer functions

To read a high-resolution timer from the device, HIP provides the following built-in functions:

- Returning the incremental counter value for every clock cycle on a device:

```
clock_t clock()
long long int clock64()
```

The difference between the values that are returned represents the cycles used.

- Returning the wall clock count at a constant frequency on the device:

```
long long int wall_clock64()
```

This can be queried using the HIP API with the `hipDeviceAttributeWallClockRate` attribute of the device in HIP application code. For example:

```
int wallClkRate = 0; //in kilohertz
HIPCHECK(hipDeviceGetAttribute(&wallClkRate, hipDeviceAttributeWallClockRate,
↪ deviceId));
```

Where `hipDeviceAttributeWallClockRate` is a device attribute. Note that wall clock frequency is a per-device attribute.

17.13 Atomic functions

Atomic functions are run as read-modify-write (RMW) operations that reside in global or shared memory. No other device or thread can observe or modify the memory location during an atomic operation. If multiple instructions from different devices or threads target the same memory location, the instructions are serialized in an undefined order.

To support system scope atomic operations, you can use the HIP APIs that contain the `_system` suffix. For example:

- `atomicAnd`: This function is atomic and coherent within the GPU device running the function
- `atomicAnd_system`: This function extends the atomic operation from the GPU device to other CPUs and GPU devices in the system.

HIP supports the following atomic operations.

Table 17.6: Atomic operations

Function
<code>int atomicAdd(int* address, int val)</code>
<code>int atomicAdd_system(int* address, int val)</code>
<code>unsigned int atomicAdd(unsigned int* address, unsigned int val)</code>
<code>unsigned int atomicAdd_system(unsigned int* address, unsigned int val)</code>
<code>unsigned long long atomicAdd(unsigned long long* address, unsigned long long val)</code>
<code>unsigned long long atomicAdd_system(unsigned long long* address, unsigned long long val)</code>

Table 17.6 – continued from previous page

```

float atomicAdd(float* address, float val)
float atomicAdd_system(float* address, float val)
double atomicAdd(double* address, double val)
double atomicAdd_system(double* address, double val)
float unsafeAtomicAdd(float* address, float val)
float safeAtomicAdd(float* address, float val)
double unsafeAtomicAdd(double* address, double val)
double safeAtomicAdd(double* address, double val)
int atomicSub(int* address, int val)
int atomicSub_system(int* address, int val)
unsigned int atomicSub(unsigned int* address, unsigned int val)
unsigned int atomicSub_system(unsigned int* address, unsigned int val)
int atomicExch(int* address, int val)
int atomicExch_system(int* address, int val)
unsigned int atomicExch(unsigned int* address, unsigned int val)
unsigned int atomicExch_system(unsigned int* address, unsigned int val)
unsigned long long atomicExch(unsigned long long int* address, unsigned long long int val)
unsigned long long atomicExch_system(unsigned long long* address, unsigned long long val)
unsigned long long atomicExch_system(unsigned long long* address, unsigned long long val)
float atomicExch(float* address, float val)
int atomicMin(int* address, int val)
int atomicMin_system(int* address, int val)
unsigned int atomicMin(unsigned int* address, unsigned int val)
unsigned int atomicMin_system(unsigned int* address, unsigned int val)
unsigned long long atomicMin(unsigned long long* address, unsigned long long val)
int atomicMax(int* address, int val)
int atomicMax_system(int* address, int val)
unsigned int atomicMax(unsigned int* address, unsigned int val)
unsigned int atomicMax_system(unsigned int* address, unsigned int val)
unsigned long long atomicMax(unsigned long long* address, unsigned long long val)
unsigned int atomicInc(unsigned int* address)
unsigned int atomicDec(unsigned int* address)
int atomicCAS(int* address, int compare, int val)
int atomicCAS_system(int* address, int compare, int val)
unsigned int atomicCAS(unsigned int* address, unsigned int compare, unsigned int val)
unsigned int atomicCAS_system(unsigned int* address, unsigned int compare, unsigned int val)
unsigned long long atomicCAS(unsigned long long* address, unsigned long long compare, unsigned long long val)
unsigned long long atomicCAS_system(unsigned long long* address, unsigned long long compare, unsigned long long val)
int atomicAnd(int* address, int val)
int atomicAnd_system(int* address, int val)
unsigned int atomicAnd(unsigned int* address, unsigned int val)
unsigned int atomicAnd_system(unsigned int* address, unsigned int val)
unsigned long long atomicAnd(unsigned long long* address, unsigned long long val)
unsigned long long atomicAnd_system(unsigned long long* address, unsigned long long val)
int atomicOr(int* address, int val)
int atomicOr_system(int* address, int val)
unsigned int atomicOr(unsigned int* address, unsigned int val)
unsigned int atomicOr_system(unsigned int* address, unsigned int val)
unsigned int atomicOr_system(unsigned int* address, unsigned int val)
unsigned long long atomicOr(unsigned long long int* address, unsigned long long val)
unsigned long long atomicOr_system(unsigned long long* address, unsigned long long val)
int atomicXor(int* address, int val)

```

Table 17.6 – continued from previous page

```

int atomicXor_system(int* address, int val)
unsigned int atomicXor(unsigned int* address, unsigned int val)
unsigned int atomicXor_system(unsigned int* address, unsigned int val)
unsigned long long atomicXor(unsigned long long* address, unsigned long long val)
unsigned long long atomicXor_system(unsigned long long* address, unsigned long long val)

```

17.13.1 Unsafe floating-point atomic RMW operations

Some HIP devices support fast atomic RMW operations on floating-point values. For example, `atomicAdd` on single- or double-precision floating-point values may generate a hardware RMW instruction that is faster than emulating the atomic operation using an atomic compare-and-swap (CAS) loop.

On some devices, fast atomic RMW instructions can produce results that differ from the same functions implemented with atomic CAS loops. For example, some devices will use different rounding or denormal modes, and some devices produce incorrect answers if fast floating-point atomic RMW instructions target fine-grained memory allocations.

The HIP-Clang compiler offers a compile-time option, so you can choose fast—but potentially unsafe—atomic instructions for your code. On devices that support these instructions, you can include the `-munsafe-fp-atomics` option. This flag indicates to the compiler that all floating-point atomic function calls are allowed to use an unsafe version, if one exists. For example, on some devices, this flag indicates to the compiler that no floating-point `atomicAdd` function can target fine-grained memory.

If you want to avoid using unsafe use a floating-point atomic RMW operations, you can use the `-mno-unsafe-fp-atomics` option. Note that the compiler default is to not produce unsafe floating-point atomic RMW instructions, so the `-mno-unsafe-fp-atomics` option is not necessarily required. However, passing this option to the compiler is good practice.

When you pass `-munsafe-fp-atomics` or `-mno-unsafe-fp-atomics` to the compiler's command line, the option is applied globally for the entire compilation. Note that if some of the atomic RMW function calls cannot safely use the faster floating-point atomic RMW instructions, you must use `-mno-unsafe-fp-atomics` in order to ensure that your atomic RMW function calls produce correct results.

HIP has four extra functions that you can use to more precisely control which floating-point atomic RMW functions produce unsafe atomic RMW instructions:

- `float unsafeAtomicAdd(float* address, float val)`
- `double unsafeAtomicAdd(double* address, double val)` (Always produces fast atomic RMW instructions on devices that have them, even when `-mno-unsafe-fp-atomics` is used)
- `float safeAtomicAdd(float* address, float val)`
- `double safeAtomicAdd(double* address, double val)` (Always produces safe atomic RMW operations, even when `-munsafe-fp-atomics` is used)

17.14 Warp cross-lane functions

Threads in a warp are referred to as lanes and are numbered from 0 to `warpSize - 1`. Warp cross-lane functions operate across all lanes in a warp. The hardware guarantees that all warp lanes will execute in lockstep, so additional synchronization is unnecessary, and the instructions use no shared memory.

Note that NVIDIA and AMD devices have different warp sizes. You can use `warpSize` built-ins in your portable code to query the warp size.

Tip: Be sure to review HIP code generated from the CUDA path to ensure that it doesn't assume a `waveSize` of 32. "Wave-aware" code that assumes a `waveSize` of 32 can run on a wave-64 machine, but it only utilizes half of the machine's resources.

To get the default warp size of a GPU device, use `hipGetDeviceProperties` in your host functions.

```
cudaDeviceProp props;
cudaGetDeviceProperties(&props, deviceID);
int w = props.warpSize;
    // implement portable algorithm based on w (rather than assume 32 or 64)
```

Only use `warpSize` built-ins in device functions, and don't assume `warpSize` to be a compile-time constant.

Note that assembly kernels may be built for a warp size that is different from the default. All mask values either returned or accepted by these builtins are 64-bit unsigned integer values, even when compiled for a wave-32 device, where all the higher bits are unused. CUDA code ported to HIP requires changes to ensure that the correct type is used.

Note that the `__sync` variants are made available in ROCm 6.2, but disabled by default to help with the transition to 64-bit masks. They can be enabled by setting the preprocessor macro `HIP_ENABLE_WARP_SYNC_BUILTINS`. These builtins will be enabled unconditionally in the next ROCm release. Wherever possible, the implementation includes a static assert to check that the program source uses the correct type for the mask.

17.14.1 Warp vote and ballot functions

```
int __all(int predicate)
int __any(int predicate)
unsigned long long __ballot(int predicate)
unsigned long long __activemask()

int __all_sync(unsigned long long mask, int predicate)
int __any_sync(unsigned long long mask, int predicate)
unsigned long long __ballot_sync(unsigned long long mask, int predicate)
```

You can use `__any` and `__all` to get a summary view of the predicates evaluated by the participating lanes.

- `__any()`: Returns 1 if the predicate is non-zero for any participating lane, otherwise it returns 0.
- `__all()`: Returns 1 if the predicate is non-zero for all participating lanes, otherwise it returns 0.

To determine if the target platform supports the any/all instruction, you can use the `hasWarpVote` device property or the `HIP_ARCH_HAS_WARP_VOTE` compiler definition.

`__ballot` returns a bit mask containing the 1-bit predicate value from each lane. The *n*th bit of the result contains the 1 bit contributed by the *n*th warp lane.

`__activemask()` returns a bit mask of currently active warp lanes. The *n*th bit of the result is 1 if the *n*th warp lane is active.

Note that the `__ballot` and `__activemask` builtins in HIP have a 64-bit return value (unlike the 32-bit value returned by the CUDA builtins). Code ported from CUDA should be adapted to support the larger warp sizes that the HIP version requires.

Applications can test whether the target platform supports the `__ballot` or `__activemask` instructions using the `hasWarpBallot` device property in host code or the `HIP_ARCH_HAS_WARP_BALLOT` macro defined by the compiler for device code.

The `_sync` variants require a 64-bit unsigned integer mask argument that specifies the lanes in the warp that will participate in cross-lane communication with the calling lane. Each participating thread must have its own bit set in its mask argument, and all active threads specified in any mask argument must execute the same call with the same mask, otherwise the result is undefined.

17.14.2 Warp match functions

```

unsigned long long __match_any(T value)
unsigned long long __match_all(T value, int *pred)

unsigned long long __match_any_sync(unsigned long long mask, T value)
unsigned long long __match_all_sync(unsigned long long mask, T value, int *pred)

```

T can be a 32-bit integer type, 64-bit integer type or a single precision or double precision floating point type.

`__match_any` returns a bit mask containing a 1-bit for every participating lane if and only if that lane has the same value in `value` as the current lane, and a 0-bit for all other lanes.

`__match_all` returns a bit mask containing a 1-bit for every participating lane if and only if they all have the same value in `value` as the current lane, and a 0-bit for all other lanes. The predicate `pred` is set to true if and only if all participating threads have the same value in `value`.

The `_sync` variants require a 64-bit unsigned integer mask argument that specifies the lanes in the warp that will participate in cross-lane communication with the calling lane. Each participating thread must have its own bit set in its mask argument, and all active threads specified in any mask argument must execute the same call with the same mask, otherwise the result is undefined.

17.14.3 Warp shuffle functions

The default width is `warpSize` (see *Warp cross-lane functions*). Half-float shuffles are not supported.

```

T __shfl      (T var, int srcLane, int width=warpSize);
T __shfl_up   (T var, unsigned int delta, int width=warpSize);
T __shfl_down (T var, unsigned int delta, int width=warpSize);
T __shfl_xor  (T var, int laneMask, int width=warpSize);

T __shfl_sync    (unsigned long long mask, T var, int srcLane, int width=warpSize);
T __shfl_up_sync (unsigned long long mask, T var, unsigned int delta, int width=warpSize);
T __shfl_down_sync (unsigned long long mask, T var, unsigned int delta, int width=warpSize);
T __shfl_xor_sync (unsigned long long mask, T var, int laneMask, int width=warpSize);

```

T can be a 32-bit integer type, 64-bit integer type or a single precision or double precision floating point type.

The `_sync` variants require a 64-bit unsigned integer mask argument that specifies the lanes in the warp that will participate in cross-lane communication with the calling lane. Each participating thread must have its own bit set in its mask argument, and all active threads specified in any mask argument must execute the same call with the same mask, otherwise the result is undefined.

17.15 Cooperative groups functions

You can use cooperative groups to synchronize groups of threads. Cooperative groups also provide a way of communicating between groups of threads at a granularity that is different from the block.

HIP supports the following kernel language cooperative groups types and functions:

Table 17.7: Cooperative groups functions

Function	Supported in HIP	Supported in CUDA
<code>void thread_group.sync();</code>	✓	✓
<code>unsigned thread_group.size();</code>	✓	✓
<code>unsigned thread_group.thread_rank();</code>	✓	✓
<code>bool thread_group.is_valid();</code>	✓	✓
<code>grid_group this_grid();</code>	✓	✓
<code>void grid_group.sync();</code>	✓	✓
<code>unsigned grid_group.size();</code>	✓	✓
<code>unsigned grid_group.thread_rank();</code>	✓	✓
<code>bool grid_group.is_valid();</code>	✓	✓
<code>multi_grid_group this_multi_grid();</code>	✓	✓
<code>void multi_grid_group.sync();</code>	✓	✓
<code>unsigned multi_grid_group.size();</code>	✓	✓
<code>unsigned multi_grid_group.thread_rank();</code>	✓	✓
<code>bool multi_grid_group.is_valid();</code>	✓	✓
<code>unsigned multi_grid_group.num_grids();</code>	✓	✓
<code>unsigned multi_grid_group.grid_rank();</code>	✓	✓
<code>thread_block this_thread_block();</code>	✓	✓
<code>multi_grid_group this_multi_grid();</code>	✓	✓
<code>void multi_grid_group.sync();</code>	✓	✓
<code>void thread_block.sync();</code>	✓	✓
<code>unsigned thread_block.size();</code>	✓	✓
<code>unsigned thread_block.thread_rank();</code>	✓	✓
<code>bool thread_block.is_valid();</code>	✓	✓
<code>dim3 thread_block.group_index();</code>	✓	✓
<code>dim3 thread_block.thread_index();</code>	✓	✓

17.16 Warp matrix functions

Warp matrix functions allow a warp to cooperatively operate on small matrices that have elements spread over lanes in an unspecified manner.

HIP does not support kernel language warp matrix types or functions.

Table 17.8: Warp matrix functions

Function	Supported in HIP	Supported in CUDA
<code>void load_matrix_sync(fragment<...> &a, const T* mptr, unsigned lda)</code>		✓
<code>void load_matrix_sync(fragment<...> &a, const T* mptr, unsigned lda, layout_t layout)</code>		✓
<code>void store_matrix_sync(T* mptr, fragment<...> &a, unsigned lda, layout_t layout)</code>		✓
<code>void fill_fragment(fragment<...> &a, const T &value)</code>		✓
<code>void mma_sync(fragment<...> &d, const fragment<...> &a, const fragment<...> &b, const fragment<...> &c, bool sat)</code>		✓

17.17 Independent thread scheduling

Certain architectures that support CUDA allow threads to progress independently of each other. This independent thread scheduling makes intra-warp synchronization possible.

HIP does not support this type of scheduling.

17.18 Profiler Counter Function

The CUDA `__prof_trigger()` instruction is not supported.

17.19 Assert

The assert function is supported in HIP. Assert function is used for debugging purpose, when the input expression equals to zero, the execution will be stopped.

```
void assert(int input)
```

There are two kinds of implementations for assert functions depending on the use sceneries, - One is for the host version of assert, which is defined in `assert.h`, - Another is the device version of assert, which is implemented in `hip/hip_runtime.h`. Users need to include `assert.h` to use `assert`. For assert to work in both device and host functions, users need to include `"hip/hip_runtime.h"`.

HIP provides the function `abort()` which can be used to terminate the application when terminal failures are detected. It is implemented using the `__builtin_trap()` function.

This function produces a similar effect of using `asm("trap")` in the CUDA code.

Note: In HIP, the function terminates the entire application, while in CUDA, `asm("trap")` only terminates the dispatch and the application continues to run.

17.20 printf

printf function is supported in HIP. The following is a simple example to print information in the kernel.

```
#include <hip/hip_runtime.h>

__global__ void run_printf() { printf("Hello World\n"); }

int main() {
    run_printf<<<dim3(1), dim3(1), 0, 0>>>();
}
```

17.21 Device-Side Dynamic Global Memory Allocation

Device-side dynamic global memory allocation is under development. HIP now includes a preliminary implementation of malloc and free that can be called from device functions.

17.22 __launch_bounds__

GPU multiprocessors have a fixed pool of resources (primarily registers and shared memory) which are shared by the actively running warps. Using more resources can increase IPC of the kernel but reduces the resources available for other warps and limits the number of warps that can be simultaneously running. Thus GPUs have a complex relationship between resource usage and performance.

`__launch_bounds__` allows the application to provide usage hints that influence the resources (primarily registers) used by the generated code. It is a function attribute that must be attached to a `__global__` function:

```
__global__ void __launch_bounds__(MAX_THREADS_PER_BLOCK, MIN_WARPS_PER_EXECUTION_UNIT)
MyKernel(hipGridLaunch lp, ...)
...
```

`__launch_bounds__` supports two parameters: - `MAX_THREADS_PER_BLOCK` - The programmers guarantees that kernel will be launched with threads less than `MAX_THREADS_PER_BLOCK`. (On NVCC this maps to the `.maxntid` PTX directive). If no `launch_bounds` is specified, `MAX_THREADS_PER_BLOCK` is the maximum block size supported by the device (typically 1024 or larger). Specifying `MAX_THREADS_PER_BLOCK` less than the maximum effectively allows the compiler to use more resources than a default unconstrained compilation that supports all possible block sizes at launch time. The threads-per-block is the product of `(blockDim.x * blockDim.y * blockDim.z)`. - `MIN_WARPS_PER_EXECUTION_UNIT` - directs the compiler to minimize resource usage so that the requested number of warps can be simultaneously active on a multi-processor. Since active warps compete for the same fixed pool of resources, the compiler must reduce resources required by each warp (primarily registers). `MIN_WARPS_PER_EXECUTION_UNIT` is optional and defaults to 1 if not specified. Specifying a `MIN_WARPS_PER_EXECUTION_UNIT` greater than the default 1 effectively constrains the compiler's resource usage.

When launch kernel with HIP APIs, for example, `hipModuleLaunchKernel()`, HIP will do validation to make sure input kernel dimension size is not larger than specified `launch_bounds`. In case exceeded, HIP would return launch failure, if `AMD_LOG_LEVEL` is set with proper value (for details, please refer to `docs/markdown/hip_logging.md`), detail information will be shown in the error log message, including launch parameters of kernel dim size, launch bounds, and the name of the faulting kernel. It's helpful to figure out which is the faulting kernel, besides, the kernel dim size and launch bounds values will also assist in debugging such failures.

17.22.1 Compiler Impact

The compiler uses these parameters as follows: - The compiler uses the hints only to manage register usage, and does not automatically reduce shared memory or other resources. - Compilation fails if compiler cannot generate a kernel which meets the requirements of the specified launch bounds. - From `MAX_THREADS_PER_BLOCK`, the compiler derives the maximum number of warps/block that can be used at launch time. Values of `MAX_THREADS_PER_BLOCK` less than the default allows the compiler to use a larger pool of registers : each warp uses registers, and this hint constrains the launch to a warps/block size which is less than maximum. - From `MIN_WARPS_PER_EXECUTION_UNIT`, the compiler derives a maximum number of registers that can be used by the kernel (to meet the required #simultaneous active blocks). If `MIN_WARPS_PER_EXECUTION_UNIT` is 1, then the kernel can use all registers supported by the multiprocessor. - The compiler ensures that the registers used in the kernel is less than both allowed maximums, typically by spilling registers (to shared or global memory), or by using more instructions. - The compiler may use heuristics to increase register usage, or may simply be able to avoid spilling. The `MAX_THREADS_PER_BLOCK` is particularly useful in this cases, since it allows the compiler to use more registers and avoid situations where the compiler constrains the register usage (potentially spilling) to meet the requirements of a large block size that is never used at launch time.

17.22.2 CU and EU Definitions

A compute unit (CU) is responsible for executing the waves of a work-group. It is composed of one or more execution units (EU) which are responsible for executing waves. An EU can have enough resources to maintain the state of more than one executing wave. This allows an EU to hide latency by switching between waves in a similar way to symmetric multithreading on a CPU. In order to allow the state for multiple waves to fit on an EU, the resources used by a single wave have to be limited. Limiting such resources can allow greater latency hiding, but can result in having to spill some register state to memory. This attribute allows an advanced developer to tune the number of waves that are capable of fitting within the resources of an EU. It can be used to ensure at least a certain number will fit to help hide latency, and can also be used to ensure no more than a certain number will fit to limit cache thrashing.

17.22.3 Porting from CUDA `__launch_bounds`

CUDA defines a `__launch_bounds` which is also designed to control occupancy:

```
__launch_bounds(MAX_THREADS_PER_BLOCK, MIN_BLOCKS_PER_MULTIPROCESSOR)
```

- The second parameter `__launch_bounds` parameters must be converted to the format used `__hip_launch_bounds`, which uses warps and execution-units rather than blocks and multi-processors (this conversion is performed automatically by HIPIFY tools).

```
MIN_WARPS_PER_EXECUTION_UNIT = (MIN_BLOCKS_PER_MULTIPROCESSOR * MAX_THREADS_PER_BLOCK) / 32
```

The key differences in the interface are: - Warps (rather than blocks): The developer is trying to tell the compiler to control resource utilization to guarantee some amount of active Warps/EU for latency hiding. Specifying active warps in terms of blocks appears to hide the micro-architectural details of the warp size, but makes the interface more confusing since the developer ultimately needs to compute the number of warps to obtain the desired level of control. - Execution Units (rather than multiprocessor): The use of execution units rather than multiprocessors provides support for architectures with multiple execution units/multi-processor. For example, the AMD GCN architecture has 4 execution units per multiprocessor. The `hipDeviceProps` has a field `executionUnitsPerMultiprocessor`. Platform-specific coding techniques such as `#ifdef` can be used to specify different launch_bounds for NVCC and HIP-Clang platforms, if desired.

17.22.4 maxregcount

Unlike NVCC, HIP-Clang does not support the `--maxregcount` option. Instead, users are encouraged to use the `hip_launch_bounds` directive since the parameters are more intuitive and portable than micro-architecture details like registers, and also the directive allows per-kernel control rather than an entire file. `hip_launch_bounds` works on both HIP-Clang and NVCC targets.

17.23 Asynchronous Functions

17.23.1 Memory stream

```
typedef void (*hipStreamCallback_t)(hipStream_t stream, hipError_t status, void *userData)
```

Stream CallBack struct

```
hipError_t hipStreamCreate(hipStream_t *stream)
```

Create an asynchronous stream.

Create a new asynchronous stream. `stream` returns an opaque handle that can be used to reference the newly created stream in subsequent `hipStream*` commands. The stream is allocated on the heap and will remain allocated even if the handle goes out-of-scope. To release the memory used by the stream, application must call `hipStreamDestroy`.

See also:

[hipStreamCreateWithFlags](#), [hipStreamCreateWithPriority](#), [hipStreamSynchronize](#), [hipStreamWaitEvent](#), [hipStreamDestroy](#)

Parameters

stream – [inout] Valid pointer to `hipStream_t`. This function writes the memory with the newly created stream.

Returns

`hipSuccess`, `hipErrorInvalidValue`

Returns

`hipSuccess`, `hipErrorInvalidValue`

```
hipError_t hipStreamCreateWithFlags(hipStream_t *stream, unsigned int flags)
```

Create an asynchronous stream.

Create a new asynchronous stream. `stream` returns an opaque handle that can be used to reference the newly created stream in subsequent `hipStream*` commands. The stream is allocated on the heap and will remain allocated even if the handle goes out-of-scope. To release the memory used by the stream, application must call `hipStreamDestroy`. `Flags` controls behavior of the stream. See `hipStreamDefault`, `hipStreamNonBlocking`.

See also:

[hipStreamCreate](#), [hipStreamCreateWithPriority](#), [hipStreamSynchronize](#), [hipStreamWaitEvent](#), [hipStreamDestroy](#)

Parameters

- **stream** – [inout] Pointer to new stream
- **flags** – [in] to control stream creation.

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipStreamCreateWithPriority**(hipStream_t *stream, unsigned int flags, int priority)

Create an asynchronous stream with the specified priority.

Create a new asynchronous stream with the specified priority. **stream** returns an opaque handle that can be used to reference the newly created stream in subsequent hipStream* commands. The stream is allocated on the heap and will remain allocated even if the handle goes out-of-scope. To release the memory used by the stream, application must call hipStreamDestroy. Flags controls behavior of the stream. See hipStreamDefault, hipStreamNonBlocking.

See also:

hipStreamCreate, hipStreamSynchronize, hipStreamWaitEvent, hipStreamDestroy

Parameters

- **stream** – [inout] Pointer to new stream
- **flags** – [in] to control stream creation.
- **priority** – [in] of the stream. Lower numbers represent higher priorities.

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipDeviceGetStreamPriorityRange**(int *leastPriority, int *greatestPriority)

Returns numerical values that correspond to the least and greatest stream priority.

Returns in *leastPriority and *greatestPriority the numerical values that correspond to the least and greatest stream priority respectively. Stream priorities follow a convention where lower numbers imply greater priorities. The range of meaningful stream priorities is given by [*greatestPriority, *leastPriority]. If the user attempts to create a stream with a priority value that is outside the meaningful range as specified by this API, the priority is automatically clamped to within the valid range.

Parameters

- **leastPriority** – [inout] pointer in which value corresponding to least priority is returned.
- **greatestPriority** – [inout] pointer in which value corresponding to greatest priority is returned.

Returns

hipSuccess

hipError_t **hipStreamDestroy**(hipStream_t stream)

Destroys the specified stream.

Destroys the specified stream.

If commands are still executing on the specified stream, some may complete execution before the queue is deleted. The queue may be destroyed while some commands are still in flight, or may wait for all commands queued to the stream before destroying it.

See also:

hipStreamCreate, *hipStreamCreateWithFlags*, *hipStreamCreateWithPriority*, *hipStreamQuery*, *hipStreamWaitEvent*, *hipStreamSynchronize*

Parameters

stream – [in] stream identifier.

Returns

hipSuccess hipErrorInvalidHandle

hipError_t **hipStreamQuery**(hipStream_t stream)

Return hipSuccess if all of the operations in the specified **stream** have completed, or hipErrorNotReady if not.

This is thread-safe and returns a snapshot of the current state of the queue. However, if other host threads are sending work to the stream, the status may change immediately after the function is called. It is typically used for debug.

See also:

hipStreamCreate, *hipStreamCreateWithFlags*, *hipStreamCreateWithPriority*, *hipStreamWaitEvent*, *hipStreamSynchronize*, *hipStreamDestroy*

Parameters

stream – [in] stream to query

Returns

hipSuccess, hipErrorNotReady, hipErrorInvalidHandle

hipError_t **hipStreamSynchronize**(hipStream_t stream)

Wait for all commands in stream to complete.

This command is host-synchronous : the host will block until the specified stream is empty.

This command follows standard null-stream semantics. Specifically, specifying the null stream will cause the command to wait for other streams on the same device to complete all pending operations.

This command honors the hipDeviceLaunchBlocking flag, which controls whether the wait is active or blocking.

See also:

hipStreamCreate, *hipStreamCreateWithFlags*, *hipStreamCreateWithPriority*, *hipStreamWaitEvent*, *hipStreamDestroy*

Parameters

stream – [in] stream identifier.

Returns

hipSuccess, hipErrorInvalidHandle

hipError_t **hipStreamWaitEvent**(hipStream_t stream, hipEvent_t event, unsigned int flags)

Make the specified compute stream wait for an event.

This function inserts a wait operation into the specified stream. All future work submitted to `stream` will wait until `event` reports completion before beginning execution.

This function only waits for commands in the current stream to complete. Notably, this function does not implicitly wait for commands in the default stream to complete, even if the specified stream is created with `hipStreamNonBlocking = 0`.

See also:

hipStreamCreate, *hipStreamCreateWithFlags*, *hipStreamCreateWithPriority*, *hipStreamSynchronize*, *hipStreamDestroy*

Parameters

- **stream** – [in] stream to make wait.
- **event** – [in] event to wait on
- **flags** – [in] control operation [must be 0]

Returns

hipSuccess, hipErrorInvalidHandle

hipError_t **hipStreamGetFlags**(hipStream_t stream, unsigned int *flags)

Return flags associated with this stream.

Return flags associated with this stream in *flags.

See also:

hipStreamCreateWithFlags

Parameters

- **stream** – [in] stream to be queried
- **flags** – [inout] Pointer to an unsigned integer in which the stream's flags are returned

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidHandle

Returns

hipSuccess hipErrorInvalidValue hipErrorInvalidHandle

hipError_t **hipStreamGetPriority**(hipStream_t stream, int *priority)

Query the priority of a stream.

Query the priority of a stream. The priority is returned in in priority.

See also:

hipStreamCreateWithFlags

Parameters

- **stream** – [in] stream to be queried
- **priority** – [inout] Pointer to an unsigned integer in which the stream’s priority is returned

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidHandle

Returns

hipSuccess hipErrorInvalidValue hipErrorInvalidHandle

hipError_t **hipStreamGetDevice**(hipStream_t stream, hipDevice_t *device)

Get the device associated with the stream.

See also:

hipStreamCreate, hipStreamDestroy, hipDeviceGetStreamPriorityRange

Parameters

- **stream** – [in] stream to be queried
- **device** – [out] device associated with the stream

Returns

hipSuccess, hipErrorInvalidValue, hipErrorContextIsDestroyed, hipErrorInvalidHandle, hipErrorNotInitialized, hipErrorDeinitialized, hipErrorInvalidContext

hipError_t **hipExtStreamCreateWithCUMask**(hipStream_t *stream, uint32_t cuMaskSize, const uint32_t *cuMask)

Create an asynchronous stream with the specified CU mask.

Create a new asynchronous stream with the specified CU mask. **stream** returns an opaque handle that can be used to reference the newly created stream in subsequent hipStream* commands. The stream is allocated on the heap and will remain allocated even if the handle goes out-of-scope. To release the memory used by the stream, application must call hipStreamDestroy.

See also:

hipStreamCreate, hipStreamSynchronize, hipStreamWaitEvent, hipStreamDestroy

Parameters

- **stream** – [inout] Pointer to new stream
- **cuMaskSize** – [in] Size of CU mask bit array passed in.
- **cuMask** – [in] Bit-vector representing the CU mask. Each active bit represents using one CU. The first 32 bits represent the first 32 CUs, and so on. If its size is greater than physical

CU number (i.e., multiProcessorCount member of hipDeviceProp_t), the extra elements are ignored. It is user's responsibility to make sure the input is meaningful.

Returns

hipSuccess, hipErrorInvalidHandle, hipErrorInvalidValue

hipError_t **hipExtStreamGetCUMask**(hipStream_t stream, uint32_t cuMaskSize, uint32_t *cuMask)

Get CU mask associated with an asynchronous stream.

See also:

hipStreamCreate, hipStreamSynchronize, hipStreamWaitEvent, hipStreamDestroy

Parameters

- **stream** – [in] stream to be queried
- **cuMaskSize** – [in] number of the block of memories (uint32_t *) allocated by user
- **cuMask** – [out] Pointer to a pre-allocated block of memories (uint32_t *) in which the stream's CU mask is returned. The CU mask is returned in a chunk of 32 bits where each active bit represents one active CU

Returns

hipSuccess, hipErrorInvalidHandle, hipErrorInvalidValue

hipError_t **hipStreamAddCallback**(hipStream_t stream, *hipStreamCallback_t* callback, void *userData, unsigned int flags)

Adds a callback to be called on the host after all currently enqueued items in the stream have completed. For each hipStreamAddCallback call, a callback will be executed exactly once. The callback will block later work in the stream until it is finished.

See also:

hipStreamCreate, hipStreamCreateWithFlags, hipStreamQuery, hipStreamSynchronize, hipStreamWaitEvent, hipStreamDestroy, hipStreamCreateWithPriority

Parameters

- **stream** – [in] - Stream to add callback to
- **callback** – [in] - The function to call once preceding stream operations are complete
- **userData** – [in] - User specified data to be passed to the callback function
- **flags** – [in] - Reserved for future use, must be 0

Returns

hipSuccess, hipErrorInvalidHandle, hipErrorNotSupported

hipError_t **hipMallocAsync**(void **dev_ptr, size_t size, hipStream_t stream)

Allocates memory with stream ordered semantics.

Inserts a memory allocation operation into **stream**. A pointer to the allocated memory is returned immediately in *dptr. The allocation must not be accessed until the allocation operation completes. The allocation comes from the memory pool associated with the stream's device.

See also:

hipMallocFromPoolAsync, *hipFreeAsync*, *hipMemPoolTrimTo*, *hipMemPoolGetAttribute*, *hipDeviceSetMemPool*, *hipMemPoolSetAttribute*, *hipMemPoolSetAccess*, *hipMemPoolGetAccess*

Note: The default memory pool of a device contains device memory from that device.

Note: Basic stream ordering allows future work submitted into the same stream to use the allocation. Stream query, stream synchronize, and HIP events can be used to guarantee that the allocation operation completes before work submitted in a separate stream runs.

Note: During stream capture, this function results in the creation of an allocation node. In this case, the allocation is owned by the graph instead of the memory pool. The memory pool's properties are used to set the node's creation parameters.

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **dev_ptr** – [out] Returned device pointer of memory allocation
- **size** – [in] Number of bytes to allocate
- **stream** – [in] The stream establishing the stream ordering contract and the memory pool to allocate from

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotSupported, hipErrorOutOfMemory

hipError_t **hipFreeAsync**(void *dev_ptr, hipStream_t stream)

Frees memory with stream ordered semantics.

Inserts a free operation into *stream*. The allocation must not be used after stream execution reaches the free. After this API returns, accessing the memory from any subsequent work launched on the GPU or querying its pointer attributes results in undefined behavior.

See also:

hipMallocFromPoolAsync, *hipMallocAsync*, *hipMemPoolTrimTo*, *hipMemPoolGetAttribute*, *hipDeviceSetMemPool*, *hipMemPoolSetAttribute*, *hipMemPoolSetAccess*, *hipMemPoolGetAccess*

Note: During stream capture, this function results in the creation of a free node and must therefore be passed the address of a graph allocation.

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **dev_ptr** – [in] Pointer to device memory to free
- **stream** – [in] The stream, where the destrucion will occur according to the execution order

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotSupported

hipError_t **hipMemPoolTrimTo**(hipMemPool_t mem_pool, size_t min_bytes_to_hold)

Releases freed memory back to the OS.

Releases memory back to the OS until the pool contains fewer than `min_bytes_to_keep` reserved bytes, or there is no more memory that the allocator can safely release. The allocator cannot release OS allocations that back outstanding asynchronous allocations. The OS allocations may happen at different granularity from the user allocations.

See also:

hipMallocFromPoolAsync, hipMallocAsync, hipFreeAsync, hipMemPoolGetAttribute, hipDeviceSetMemPool, hipMemPoolSetAttribute, hipMemPoolSetAccess, hipMemPoolGetAccess

Note: : Allocations that have not been freed count as outstanding.

Note: : Allocations that have been asynchronously freed but whose completion has not been observed on the host (eg. by a synchronize) can count as outstanding.

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **mem_pool** – [in] The memory pool to trim allocations
- **min_bytes_to_hold** – [in] If the pool has less than `min_bytes_to_hold` reserved, then the TrimTo operation is a no-op. Otherwise the memory pool will contain at least `min_bytes_to_hold` bytes reserved after the operation.

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemPoolSetAttribute**(hipMemPool_t mem_pool, hipMemPoolAttr attr, void *value)

Sets attributes of a memory pool.

Supported attributes are:

- **hipMemPoolAttrReleaseThreshold:** (value type = cuuint64_t) Amount of reserved memory in bytes to hold onto before trying to release memory back to the OS. When more than the release threshold bytes of memory are held by the memory pool, the allocator will try to release memory back to the OS on the next call to stream, event or context synchronize. (default 0)
- **hipMemPoolReuseFollowEventDependencies:** (value type = int) Allow **hipMallocAsync** to use memory asynchronously freed in another stream as long as a stream ordering dependency of the allocating stream on the free action exists. HIP events and null stream interactions can create the required stream ordered dependencies. (default enabled)
- **hipMemPoolReuseAllowOpportunistic:** (value type = int) Allow reuse of already completed frees when there is no dependency between the free and allocation. (default enabled)
- **hipMemPoolReuseAllowInternalDependencies:** (value type = int) Allow **hipMallocAsync** to insert new stream dependencies in order to establish the stream ordering required to reuse a piece of memory released by **hipFreeAsync** (default enabled).

See also:

hipMallocFromPoolAsync, hipMallocAsync, hipFreeAsync, hipMemPoolGetAttribute, hipMemPoolTrimTo, hipDeviceSetMemPool, hipMemPoolSetAccess, hipMemPoolGetAccess

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **mem_pool** – [in] The memory pool to modify
- **attr** – [in] The attribute to modify
- **value** – [in] Pointer to the value to assign

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemPoolGetAttribute**(hipMemPool_t mem_pool, hipMemPoolAttr attr, void *value)

Gets attributes of a memory pool.

Supported attributes are:

- **hipMemPoolAttrReleaseThreshold:** (value type = cuuint64_t) Amount of reserved memory in bytes to hold onto before trying to release memory back to the OS. When more than the release threshold bytes of memory are held by the memory pool, the allocator will try to release memory back to the OS on the next call to stream, event or context synchronize. (default 0)
- **hipMemPoolReuseFollowEventDependencies:** (value type = int) Allow **hipMallocAsync** to use memory asynchronously freed in another stream as long as a stream ordering dependency of the allocating

stream on the free action exists. HIP events and null stream interactions can create the required stream ordered dependencies. (default enabled)

- **hipMemPoolReuseAllowOpportunistic**: (value type = int) Allow reuse of already completed frees when there is no dependency between the free and allocation. (default enabled)
- **hipMemPoolReuseAllowInternalDependencies**: (value type = int) Allow **hipMallocAsync** to insert new stream dependencies in order to establish the stream ordering required to reuse a piece of memory released by **hipFreeAsync** (default enabled).

See also:

hipMallocFromPoolAsync, hipMallocAsync, hipFreeAsync, hipMemPoolTrimTo, hipDeviceSetMemPool, hipMemPoolSetAttribute, hipMemPoolSetAccess, hipMemPoolGetAccess

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **mem_pool** – [in] The memory pool to get attributes of
- **attr** – [in] The attribute to get
- **value** – [in] Retrieved value

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemPoolSetAccess**(hipMemPool_t mem_pool, const hipMemAccessDesc *desc_list, size_t count)
Controls visibility of the specified pool between devices.

See also:

hipMallocFromPoolAsync, hipMallocAsync, hipFreeAsync, hipMemPoolGetAttribute, hipMemPoolTrimTo, hipDeviceSetMemPool, hipMemPoolSetAttribute, hipMemPoolGetAccess

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **mem_pool** – [in] Memory pool for access change
- **desc_list** – [in] Array of access descriptors. Each descriptor instructs the access to enable for a single gpu

- **count** – [in] Number of descriptors in the map array.

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemPoolGetAccess**(hipMemAccessFlags *flags, hipMemPool_t mem_pool, hipMemLocation *location)

Returns the accessibility of a pool from a device.

Returns the accessibility of the pool's memory from the specified location.

See also:

hipMallocFromPoolAsync, hipMallocAsync, hipFreeAsync, hipMemPoolGetAttribute, hipMemPoolTrimTo, hipDeviceSetMemPool, hipMemPoolSetAttribute, hipMemPoolSetAccess

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **flags** – [out] Accessibility of the memory pool from the specified location/device
- **mem_pool** – [in] Memory pool being queried
- **location** – [in] Location/device for memory pool access

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemPoolCreate**(hipMemPool_t *mem_pool, const hipMemPoolProps *pool_props)

Creates a memory pool.

Creates a HIP memory pool and returns the handle in `mem_pool`. The `pool_props` determines the properties of the pool such as the backing device and IPC capabilities.

By default, the memory pool will be accessible from the device it is allocated on.

See also:

hipMallocFromPoolAsync, hipMallocAsync, hipFreeAsync, hipMemPoolGetAttribute, hipMemPoolDestroy, hipMemPoolTrimTo, hipDeviceSetMemPool, hipMemPoolSetAttribute, hipMemPoolSetAccess, hipMemPoolGetAccess

Note: Specifying `hipMemHandleTypeNone` creates a memory pool that will not support IPC.

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **mem_pool** – [out] Contains created memory pool
- **pool_props** – [in] Memory pool properties

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotSupported

hipError_t **hipMemPoolDestroy**(hipMemPool_t mem_pool)

Destroys the specified memory pool.

If any pointers obtained from this pool haven't been freed or the pool has free operations that haven't completed when **hipMemPoolDestroy** is invoked, the function will return immediately and the resources associated with the pool will be released automatically once there are no more outstanding allocations.

Destroying the current mempool of a device sets the default mempool of that device as the current mempool for that device.

See also:

hipMallocFromPoolAsync, hipMallocAsync, hipFreeAsync, hipMemPoolGetAttribute, hipMemPoolCreate, hipMemPoolTrimTo, hipDeviceSetMemPool, hipMemPoolSetAttribute, hipMemPoolSetAccess, hipMemPoolGetAccess

Note: A device's default memory pool cannot be destroyed.

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

mem_pool – [in] Memory pool for destruction

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMallocFromPoolAsync**(void **dev_ptr, size_t size, hipMemPool_t mem_pool, hipStream_t stream)

Allocates memory from a specified pool with stream ordered semantics.

Inserts an allocation operation into **stream**. A pointer to the allocated memory is returned immediately in **dev_ptr**. The allocation must not be accessed until the allocation operation completes. The allocation comes from the specified memory pool.

Basic stream ordering allows future work submitted into the same stream to use the allocation. Stream query, stream synchronize, and HIP events can be used to guarantee that the allocation operation completes before work submitted in a separate stream runs.

See also:

hipMallocAsync, *hipFreeAsync*, *hipMemPoolGetAttribute*, *hipMemPoolCreate*, *hipMemPoolTrimTo*, *hipDeviceSetMemPool*, *hipMemPoolSetAttribute*, *hipMemPoolSetAccess*, *hipMemPoolGetAccess*,

Note: The specified memory pool may be from a device different than that of the specified `stream`.

Note: During stream capture, this function results in the creation of an allocation node. In this case, the allocation is owned by the graph instead of the memory pool. The memory pool's properties are used to set the node's creation parameters.

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **dev_ptr** – [out] Returned device pointer
- **size** – [in] Number of bytes to allocate
- **mem_pool** – [in] The pool to allocate from
- **stream** – [in] The stream establishing the stream ordering semantic

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotSupported, hipErrorOutOfMemory

hipError_t **hipMemPoolExportToShareableHandle**(void *shared_handle, hipMemPool_t mem_pool,
hipMemAllocationHandleType handle_type, unsigned int
flags)

Exports a memory pool to the requested handle type.

Given an IPC capable mempool, create an OS handle to share the pool with another process. A recipient process can convert the shareable handle into a mempool with `hipMemPoolImportFromShareableHandle`. Individual pointers can then be shared with the `hipMemPoolExportPointer` and `hipMemPoolImportPointer` APIs. The implementation of what the shareable handle is and how it can be transferred is defined by the requested handle type.

See also:

hipMemPoolImportFromShareableHandle

Note: : To create an IPC capable mempool, create a mempool with a `hipMemAllocationHandleType` other than `hipMemHandleTypeNone`.

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **shared_handle** – [out] Pointer to the location in which to store the requested handle
- **mem_pool** – [in] Pool to export
- **handle_type** – [in] The type of handle to create
- **flags** – [in] Must be 0

Returns

`hipSuccess`, `hipErrorInvalidValue`, `hipErrorOutOfMemory`

`hipError_t` **hipMemPoolImportFromShareableHandle**(`hipMemPool_t` *`mem_pool`, `void` *`shared_handle`, `hipMemAllocationHandleType` `handle_type`, unsigned int `flags`)

Imports a memory pool from a shared handle.

Specific allocations can be imported from the imported pool with `hipMemPoolImportPointer`.

See also:

[hipMemPoolExportToShareableHandle](#)

Note: Imported memory pools do not support creating new allocations. As such imported memory pools may not be used in `hipDeviceSetMemPool` or `hipMallocFromPoolAsync` calls.

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **mem_pool** – [out] Returned memory pool
- **shared_handle** – [in] OS handle of the pool to open
- **handle_type** – [in] The type of handle being imported
- **flags** – [in] Must be 0

Returns

hipSuccess, hipErrorInvalidValue, hipErrorOutOfMemory

hipError_t **hipMemPoolExportPointer**(hipMemPoolPtrExportData *export_data, void *dev_ptr)

Export data to share a memory pool allocation between processes.

Constructs `export_data` for sharing a specific allocation from an already shared memory pool. The recipient process can import the allocation with the `hipMemPoolImportPointer` api. The data is not a handle and may be shared through any IPC mechanism.

See also:

hipMemPoolImportPointer

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **export_data** – [out] Returned export data
- **dev_ptr** – [in] Pointer to memory being exported

Returns

hipSuccess, hipErrorInvalidValue, hipErrorOutOfMemory

hipError_t **hipMemPoolImportPointer**(void **dev_ptr, hipMemPool_t mem_pool, hipMemPoolPtrExportData *export_data)

Import a memory pool allocation from another process.

Returns in `dev_ptr` a pointer to the imported memory. The imported memory must not be accessed before the allocation operation completes in the exporting process. The imported memory must be freed from all importing processes before being freed in the exporting process. The pointer may be freed with `hipFree` or `hipFreeAsync`. If `hipFreeAsync` is used, the free must be completed on the importing process before the free operation on the exporting process.

See also:

hipMemPoolExportPointer

Note: The `hipFreeAsync` api may be used in the exporting process before the `hipFreeAsync` operation completes in its stream as long as the `hipFreeAsync` in the exporting process specifies a stream with a stream dependency on the importing process's `hipFreeAsync`.

Note: This API is implemented on Linux, under development on Windows.

Warning: : This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **dev_ptr** – [out] Pointer to imported memory
- **mem_pool** – [in] Memory pool from which to import a pointer
- **export_data** – [in] Data specifying the memory to import

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotInitialized, hipErrorOutOfMemory

17.23.2 Peer to peer

hipError_t **hipDeviceCanAccessPeer**(int *canAccessPeer, int deviceId, int peerDeviceId)

Determine if a device can access a peer's memory.

Returns "1" in `canAccessPeer` if the specified device is capable of directly accessing memory physically located on `peerDevice`, or "0" if not.

Returns "0" in `canAccessPeer` if `deviceId == peerDeviceId`, and both are valid devices : a device is not a peer of itself.

Parameters

- **canAccessPeer** – [out] Returns the peer access capability (0 or 1)
- **deviceId** – [in] - device from where memory may be accessed.
- **peerDeviceId** – [in] - device where memory is physically located

Returns

hipSuccess,

Returns

hipErrorInvalidDevice if `deviceId` or `peerDeviceId` are not valid devices

hipError_t **hipDeviceEnablePeerAccess**(int peerDeviceId, unsigned int flags)

Enable direct access from current device's virtual address space to memory allocations physically located on a peer device.

Memory which already allocated on peer device will be mapped into the address space of the current device. In addition, all future memory allocations on `peerDeviceId` will be mapped into the address space of the current device when the memory is allocated. The peer memory remains accessible from the current device until a call to `hipDeviceDisablePeerAccess` or `hipDeviceReset`.

Returns hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue,

Parameters

- **peerDeviceId** – [in] Peer device to enable direct access to from the current device
- **flags** – [in] Reserved for future use, must be zero

Returns

hipErrorPeerAccessAlreadyEnabled if peer access is already enabled for this device.

hipError_t **hipDeviceDisablePeerAccess**(int peerDeviceId)

Disable direct access from current device's virtual address space to memory allocations physically located on a peer device.

Returns hipErrorPeerAccessNotEnabled if direct access to memory on peerDevice has not yet been enabled from the current device.

Parameters

peerDeviceId – [in] Peer device to disable direct access to

Returns

hipSuccess, hipErrorPeerAccessNotEnabled

hipError_t **hipMemGetAddressRange**(hipDeviceptr_t *pbase, size_t *psize, hipDeviceptr_t dptr)

Get information on memory allocations.

See also:

hipCtxCreate, hipCtxDestroy, hipCtxGetFlags, hipCtxPopCurrent, hipCtxGetCurrent, hipCtxSetCurrent, hipCtxPushCurrent, hipCtxSetCacheConfig, hipCtxSynchronize, hipCtxGetDevice

Parameters

- **pbase** – [out] - Base pointer address
- **psize** – [out] - Size of allocation
- **dptr** – [in] Device Pointer

Returns

hipSuccess, hipErrorNotFound

hipError_t **hipMemcpyPeer**(void *dst, int dstDeviceId, const void *src, int srcDeviceId, size_t sizeBytes)

Copies memory from one device to memory on another device.

Parameters

- **dst** – [out] - Destination device pointer.
- **dstDeviceId** – [in] - Destination device
- **src** – [in] - Source device pointer
- **srcDeviceId** – [in] - Source device
- **sizeBytes** – [in] - Size of memory copy in bytes

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidDevice

hipError_t **hipMemcpyPeerAsync**(void *dst, int dstDeviceId, const void *src, int srcDevice, size_t sizeBytes, hipStream_t stream)

Copies memory from one device to memory on another device.

Parameters

- **dst** – [out] - Destination device pointer.
- **dstDeviceId** – [in] - Destination device
- **src** – [in] - Source device pointer
- **srcDevice** – [in] - Source device

- **sizeBytes** – [in] - Size of memory copy in bytes
- **stream** – [in] - Stream identifier

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidDevice

USE_PEER_NON_UNIFIED

17.23.3 Memory management

hipError_t **hipPointerSetAttribute**(const void *value, hipPointer_attribute attribute, hipDeviceptr_t ptr)

Sets information on the specified pointer.[BETA].

Warning: This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **value** – [in] Sets pointer attribute value
- **attribute** – [in] Attribute to set
- **ptr** – [in] Pointer to set attributes for

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipPointerGetAttributes**(hipPointerAttribute_t *attributes, const void *ptr)

Returns attributes for the specified pointer.

The output parameter ‘attributes’ has a member named ‘type’ that describes what memory the pointer is associated with, such as device memory, host memory, managed memory, and others. Otherwise, the API cannot handle the pointer and returns hipErrorInvalidValue.

See also:

[hipPointerGetAttribute](#)

Note: The unrecognized memory type is unsupported to keep the HIP functionality backward compatibility due to hipMemoryType enum values.

Note: The current behavior of this HIP API corresponds to the CUDA API before version 11.0.

Parameters

- **attributes** – [out] attributes for the specified pointer
- **ptr** – [in] pointer to get attributes for

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipPointerGetAttribute**(void *data, hipPointer_attribute attribute, hipDeviceptr_t ptr)

Returns information about the specified pointer.[BETA].

See also:

hipPointerGetAttributes

Warning: This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **data** – [inout] Returned pointer attribute value
- **attribute** – [in] Attribute to query for
- **ptr** – [in] Pointer to get attributes for

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipDrvPointerGetAttributes**(unsigned int numAttributes, hipPointer_attribute *attributes, void **data, hipDeviceptr_t ptr)

Returns information about the specified pointer.[BETA].

See also:

hipPointerGetAttribute

Warning: This API is marked as beta, meaning, while this is feature complete, it is still open to changes and may have outstanding issues.

Parameters

- **numAttributes** – [in] number of attributes to query for
- **attributes** – [in] attributes to query for
- **data** – [inout] a two-dimensional containing pointers to memory locations where the result of each attribute query will be written to
- **ptr** – [in] pointer to get attributes for

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipMalloc**(void **ptr, size_t size)

Allocate memory on the default accelerator.

If size is 0, no memory is allocated, *ptr returns nullptr, and hipSuccess is returned.

See also:

hipMallocPitch, *hipFree*, *hipMallocArray*, *hipFreeArray*, *hipMalloc3D*, *hipMalloc3DArray*, *hipHostFree*, *hipHostMalloc*

Parameters

- **ptr** – [out] Pointer to the allocated memory
- **size** – [in] Requested memory size

Returns

hipSuccess, hipErrorOutOfMemory, hipErrorInvalidValue (bad context, null *ptr)

hipError_t **hipExtMallocWithFlags**(void **ptr, size_t sizeBytes, unsigned int flags)

Allocate memory on the default accelerator.

If requested memory size is 0, no memory is allocated, *ptr returns nullptr, and hipSuccess is returned.

The memory allocation flag should be either hipDeviceMallocDefault, hipDeviceMallocFinegrained, hipDeviceMallocUncached, or hipMallocSignalMemory. If the flag is any other value, the API returns hipErrorInvalidValue.

See also:

hipMallocPitch, *hipFree*, *hipMallocArray*, *hipFreeArray*, *hipMalloc3D*, *hipMalloc3DArray*, *hipHostFree*, *hipHostMalloc*

Parameters

- **ptr** – [out] Pointer to the allocated memory
- **sizeBytes** – [in] Requested memory size
- **flags** – [in] Type of memory allocation

Returns

hipSuccess, hipErrorOutOfMemory, hipErrorInvalidValue (bad context, null *ptr)

hipError_t **hipHostMalloc**(void **ptr, size_t size, unsigned int flags)

Allocates device accessible page locked (pinned) host memory.

This API allocates pinned host memory which is mapped into the address space of all GPUs in the system, the memory can be accessed directly by the GPU device, and can be read or written with much higher bandwidth than pageable memory obtained with functions such as malloc().

Using the pinned host memory, applications can implement faster data transfers for HostToDevice and DeviceToHost. The runtime tracks the hipHostMalloc allocations and can avoid some of the setup required for regular unpinned memory.

When the memory accesses are infrequent, zero-copy memory can be a good choice, for coherent allocation. GPU can directly access the host memory over the CPU/GPU interconnect, without need to copy the data.

Currently the allocation granularity is 4KB for the API.

Developers need to choose proper allocation flag with consideration of synchronization.

If no input for flags, it will be the default pinned memory allocation on the host.

See also:

hipSetDeviceFlags, hipHostFree

Parameters

- **ptr** – [out] Pointer to the allocated host pinned memory
- **size** – [in] Requested memory size in bytes If size is 0, no memory is allocated, *ptr returns nullptr, and hipSuccess is returned.
- **flags** – [in] Type of host memory allocation

Returns

hipSuccess, hipErrorOutOfMemory

hipError_t **hipHostAlloc**(void **ptr, size_t size, unsigned int flags)

Allocate device accessible page locked host memory.

If size is 0, no memory is allocated, *ptr returns nullptr, and hipSuccess is returned.

Flags:

- hipHostAllocDefault Default pinned memory allocation on the host.
- hipHostAllocPortable Memory is considered allocated by all contexts.
- hipHostAllocMapped Map the allocation into the address space for the current device.
- hipHostAllocWriteCombined Allocates the memory as write-combined.

Parameters

- **ptr** – [out] Pointer to the allocated host pinned memory
- **size** – [in] Requested memory size in bytes
- **flags** – [in] Type of host memory allocation see below

Returns

hipSuccess, hipErrorOutOfMemory, hipErrorInvalidValue

hipError_t **hipHostGetDevicePointer**(void **devPtr, void *hstPtr, unsigned int flags)

Get Device pointer from Host Pointer allocated through hipHostMalloc.

See also:

hipSetDeviceFlags, [hipHostMalloc](#)

Parameters

- **devPtr** – [out] Device Pointer mapped to passed host pointer
- **hstPtr** – [in] Host Pointer allocated through hipHostMalloc
- **flags** – [in] Flags to be passed for extension

Returns

hipSuccess, hipErrorInvalidValue, hipErrorOutOfMemory

hipError_t **hipHostGetFlags**(unsigned int *flagsPtr, void *hostPtr)

Return flags associated with host pointer.

See also:

hipHostMalloc

Parameters

- **flagsPtr** – [out] Memory location to store flags
- **hostPtr** – [in] Host Pointer allocated through hipHostMalloc

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipHostRegister**(void *hostPtr, size_t sizeBytes, unsigned int flags)

Register host memory so it can be accessed from the current device.

Flags:

- **hipHostRegisterDefault** Memory is Mapped and Portable
- **hipHostRegisterPortable** Memory is considered registered by all contexts. HIP only supports one context so this is always assumed true.
- **hipHostRegisterMapped** Map the allocation into the address space for the current device. The device pointer can be obtained with *hipHostGetDevicePointer*.

After registering the memory, use *hipHostGetDevicePointer* to obtain the mapped device pointer. On many systems, the mapped device pointer will have a different value than the mapped host pointer. Applications must use the device pointer in device code, and the host pointer in host code.

On some systems, registered memory is pinned. On some systems, registered memory may not be actually be pinned but uses OS or hardware facilities to all GPU access to the host memory.

Developers are strongly encouraged to register memory blocks which are aligned to the host cache-line size. (typically 64-bytes but can be obtains from the CPUID instruction).

If registering non-aligned pointers, the application must take care when register pointers from the same cache line on different devices. HIP's coarse-grained synchronization model does not guarantee correct results if different devices write to different parts of the same cache block - typically one of the writes will “win” and overwrite data from the other registered memory region.

See also:

hipHostUnregister, *hipHostGetFlags*, *hipHostGetDevicePointer*

Parameters

- **hostPtr** – [out] Pointer to host memory to be registered.
- **sizeBytes** – [in] Size of the host memory
- **flags** – [in] See below.

Returns

hipSuccess, hipErrorOutOfMemory

hipError_t **hipHostUnregister**(void *hostPtr)

Un-register host pointer.

See also:

hipHostRegister

Parameters

hostPtr – [in] Host pointer previously registered with *hipHostRegister*

Returns

Error code

hipError_t **hipMallocPitch**(void **ptr, size_t *pitch, size_t width, size_t height)

Allocates at least width (in bytes) * height bytes of linear memory. Padding may occur to ensure alignment requirements are met for the given row. The change in width size due to padding will be returned in *pitch. Currently the alignment is set to 128 bytes.

If size is 0, no memory is allocated, *ptr returns nullptr, and hipSuccess is returned.

See also:

hipMalloc, *hipFree*, *hipMallocArray*, *hipFreeArray*, *hipHostFree*, *hipMalloc3D*, *hipMalloc3DArray*, *hipHostMalloc*

Parameters

- **ptr** – [out] Pointer to the allocated device memory
- **pitch** – [out] Pitch for allocation (in bytes)
- **width** – [in] Requested pitched allocation width (in bytes)
- **height** – [in] Requested pitched allocation height

Returns

Error code

hipError_t **hipMemAllocPitch**(hipDeviceptr_t *dptr, size_t *pitch, size_t widthInBytes, size_t height, unsigned int elementSizeBytes)

Allocates at least width (in bytes) * height bytes of linear memory. Padding may occur to ensure alignment requirements are met for the given row. The change in width size due to padding will be returned in *pitch. Currently the alignment is set to 128 bytes.

If size is 0, no memory is allocated, *ptr* returns *nullptr*, and *hipSuccess* is returned. The intended usage of *pitch* is as a separate parameter of the allocation, used to compute addresses within the 2D array. Given the row and column of an array element of type *T*, the address is computed as: $pElement = (T*)((char*)BaseAddress + Row * Pitch) + Column;$

See also:

hipMalloc, *hipFree*, *hipMallocArray*, *hipFreeArray*, *hipHostFree*, *hipMalloc3D*, *hipMalloc3DArray*, *hipHostMalloc*

Parameters

- **dptr** – [out] Pointer to the allocated device memory
- **pitch** – [out] Pitch for allocation (in bytes)
- **widthInBytes** – [in] Requested pitched allocation width (in bytes)
- **height** – [in] Requested pitched allocation height
- **elementSizeBytes** – [in] The size of element bytes, should be 4, 8 or 16

Returns

Error code

hipError_t **hipFree**(void *ptr)

Free memory allocated by the hcc hip memory allocation API. This API performs an implicit hipDeviceSynchronize() call. If pointer is NULL, the hip runtime is initialized and hipSuccess is returned.

See also:

hipMalloc, *hipMallocPitch*, *hipMallocArray*, *hipFreeArray*, *hipHostFree*, *hipMalloc3D*, *hipMalloc3DArray*, *hipHostMalloc*

Parameters

ptr – [in] Pointer to memory to be freed

Returns

hipSuccess

Returns

hipErrorInvalidDevicePointer (if pointer is invalid, including host pointers allocated with hipHostMalloc)

hipError_t **hipFreeHost**(void *ptr)

Frees page-locked memory This API performs an implicit hipDeviceSynchronize() call. If pointer is NULL, the hip runtime is initialized and hipSuccess is returned.

Parameters

ptr – [in] Pointer to memory to be freed

Returns

hipSuccess, hipErrorInvalidValue (if pointer is invalid, including device pointers allocated with hipMalloc)

hipError_t **hipMemcpy**(void *dst, const void *src, size_t sizeBytes, hipMemcpyKind kind)

Copy data from src to dst.

It supports memory from host to device, device to host, device to device and host to host The src and dst must not overlap.

For hipMemcpy, the copy is always performed by the current device (set by hipSetDevice). For multi-gpu or peer-to-peer configurations, it is recommended to set the current device to the device where the src data is physically located. For optimal peer-to-peer copies, the copy device must be able to access the src and dst pointers (by calling hipDeviceEnablePeerAccess with copy agent as the current device and src/dest as the peerDevice argument. if

this is not done, the `hipMemcpy` will still work, but will perform the copy using a staging buffer on the host. Calling `hipMemcpy` with `dst` and `src` pointers that do not match the `hipMemcpyKind` results in undefined behavior.

See also:

hipArrayCreate, *hipArrayDestroy*, *hipArrayGetDescriptor*, *hipMemAlloc*, *hipMemAllocHost*, *hipMemAllocPitch*, *hipMemcpy2D*, *hipMemcpy2DAsync*, *hipMemcpy2DUnaligned*, *hipMemcpyAtoA*, *hipMemcpyAtoD*, *hipMemcpyAtoH*, *hipMemcpyAtoHAsync*, *hipMemcpyDtoA*, *hipMemcpyDtoD*, *hipMemcpyDtoDAsync*, *hipMemcpyDtoH*, *hipMemcpyDtoHAsync*, *hipMemcpyHtoA*, *hipMemcpyHtoAAsync*, *hipMemcpyHtoDAsync*, *hipMemFree*, *hipMemFreeHost*, *hipMemGetAddressRange*, *hipMemGetInfo*, *hipMemHostAlloc*, *hipMemHostGetDevicePointer*

Parameters

- **dst** – [out] Data being copy to
- **src** – [in] Data being copy from
- **sizeBytes** – [in] Data size in bytes
- **kind** – [in] Kind of transfer

Returns

`hipSuccess`, `hipErrorInvalidValue`, `hipErrorUnknown`

`hipError_t` **hipMemcpyWithStream**(void *dst, const void *src, size_t sizeBytes, hipMemcpyKind kind, hipStream_t stream)

Memory copy on the stream. It allows single or multiple devices to do memory copy on single or multiple streams.

See also:

hipMemcpy, *hipStreamCreate*, *hipStreamSynchronize*, *hipStreamDestroy*, *hipSetDevice*, *hipLaunchKernelGGL*

Parameters

- **dst** – [out] Data being copy to
- **src** – [in] Data being copy from
- **sizeBytes** – [in] Data size in bytes
- **kind** – [in] Kind of transfer
- **stream** – [in] Valid stream

Returns

`hipSuccess`, `hipErrorInvalidValue`, `hipErrorUnknown`, `hipErrorContextIsDestroyed`

`hipError_t` **hipMemcpyHtoD**(hipDeviceptr_t dst, void *src, size_t sizeBytes)

Copy data from Host to Device.

See also:

hipArrayCreate, *hipArrayDestroy*, *hipArrayGetDescriptor*, *hipMemAlloc*, *hipMemAllocHost*, *hipMemAllocPitch*, *hipMemcpy2D*, *hipMemcpy2DAsync*, *hipMemcpy2DUnaligned*, *hipMemcpyAtoA*, *hipMemcpyAtoD*,

hipMemcpyToH, hipMemcpyToHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer

Parameters

- **dst** – [out] Data being copy to
- **src** – [in] Data being copy from
- **sizeBytes** – [in] Data size in bytes

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue

hipError_t **hipMemcpyToH**(void *dst, hipDeviceptr_t src, size_t sizeBytes)

Copy data from Device to Host.

See also:

hipArrayCreate, hipArrayDestroy, hipArrayGetDescriptor, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpyToA, hipMemcpyToD, hipMemcpyToH, hipMemcpyToHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer

Parameters

- **dst** – [out] Data being copy to
- **src** – [in] Data being copy from
- **sizeBytes** – [in] Data size in bytes

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue

hipError_t **hipMemcpyDtoD**(hipDeviceptr_t dst, hipDeviceptr_t src, size_t sizeBytes)

Copy data from Device to Device.

See also:

hipArrayCreate, hipArrayDestroy, hipArrayGetDescriptor, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpyToA, hipMemcpyToD, hipMemcpyToH, hipMemcpyToHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer

Parameters

- **dst** – [out] Data being copy to

- **src** – [in] Data being copy from
- **sizeBytes** – [in] Data size in bytes

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue

hipError_t **hipMemcpyAtoD**(hipDeviceptr_t dstDevice, hipArray_t srcArray, size_t srcOffset, size_t ByteCount)

Copies from one 1D array to device memory.

See also:

hipArrayCreate, hipArrayDestroy, hipArrayGetDescriptor, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpyAtoA, hipMemcpyAtoD, hipMemcpyAtoH, hipMemcpyAtoHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer

Parameters

- **dstDevice** – [out] Destination device pointer
- **srcArray** – [in] Source array
- **srcOffset** – [in] Offset in bytes of source array
- **ByteCount** – [in] Size of memory copy in bytes

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue

hipError_t **hipMemcpyDtoA**(hipArray_t dstArray, size_t dstOffset, hipDeviceptr_t srcDevice, size_t ByteCount)

Copies from device memory to a 1D array.

See also:

hipArrayCreate, hipArrayDestroy, hipArrayGetDescriptor, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpyAtoA, hipMemcpyAtoD, hipMemcpyAtoH, hipMemcpyAtoHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer

Parameters

- **dstArray** – [out] Destination array
- **dstOffset** – [in] Offset in bytes of destination array
- **srcDevice** – [in] Source device pointer
- **ByteCount** – [in] Size of memory copy in bytes

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue

hipError_t **hipMemcpyAtoA**(hipArray_t dstArray, size_t dstOffset, hipArray_t srcArray, size_t srcOffset, size_t ByteCount)

Copies from one 1D array to another.

See also:

hipArrayCreate, hipArrayDestroy, hipArrayGetDescriptor, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpyAtoA, hipMemcpyAtoD, hipMemcpyAtoH, hipMemcpyAtoHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer

Parameters

- **dstArray** – [out] Destination array
- **dstOffset** – [in] Offset in bytes of destination array
- **srcArray** – [in] Source array
- **srcOffset** – [in] Offset in bytes of source array
- **ByteCount** – [in] Size of memory copy in bytes

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue

hipError_t **hipMemcpyHtoDAsync**(hipDeviceptr_t dst, void *src, size_t sizeBytes, hipStream_t stream)

Copy data from Host to Device asynchronously.

See also:

hipArrayCreate, hipArrayDestroy, hipArrayGetDescriptor, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpyAtoA, hipMemcpyAtoD, hipMemcpyAtoH, hipMemcpyAtoHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer

Parameters

- **dst** – [out] Data being copy to
- **src** – [in] Data being copy from
- **sizeBytes** – [in] Data size in bytes
- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue

hipError_t **hipMemcpyDtoHAsync**(void *dst, hipDeviceptr_t src, size_t sizeBytes, hipStream_t stream)

Copy data from Device to Host asynchronously.

See also:

hipArrayCreate, hipArrayDestroy, hipArrayGetDescriptor, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpyAtoA, hipMemcpyAtoD, hipMemcpyAtoH, hipMemcpyAtoHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer

Parameters

- **dst** – [out] Data being copy to
- **src** – [in] Data being copy from
- **sizeBytes** – [in] Data size in bytes
- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue

hipError_t **hipMemcpyDtoDAsync**(hipDeviceptr_t dst, hipDeviceptr_t src, size_t sizeBytes, hipStream_t stream)

Copy data from Device to Device asynchronously.

See also:

hipArrayCreate, hipArrayDestroy, hipArrayGetDescriptor, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpyAtoA, hipMemcpyAtoD, hipMemcpyAtoH, hipMemcpyAtoHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer

Parameters

- **dst** – [out] Data being copy to
- **src** – [in] Data being copy from
- **sizeBytes** – [in] Data size in bytes
- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue

hipError_t **hipMemcpyAtoHAsync**(void *dstHost, hipArray_t srcArray, size_t srcOffset, size_t ByteCount, hipStream_t stream)

Copies from one 1D array to host memory.

See also:

hipArrayCreate, *hipArrayDestroy*, *hipArrayGetDescriptor*, *hipMemAlloc*, *hipMemAllocHost*, *hipMemAllocPitch*, *hipMemcpy2D*, *hipMemcpy2DAsync*, *hipMemcpy2DUnaligned*, *hipMemcpyAtoA*, *hipMemcpyAtoD*, *hipMemcpyAtoH*, *hipMemcpyAtoHAsync*, *hipMemcpyDtoA*, *hipMemcpyDtoD*, *hipMemcpyDtoDAsync*, *hipMemcpyDtoH*, *hipMemcpyDtoHAsync*, *hipMemcpyHtoA*, *hipMemcpyHtoAAsync*, *hipMemcpyHtoDAsync*, *hipMemFree*, *hipMemFreeHost*, *hipMemGetAddressRange*, *hipMemGetInfo*, *hipMemHostAlloc*, *hipMemHostGetDevicePointer*

Parameters

- **dstHost** – [out] Destination pointer
- **srcArray** – [in] Source array
- **srcOffset** – [in] Offset in bytes of source array
- **ByteCount** – [in] Size of memory copy in bytes
- **stream** – [in] Stream identifier

Returns

hipSuccess, *hipErrorDeinitialized*, *hipErrorNotInitialized*, *hipErrorInvalidContext*, *hipErrorInvalidValue*

hipError_t **hipMemcpyHtoAAsync**(*hipArray_t* dstArray, *size_t* dstOffset, const void *srcHost, *size_t* ByteCount, *hipStream_t* stream)

Copies from host memory to a 1D array.

See also:

hipArrayCreate, *hipArrayDestroy*, *hipArrayGetDescriptor*, *hipMemAlloc*, *hipMemAllocHost*, *hipMemAllocPitch*, *hipMemcpy2D*, *hipMemcpy2DAsync*, *hipMemcpy2DUnaligned*, *hipMemcpyAtoA*, *hipMemcpyAtoD*, *hipMemcpyAtoH*, *hipMemcpyAtoHAsync*, *hipMemcpyDtoA*, *hipMemcpyDtoD*, *hipMemcpyDtoDAsync*, *hipMemcpyDtoH*, *hipMemcpyDtoHAsync*, *hipMemcpyHtoA*, *hipMemcpyHtoAAsync*, *hipMemcpyHtoDAsync*, *hipMemFree*, *hipMemFreeHost*, *hipMemGetAddressRange*, *hipMemGetInfo*, *hipMemHostAlloc*, *hipMemHostGetDevicePointer*

Parameters

- **dstArray** – [out] Destination array
- **dstOffset** – [in] Offset in bytes of destination array
- **srcHost** – [in] Source host pointer
- **ByteCount** – [in] Size of memory copy in bytes
- **stream** – [in] Stream identifier

Returns

hipSuccess, *hipErrorDeinitialized*, *hipErrorNotInitialized*, *hipErrorInvalidContext*, *hipErrorInvalidValue*

hipError_t **hipModuleGetGlobal**(*hipDeviceptr_t* *dptr, *size_t* *bytes, *hipModule_t* hmod, const char *name)

Returns a global pointer from a module. Returns in *dptr and *bytes the pointer and size of the global of name name located in module hmod. If no variable of that name exists, it returns *hipErrorNotFound*. Both parameters dptr and bytes are optional. If one of them is NULL, it is ignored and *hipSuccess* is returned.

Parameters

- **dptr** – [out] Returns global device pointer
- **bytes** – [out] Returns global size in bytes
- **hmod** – [in] Module to retrieve global from
- **name** – [in] Name of global to retrieve

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotFound, hipErrorInvalidContext

hipError_t **hipGetSymbolAddress**(void **devPtr, const void *symbol)

Gets device pointer associated with symbol on the device.

Parameters

- **devPtr** – [out] pointer to the device associated the symbole
- **symbol** – [in] pointer to the symbole of the device

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipGetSymbolSize**(size_t *size, const void *symbol)

Gets the size of the given symbol on the device.

Parameters

- **symbol** – [in] pointer to the device symbole
- **size** – [out] pointer to the size

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipGetProcAddress**(const char *symbol, void **pfn, int hipVersion, uint64_t flags, hipDriverProcAddressQueryResult *symbolStatus)

Gets the pointer of requested HIP driver function.

Returns hipSuccess if the returned pfn is addressed to the pointer of found driver function.

Parameters

- **symbol** – [in] The Symbol name of the driver function to request.
- **pfn** – [out] Output pointer to the requested driver function.
- **hipVersion** – [in] The HIP version for the requested driver function symbol. HIP version is defined as 100*version_major + version_minor. For example, in HIP 6.1, the hipversion is 601, for the symbol function “hipGetDeviceProperties”, the specified hipVersion 601 is greater or equal to the version 600, the symbol function will be handle properly as backend compatible function.
- **flags** – [in] Currently only default flag is supported.
- **symbolStatus** – [out] Optional enumeration for returned status of searching for symbol driver function based on the input hipVersion.

Returns

hipSuccess, hipErrorInvalidValue.

hipError_t **hipMemcpyToSymbol**(const void *symbol, const void *src, size_t sizeBytes, size_t offset, hipMemcpyKind kind)

Copies data to the given symbol on the device. Symbol HIP APIs allow a kernel to define a device-side data symbol which can be accessed on the host side. The symbol can be in `__constant` or device space. Note that the symbol name needs to be encased in the `HIP_SYMBOL` macro. This also applies to `hipMemcpyFromSymbol`, `hipGetSymbolAddress`, and `hipGetSymbolSize`. For detailed usage, see the [memcpyToSymbol example](#) in the HIP Porting Guide.

Parameters

- **symbol** – [out] pointer to the device symbole
- **src** – [in] pointer to the source address
- **sizeBytes** – [in] size in bytes to copy
- **offset** – [in] offset in bytes from start of symbole
- **kind** – [in] type of memory transfer

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemcpyToSymbolAsync**(const void *symbol, const void *src, size_t sizeBytes, size_t offset, hipMemcpyKind kind, hipStream_t stream)

Copies data to the given symbol on the device asynchronously.

Parameters

- **symbol** – [out] pointer to the device symbole
- **src** – [in] pointer to the source address
- **sizeBytes** – [in] size in bytes to copy
- **offset** – [in] offset in bytes from start of symbole
- **kind** – [in] type of memory transfer
- **stream** – [in] stream identifier

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemcpyFromSymbol**(void *dst, const void *symbol, size_t sizeBytes, size_t offset, hipMemcpyKind kind)

Copies data from the given symbol on the device.

Parameters

- **dst** – [out] Returns pointer to destinition memory address
- **symbol** – [in] Pointer to the symbole address on the device
- **sizeBytes** – [in] Size in bytes to copy
- **offset** – [in] Offset in bytes from the start of symbole
- **kind** – [in] Type of memory transfer

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemcpyFromSymbolAsync**(void *dst, const void *symbol, size_t sizeBytes, size_t offset, hipMemcpyKind kind, hipStream_t stream)

Copies data from the given symbol on the device asynchronously.

Parameters

- **dst** – [out] Returns pointer to destination memory address
- **symbol** – [in] pointer to the symbol address on the device
- **sizeBytes** – [in] size in bytes to copy
- **offset** – [in] offset in bytes from the start of symbol
- **kind** – [in] type of memory transfer
- **stream** – [in] stream identifier

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemcpyAsync**(void *dst, const void *src, size_t sizeBytes, hipMemcpyKind kind, hipStream_t stream)

Copy data from src to dst asynchronously.

For multi-gpu or peer-to-peer configurations, it is recommended to use a stream which is attached to the device where the src data is physically located. For optimal peer-to-peer copies, the copy device must be able to access the src and dst pointers (by calling `hipDeviceEnablePeerAccess` with copy agent as the current device and src/dst as the peerDevice argument. if this is not done, the `hipMemcpy` will still work, but will perform the copy using a staging buffer on the host.

See also:

hipMemcpy, *hipMemcpy2D*, *hipMemcpyToArray*, *hipMemcpy2DToArray*, *hipMemcpyFromArray*, *hipMemcpy2DFromArray*, *hipMemcpyArrayToArray*, *hipMemcpy2DArrayToArray*, *hipMemcpyToSymbol*, *hipMemcpyFromSymbol*, *hipMemcpy2DAsync*, *hipMemcpyToArrayAsync*, *hipMemcpy2DToArrayAsync*, *hipMemcpyFromArrayAsync*, *hipMemcpy2DFromArrayAsync*, *hipMemcpyToSymbolAsync*, *hipMemcpyFromSymbolAsync*

Warning: If host or dest are not pinned, the memory copy will be performed synchronously. For best performance, use `hipHostMalloc` to allocate host memory that is transferred asynchronously.

Warning: on HCC `hipMemcpyAsync` does not support overlapped H2D and D2H copies. For `hipMemcpy`, the copy is always performed by the device associated with the specified stream.

Parameters

- **dst** – [out] Data being copy to
- **src** – [in] Data being copy from
- **sizeBytes** – [in] Data size in bytes
- **kind** – [in] Type of memory transfer

- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorInvalidValue, hipErrorUnknown

hipError_t **hipMemset**(void *dst, int value, size_t sizeBytes)

Fills the first sizeBytes bytes of the memory area pointed to by dest with the constant byte value value.

Parameters

- **dst** – [out] Data being filled
- **value** – [in] Value to be set
- **sizeBytes** – [in] Data size in bytes

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotInitialized

hipError_t **hipMemsetD8**(hipDeviceptr_t dest, unsigned char value, size_t count)

Fills the first sizeBytes bytes of the memory area pointed to by dest with the constant byte value value.

Parameters

- **dest** – [out] Data ptr to be filled
- **value** – [in] Value to be set
- **count** – [in] Number of values to be set

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotInitialized

hipError_t **hipMemsetD8Async**(hipDeviceptr_t dest, unsigned char value, size_t count, hipStream_t stream)

Fills the first sizeBytes bytes of the memory area pointed to by dest with the constant byte value value.

hipMemsetD8Async() is asynchronous with respect to the host, so the call may return before the memset is complete. The operation can optionally be associated to a stream by passing a non-zero stream argument. If stream is non-zero, the operation may overlap with operations in other streams.

Parameters

- **dest** – [out] Data ptr to be filled
- **value** – [in] Constant value to be set
- **count** – [in] Number of values to be set
- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotInitialized

hipError_t **hipMemsetD16**(hipDeviceptr_t dest, unsigned short value, size_t count)

Fills the first sizeBytes bytes of the memory area pointed to by dest with the constant short value value.

Parameters

- **dest** – [out] Data ptr to be filled
- **value** – [in] Constant value to be set
- **count** – [in] Number of values to be set

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotInitialized

hipError_t **hipMemsetD16Async**(hipDeviceptr_t dest, unsigned short value, size_t count, hipStream_t stream)

Fills the first sizeBytes bytes of the memory area pointed to by dest with the constant short value value.

hipMemsetD16Async() is asynchronous with respect to the host, so the call may return before the memset is complete. The operation can optionally be associated to a stream by passing a non-zero stream argument. If stream is non-zero, the operation may overlap with operations in other streams.

Parameters

- **dest** – [out] Data ptr to be filled
- **value** – [in] Constant value to be set
- **count** – [in] Number of values to be set
- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotInitialized

hipError_t **hipMemsetD32**(hipDeviceptr_t dest, int value, size_t count)

Fills the memory area pointed to by dest with the constant integer value for specified number of times.

Parameters

- **dest** – [out] Data being filled
- **value** – [in] Constant value to be set
- **count** – [in] Number of values to be set

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotInitialized

hipError_t **hipMemsetAsync**(void *dst, int value, size_t sizeBytes, hipStream_t stream)

Fills the first sizeBytes bytes of the memory area pointed to by dev with the constant byte value value.

hipMemsetAsync() is asynchronous with respect to the host, so the call may return before the memset is complete. The operation can optionally be associated to a stream by passing a non-zero stream argument. If stream is non-zero, the operation may overlap with operations in other streams.

Parameters

- **dst** – [out] Pointer to device memory
- **value** – [in] Value to set for each byte of specified memory
- **sizeBytes** – [in] Size in bytes to set
- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemsetD32Async**(hipDeviceptr_t dst, int value, size_t count, hipStream_t stream)

Fills the memory area pointed to by dev with the constant integer value for specified number of times.

hipMemsetD32Async() is asynchronous with respect to the host, so the call may return before the memset is complete. The operation can optionally be associated to a stream by passing a non-zero stream argument. If stream is non-zero, the operation may overlap with operations in other streams.

Parameters

- **dst** – [out] Pointer to device memory
- **value** – [in] Value to set for each byte of specified memory

- **count** – [in] Number of values to be set
- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemset2D**(void *dst, size_t pitch, int value, size_t width, size_t height)

Fills the memory area pointed to by dst with the constant value.

Parameters

- **dst** – [out] Pointer to device memory
- **pitch** – [in] Data size in bytes
- **value** – [in] Constant value to be set
- **width** – [in]
- **height** – [in]

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemset2DAsync**(void *dst, size_t pitch, int value, size_t width, size_t height, hipStream_t stream)

Fills asynchronously the memory area pointed to by dst with the constant value.

Parameters

- **dst** – [in] Pointer to 2D device memory
- **pitch** – [in] Pitch size in bytes
- **value** – [in] Value to be set for each byte of specified memory
- **width** – [in] Width of matrix set columns in bytes
- **height** – [in] Height of matrix set rows in bytes
- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemset3D**(hipPitchedPtr pitchedDevPtr, int value, hipExtent extent)

Fills synchronously the memory area pointed to by pitchedDevPtr with the constant value.

Parameters

- **pitchedDevPtr** – [in] Pointer to pitched device memory
- **value** – [in] Value to set for each byte of specified memory
- **extent** – [in] Size parameters for width field in bytes in device memory

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemset3DAsync**(hipPitchedPtr pitchedDevPtr, int value, hipExtent extent, hipStream_t stream)

Fills asynchronously the memory area pointed to by pitchedDevPtr with the constant value.

Parameters

- **pitchedDevPtr** – [in] Pointer to pitched device memory
- **value** – [in] Value to set for each byte of specified memory

- **extent** – [in] Size parameters for width field in bytes in device memory
- **stream** – [in] Stream identifier

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMemGetInfo**(size_t *free, size_t *total)

Query memory info.

On ROCm, this function gets the actual free memory left on the current device, so supports the cases while running multi-workload (such as multiple processes, multiple threads, and multiple GPUs).

Warning: On Windows, the free memory only accounts for memory allocated by this process and may be optimistic.

Parameters

- **free** – [out] Returns free memory on the current device in bytes
- **total** – [out] Returns total allocatable memory on the current device in bytes

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipMemPtrGetInfo**(void *ptr, size_t *size)

Get allocated memory size via memory pointer.

This function gets the allocated shared virtual memory size from memory pointer.

Parameters

- **ptr** – [in] Pointer to allocated memory
- **size** – [out] Returns the allocated memory size in bytes

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipMallocArray**(hipArray_t *array, const hipChannelFormatDesc *desc, size_t width, size_t height, unsigned int flags)

Allocate an array on the device.

See also:

[*hipMalloc*](#), [*hipMallocPitch*](#), [*hipFree*](#), [*hipFreeArray*](#), [*hipHostMalloc*](#), [*hipHostFree*](#)

Parameters

- **array** – [out] Pointer to allocated array in device memory
- **desc** – [in] Requested channel format
- **width** – [in] Requested array allocation width
- **height** – [in] Requested array allocation height
- **flags** – [in] Requested properties of allocated array

Returns

hipSuccess, hipErrorOutOfMemory

hipError_t **hipArrayCreate**(hipArray_t *pHandle, const HIP_ARRAY_DESCRIPTOR *pAllocateArray)

Create an array memory pointer on the device.

See also:

hipMallocArray, hipArrayDestroy, hipFreeArray

Parameters

- **pHandle** – [out] Pointer to the array memory
- **pAllocateArray** – [in] Requested array descriptor

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotSupported

hipError_t **hipArrayDestroy**(hipArray_t array)

Destroy an array memory pointer on the device.

See also:

hipArrayCreate, hipArrayDestroy, hipFreeArray

Parameters

array – [in] Pointer to the array memory

Returns

hipSuccess, hipErrorInvalidValue

hipError_t **hipArray3DCreate**(hipArray_t *array, const HIP_ARRAY3D_DESCRIPTOR *pAllocateArray)

Create a 3D array memory pointer on the device.

See also:

hipMallocArray, hipArrayDestroy, hipFreeArray

Parameters

- **array** – [out] Pointer to the 3D array memory
- **pAllocateArray** – [in] Requested array descriptor

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotSupported

hipError_t **hipMalloc3D**(hipPitchedPtr *pitchedDevPtr, hipExtent extent)

Create a 3D memory pointer on the device.

See also:

hipMallocPitch, hipMemGetInfo, hipFree

Parameters

- **pitchedDevPtr** – [out] Pointer to the 3D memory

- **extent** – [in] Requested extent

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotSupported

hipError_t **hipFreeArray**(hipArray_t array)

Frees an array on the device.

See also:

hipMalloc, hipMallocPitch, hipFree, hipMallocArray, hipHostMalloc, hipHostFree

Parameters

array – [in] Pointer to array to free

Returns

hipSuccess, hipErrorInvalidValue, hipErrorNotInitialized

hipError_t **hipMalloc3DArray**(hipArray_t *array, const struct hipChannelFormatDesc *desc, struct hipExtent extent, unsigned int flags)

Allocate an array on the device.

See also:

hipMalloc, hipMallocPitch, hipFree, hipFreeArray, hipHostMalloc, hipHostFree

Parameters

- **array** – [out] Pointer to allocated array in device memory
- **desc** – [in] Requested channel format
- **extent** – [in] Requested array allocation width, height and depth
- **flags** – [in] Requested properties of allocated array

Returns

hipSuccess, hipErrorOutOfMemory

hipError_t **hipArrayGetInfo**(hipChannelFormatDesc *desc, hipExtent *extent, unsigned int *flags, hipArray_t array)

Gets info about the specified array.

See also:

hipArrayGetDescriptor, hipArray3DGetDescriptor

Parameters

- **desc** – [out] - Returned array type
- **extent** – [out] - Returned array shape. 2D arrays will have depth of zero
- **flags** – [out] - Returned array flags
- **array** – [in] - The HIP array to get info for

Returns

hipSuccess, hipErrorInvalidValue hipErrorInvalidHandle

hipError_t **hipArrayGetDescriptor**(HIP_ARRAY_DESCRIPTOR *pArrayDescriptor, hipArray_t array)

Gets a 1D or 2D array descriptor.

See also:

hipArray3DCreate, hipArray3DGetDescriptor, hipArrayCreate, hipArrayDestroy, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpy3D, hipMemcpy3DAsync, hipMemcpyAtoA, hipMemcpyAtoD, hipMemcpyAtoH, hipMemcpyAtoHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoD, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer, hipMemsetD8, hipMemsetD16, hipMemsetD32, hipArrayGetInfo

Parameters

- **pArrayDescriptor** – [out] - Returned array descriptor
- **array** – [in] - Array to get descriptor of

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue hipErrorInvalidHandle

hipError_t **hipArray3DGetDescriptor**(HIP_ARRAY3D_DESCRIPTOR *pArrayDescriptor, hipArray_t array)

Gets a 3D array descriptor.

See also:

hipArray3DCreate, hipArrayCreate, hipArrayDestroy, hipArrayGetDescriptor, hipMemAlloc, hipMemAllocHost, hipMemAllocPitch, hipMemcpy2D, hipMemcpy2DAsync, hipMemcpy2DUnaligned, hipMemcpy3D, hipMemcpy3DAsync, hipMemcpyAtoA, hipMemcpyAtoD, hipMemcpyAtoH, hipMemcpyAtoHAsync, hipMemcpyDtoA, hipMemcpyDtoD, hipMemcpyDtoDAsync, hipMemcpyDtoH, hipMemcpyDtoHAsync, hipMemcpyHtoA, hipMemcpyHtoAAsync, hipMemcpyHtoD, hipMemcpyHtoDAsync, hipMemFree, hipMemFreeHost, hipMemGetAddressRange, hipMemGetInfo, hipMemHostAlloc, hipMemHostGetDevicePointer, hipMemsetD8, hipMemsetD16, hipMemsetD32, hipArrayGetInfo

Parameters

- **pArrayDescriptor** – [out] - Returned 3D array descriptor
- **array** – [in] - 3D array to get descriptor of

Returns

hipSuccess, hipErrorDeinitialized, hipErrorNotInitialized, hipErrorInvalidContext, hipErrorInvalidValue hipErrorInvalidHandle, hipErrorContextIsDestroyed

hipError_t **hipMemcpy2D**(void *dst, size_t dpitch, const void *src, size_t spitch, size_t width, size_t height, hipMemcpyKind kind)

Copies data between host and device.

See also:

[*hipMemcpy*](#), [*hipMemcpyToArray*](#), [*hipMemcpy2DToArray*](#), [*hipMemcpyFromArray*](#), [*hipMemcpyToSymbol*](#), [*hipMemcpyAsync*](#)

Parameters

- **dst** – [in] Destination memory address
- **dpitch** – [in] Pitch of destination memory
- **src** – [in] Source memory address
- **spitch** – [in] Pitch of source memory
- **width** – [in] Width of matrix transfer (columns in bytes)
- **height** – [in] Height of matrix transfer (rows)
- **kind** – [in] Type of transfer

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpyParam2D**(const hip_Memcpy2D *pCopy)

Copies memory for 2D arrays.

See also:

[*hipMemcpy*](#), [*hipMemcpy2D*](#), [*hipMemcpyToArray*](#), [*hipMemcpy2DToArray*](#), [*hipMemcpyFromArray*](#), [*hipMemcpyToSymbol*](#), [*hipMemcpyAsync*](#)

Parameters

pCopy – [in] Parameters for the memory copy

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpyParam2DAsync**(const hip_Memcpy2D *pCopy, hipStream_t stream)

Copies memory for 2D arrays.

See also:

[*hipMemcpy*](#), [*hipMemcpy2D*](#), [*hipMemcpyToArray*](#), [*hipMemcpy2DToArray*](#), [*hipMemcpyFromArray*](#), [*hipMemcpyToSymbol*](#), [*hipMemcpyAsync*](#)

Parameters

- **pCopy** – [in] Parameters for the memory copy
- **stream** – [in] Stream to use

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpy2DAsync**(void *dst, size_t dpitch, const void *src, size_t spitch, size_t width, size_t height, hipMemcpyKind kind, hipStream_t stream)

Copies data between host and device.

See also:

hipMemcpy, *hipMemcpyToArray*, *hipMemcpy2DToArray*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

- **dst** – [in] Destination memory address
- **dpitch** – [in] Pitch of destination memory
- **src** – [in] Source memory address
- **spitch** – [in] Pitch of source memory
- **width** – [in] Width of matrix transfer (columns in bytes)
- **height** – [in] Height of matrix transfer (rows)
- **kind** – [in] Type of transfer
- **stream** – [in] Stream to use

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpy2DToArray**(hipArray_t dst, size_t wOffset, size_t hOffset, const void *src, size_t spitch, size_t width, size_t height, hipMemcpyKind kind)

Copies data between host and device.

See also:

hipMemcpy, *hipMemcpyToArray*, *hipMemcpy2D*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

- **dst** – [in] Destination memory address
- **wOffset** – [in] Destination starting X offset
- **hOffset** – [in] Destination starting Y offset
- **src** – [in] Source memory address
- **spitch** – [in] Pitch of source memory
- **width** – [in] Width of matrix transfer (columns in bytes)
- **height** – [in] Height of matrix transfer (rows)
- **kind** – [in] Type of transfer

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpy2DToArrayAsync**(hipArray_t dst, size_t wOffset, size_t hOffset, const void *src, size_t spitch, size_t width, size_t height, hipMemcpyKind kind, hipStream_t stream)

Copies data between host and device.

See also:

hipMemcpy, *hipMemcpyToArray*, *hipMemcpy2D*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

- **dst** – [in] Destination memory address
- **wOffset** – [in] Destination starting X offset
- **hOffset** – [in] Destination starting Y offset
- **src** – [in] Source memory address
- **spitch** – [in] Pitch of source memory
- **width** – [in] Width of matrix transfer (columns in bytes)
- **height** – [in] Height of matrix transfer (rows)
- **kind** – [in] Type of transfer
- **stream** – [in] Accelerator view which the copy is being enqueued

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpy2DArrayToArray**(hipArray_t dst, size_t wOffsetDst, size_t hOffsetDst, hipArray_const_t src, size_t wOffsetSrc, size_t hOffsetSrc, size_t width, size_t height, hipMemcpyKind kind)

Copies data between host and device.

See also:

hipMemcpy, *hipMemcpyToArray*, *hipMemcpy2D*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

- **dst** – [in] Destination memory address
- **wOffsetDst** – [in] Destination starting X offset
- **hOffsetDst** – [in] Destination starting Y offset
- **src** – [in] Source memory address
- **wOffsetSrc** – [in] Source starting X offset
- **hOffsetSrc** – [in] Source starting Y offset (columns in bytes)
- **width** – [in] Width of matrix transfer (columns in bytes)
- **height** – [in] Height of matrix transfer (rows)

- **kind** – [in] Type of transfer

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpy2DFromArray**(void *dst, size_t dpitch, hipArray_const_t src, size_t wOffset, size_t hOffset, size_t width, size_t height, hipMemcpyKind kind)

Copies data between host and device.

See also:

hipMemcpy, *hipMemcpy2DToArray*, *hipMemcpy2D*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

- **dst** – [in] Destination memory address
- **dpitch** – [in] Pitch of destination memory
- **src** – [in] Source memory address
- **wOffset** – [in] Source starting X offset
- **hOffset** – [in] Source starting Y offset
- **width** – [in] Width of matrix transfer (columns in bytes)
- **height** – [in] Height of matrix transfer (rows)
- **kind** – [in] Type of transfer

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpy2DFromArrayAsync**(void *dst, size_t dpitch, hipArray_const_t src, size_t wOffset, size_t hOffset, size_t width, size_t height, hipMemcpyKind kind, hipStream_t stream)

Copies data between host and device asynchronously.

See also:

hipMemcpy, *hipMemcpy2DToArray*, *hipMemcpy2D*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

- **dst** – [in] Destination memory address
- **dpitch** – [in] Pitch of destination memory
- **src** – [in] Source memory address
- **wOffset** – [in] Source starting X offset
- **hOffset** – [in] Source starting Y offset
- **width** – [in] Width of matrix transfer (columns in bytes)
- **height** – [in] Height of matrix transfer (rows)

- **kind** – [in] Type of transfer
- **stream** – [in] Accelerator view which the copy is being enqueued

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpyAtoH**(void *dst, hipArray_t srcArray, size_t srcOffset, size_t count)

Copies data between host and device.

See also:

hipMemcpy, hipMemcpy2DToArray, hipMemcpy2D, hipMemcpyFromArray, hipMemcpyToSymbol, hipMemcpyAsync

Parameters

- **dst** – [in] Destination memory address
- **srcArray** – [in] Source array
- **srcOffset** – [in] Offset in bytes of source array
- **count** – [in] Size of memory copy in bytes

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpyHtoA**(hipArray_t dstArray, size_t dstOffset, const void *srcHost, size_t count)

Copies data between host and device.

See also:

hipMemcpy, hipMemcpy2DToArray, hipMemcpy2D, hipMemcpyFromArray, hipMemcpyToSymbol, hipMemcpyAsync

Parameters

- **dstArray** – [in] Destination memory address
- **dstOffset** – [in] Offset in bytes of destination array
- **srcHost** – [in] Source host pointer
- **count** – [in] Size of memory copy in bytes

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpy3D**(const struct hipMemcpy3DParms *p)

Copies data between host and device.

See also:

hipMemcpy, *hipMemcpy2DToArray*, *hipMemcpy2D*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

p – [in] 3D memory copy parameters

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipMemcpy3DAsync**(const struct hipMemcpy3DParms *p, hipStream_t stream)

Copies data between host and device asynchronously.

See also:

hipMemcpy, *hipMemcpy2DToArray*, *hipMemcpy2D*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

- **p** – [in] 3D memory copy parameters
- **stream** – [in] Stream to use

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipDrvMemcpy3D**(const HIP_MEMCPY3D *pCopy)

Copies data between host and device.

See also:

hipMemcpy, *hipMemcpy2DToArray*, *hipMemcpy2D*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

pCopy – [in] 3D memory copy parameters

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

hipError_t **hipDrvMemcpy3DAsync**(const HIP_MEMCPY3D *pCopy, hipStream_t stream)

Copies data between host and device asynchronously.

See also:

hipMemcpy, *hipMemcpy2DToArray*, *hipMemcpy2D*, *hipMemcpyFromArray*, *hipMemcpyToSymbol*, *hipMemcpyAsync*

Parameters

- **pCopy** – [in] 3D memory copy parameters
- **stream** – [in] Stream to use

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidPitchValue, hipErrorInvalidDevicePointer, hipErrorInvalidMemcpyDirection

17.23.4 External Resource Interoperability

hipError_t **hipImportExternalSemaphore**(hipExternalSemaphore_t *extSem_out, const
hipExternalSemaphoreHandleDesc *semHandleDesc)

Imports an external semaphore.

See also:

Note: This API is currently not supported on Linux.

Parameters

- **extSem_out** – [out] External semaphores to be waited on
- **semHandleDesc** – [in] Semaphore import handle descriptor

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipSignalExternalSemaphoresAsync**(const hipExternalSemaphore_t *extSemArray, const
hipExternalSemaphoreSignalParams *paramsArray, unsigned
int numExtSems, hipStream_t stream)

Signals a set of external semaphore objects.

See also:

Note: This API is currently not supported on Linux.

Parameters

- **extSemArray** – [in] External semaphores to be waited on
- **paramsArray** – [in] Array of semaphore parameters
- **numExtSems** – [in] Number of semaphores to wait on
- **stream** – [in] Stream to enqueue the wait operations in

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipWaitExternalSemaphoresAsync**(const hipExternalSemaphore_t *extSemArray, const
hipExternalSemaphoreWaitParams *paramsArray, unsigned int
numExtSems, hipStream_t stream)

Waits on a set of external semaphore objects.

See also:

Note: This API is currently not supported on Linux.

Parameters

- **extSemArray** – [in] External semaphores to be waited on
- **paramsArray** – [in] Array of semaphore parameters
- **numExtSems** – [in] Number of semaphores to wait on
- **stream** – [in] Stream to enqueue the wait operations in

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipDestroyExternalSemaphore**(hipExternalSemaphore_t extSem)

Destroys an external semaphore object and releases any references to the underlying resource. Any outstanding signals or waits must have completed before the semaphore is destroyed.

See also:

Note: This API is currently not supported on Linux.

Parameters

extSem – [in] handle to an external memory object

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipImportExternalMemory**(hipExternalMemory_t *extMem_out, const
hipExternalMemoryHandleDesc *memHandleDesc)

Imports an external memory object.

See also:

Parameters

- **extMem_out** – [out] Returned handle to an external memory object
- **memHandleDesc** – [in] Memory import handle descriptor

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipExternalMemoryGetMappedBuffer**(void **devPtr, hipExternalMemory_t extMem, const
hipExternalMemoryBufferDesc *bufferDesc)

Maps a buffer onto an imported memory object.

See also:

Parameters

- **devPtr** – [out] Returned device pointer to buffer
- **extMem** – [in] Handle to external memory object
- **bufferDesc** – [in] Buffer descriptor

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipDestroyExternalMemory**(hipExternalMemory_t extMem)

Destroys an external memory object.

See also:

Parameters

extMem – [in] External memory object to be destroyed

Returns

hipSuccess, hipErrorInvalidDevice, hipErrorInvalidValue

hipError_t **hipExternalMemoryGetMappedMipmappedArray**(hipMipmappedArray_t *mipmap,
hipExternalMemory_t extMem, const
hipExternalMemoryMipmappedArrayDesc
*mipmapDesc)

Maps a mipmapped array onto an external memory object.

Returned mipmapped array must be freed using `hipFreeMipmappedArray`.

See also:

hipImportExternalMemory, *hipDestroyExternalMemory*, *hipExternalMemoryGetMappedBuffer*,
hipFreeMipmappedArray

Parameters

- **mipmap** – [out] mipmapped array to return
- **extMem** – [in] external memory object handle
- **mipmapDesc** – [in] external mipmapped array descriptor

Returns

hipSuccess, hipErrorInvalidValue, hipErrorInvalidResourceHandle

17.24 Register Keyword

The register keyword is deprecated in C++, and is silently ignored by both NVCC and HIP-Clang. You can pass the option `-Wdeprecated-register` the compiler warning message.

17.25 Pragma Unroll

Unroll with a bounds that is known at compile-time is supported. For example:

```
#pragma unroll 16 /* hint to compiler to unroll next loop by 16 */
for (int i=0; i<16; i++) ...
```

```
#pragma unroll 1 /* tell compiler to never unroll the loop */
for (int i=0; i<16; i++) ...
```

```
#pragma unroll /* hint to compiler to completely unroll next loop. */
for (int i=0; i<16; i++) ...
```

17.26 In-Line Assembly

GCN ISA In-line assembly is supported.

There are some usage limitations in ROCm compiler for inline asm support, please refer to [Inline ASM statements](#) for details.

Users can get related background resources on [how to use inline assembly](#) for any usage of inline assembly features.

A short example program including an inline assembly statement can be found at [inline asm tutorial](#).

For further usage of special AMD GPU hardware features that are available through assembly, please refer to the ISA manual for [AMDGPU usage](#), in which AMD GCN is listed from gfx906 to RDNA 3.5.

17.27 C++ Support

The following C++ features are not supported:

- Run-time-type information (RTTI)
- Try/catch

Partially supported features:

- Virtual functions

Virtual functions are not supported if objects containing virtual function tables are passed between GPU's of different offload arch's, e.g. between gfx906 and gfx1030. Otherwise virtual functions are supported.

17.28 Kernel Compilation

hipcc now supports compiling C++/HIP kernels to binary code objects. The file format for binary is .co which means Code Object. The following command builds the code object using hipcc.

```
hipcc --genco --offload-arch=[TARGET GPU] [INPUT FILE] -o [OUTPUT FILE]
```

[TARGET GPU] = GPU architecture

[INPUT FILE] = Name of the file containing kernels

[OUTPUT FILE] = Name of the generated code object file

Note: When using binary code objects is that the number of arguments to the kernel is different on HIP-Clang and NVCC path. Refer to the [HIP module_api sample](#) for differences in the arguments to be passed to the kernel.

17.29 gfx-arch-specific-kernel

Clang defined ‘__gfx*__’ macros can be used to execute gfx arch specific codes inside the kernel. Refer to the sample in [HIP 14_gpu_arch sample](#).

TABLE COMPARING SYNTAX FOR DIFFERENT COMPUTE APIS

Term	CUDA	HIP	OpenCL
Device	int deviceId	int deviceId	cl_device
Queue	cudaStream_t	hipStream_t	cl_command_queue
Event	cudaEvent_t	hipEvent_t	cl_event
Memory	void *	void *	cl_mem
	grid	grid	NDRange
	block	block	work-group
	thread	thread	work-item
	warp	warp	sub-group
Thread-index	threadIdx.x	threadIdx.x	get_local_id(0)
Block-index	blockIdx.x	blockIdx.x	get_group_id(0)
Block-dim	blockDim.x	blockDim.x	get_local_size(0)
Grid-dim	gridDim.x	gridDim.x	get_num_groups(0)
Device Kernel	__global__	__global__	__kernel
Device Function	__device__	__device__	Implied in device compil
Host Function	__host__ (default)	__host__ (default)	Implied in host compil
Host + Device Function	__host__ __device__	__host__ __device__	No equivalent
Kernel Launch	<<< >>>	hipLaunchKernel/hipLaunchKernelGGL/<<< >>>	clEnqueueNDRangeK
Global Memory	__global__	__global__	__global
Group Memory	__shared__	__shared__	__local
Constant	__constant__	__constant__	__constant
	__syncthreads	__syncthreads	barrier(CLK_LOCAL
Atomic Builtins	atomicAdd	atomicAdd	atomic_add
Precise Math	cos(f)	cos(f)	cos(f)
Fast Math	__cos(f)	__cos(f)	native_cos(f)
Vector	float4	float4	float4

18.1 Notes

The indexing functions (starting with `thread-index`) show the terminology for a 1D grid. Some APIs use reverse order of `xyz / 012` indexing for 3D grids.

HSA RUNTIME API FOR ROCM

The following functions are located in the <https://github.com/ROCm/ROCR-Runtime> repository.

`hsa_status_t hsa_amd_vmem_address_reserve(void **va, size_t size, uint64_t address, uint64_t flags)`

Allocate a reserved address range.

Reserve a virtual address range. The size must be a multiple of the system page size. If it is not possible to allocate the address specified by `address`, then `va` will be a different address range. Address range should be released by calling `hsa_amd_vmem_address_free`.

Note that this API will be deprecated in a future release and replaced by `hsa_amd_vmem_address_reserve_align`

Parameters

- **va** – [out] virtual address allocated
- **size** – [in] of address range requested
- **address** – [in] requested
- **flags** – [in] currently unsupported

Return values

- `::HSA_STATUS_SUCCESS` – Address range allocated successfully
- `::HSA_STATUS_ERROR_NOT_INITIALIZED` – The HSA runtime has not been initialized.
- `::HSA_STATUS_ERROR_OUT_OF_RESOURCES` – Insufficient resources to allocate an address range of this size.

`hsa_status_t hsa_amd_vmem_address_free(void *va, size_t size)`

Free a reserved address range.

Free a previously allocated address range. The size must match the size of a previously allocated address range.

Parameters

- **va** – [out] virtual address to be freed
- **size** – [in] of address range

Return values

- `::HSA_STATUS_SUCCESS` – Address range released successfully
- `::HSA_STATUS_ERROR_INVALID_ALLOCATION` – Invalid `va` specified
- `::HSA_STATUS_ERROR_INVALID_ARGUMENT` – Invalid size specified
- `::HSA_STATUS_ERROR_RESOURCE_FREE` – Address range is still in use

- `::HSA_STATUS_ERROR` – Internal unexpected error

`hsa_status_t hsa_amd_vmem_handle_create`(`hsa_amd_memory_pool_t` pool, `size_t` size, `hsa_amd_memory_type_t` type, `uint64_t` flags, `hsa_amd_vmem_alloc_handle_t *`memory_handle)

Create a virtual memory handle.

Create a virtual memory handle within this pool size must be aligned to allocation granule size for this memory pool, see `HSA_AMD_MEMORY_POOL_INFO_RUNTIME_ALLOC_GRANULE` To minimize internal memory fragmentation, align the size to the recommended allocation granule size, see `HSA_AMD_MEMORY_POOL_INFO_RUNTIME_ALLOC_REC_GRANULE`

Parameters

- **pool** – [in] memory to use
- **size** – [in] of the memory allocation
- **type** – [in] of memory
- **flags** – [in] - currently unsupported
- **memory_handle** – [out] - handle for the allocation

Return values

- `::HSA_STATUS_SUCCESS` – memory allocated successfully
- `::HSA_STATUS_ERROR_NOT_INITIALIZED` – The HSA runtime has not been initialized.
- `::HSA_STATUS_ERROR_INVALID_ARGUMENT` – Invalid arguments
- `::HSA_STATUS_ERROR_INVALID_ALLOCATION` – This memory pool does not support allocations
- `::HSA_STATUS_ERROR_OUT_OF_RESOURCES` – Insufficient resources to allocate this memory

`hsa_status_t hsa_amd_vmem_handle_release`(`hsa_amd_vmem_alloc_handle_t` memory_handle)

Release a virtual memory handle.

Parameters

memory – [in] handle that was previously allocated

Return values

- `::HSA_STATUS_SUCCESS` – Address range allocated successfully
- `::HSA_STATUS_ERROR_INVALID_ALLOCATION` – Invalid memory handle

`hsa_status_t hsa_amd_vmem_map`(`void *`va, `size_t` size, `size_t` in_offset, `hsa_amd_vmem_alloc_handle_t` memory_handle, `uint64_t` flags)

Map a virtual memory handle.

Map a virtual memory handle to a reserved address range. The virtual address requested must be within a previously reserved address range. va and (va + size) must be within (va + size) of the previous allocated address range. size must be equal to size of the memory_handle hsa_amd_vmem_set_access needs to be called to make the memory accessible to specific agents

Parameters

- **va** – [in] virtual address range where memory will be mapped
- **size** – [in] of memory mapping
- **in_offset** – [in] offset into memory. Currently unsupported

- **memory_handle** – [in] virtual memory handle to be mapped
- **flags.** – [in] Currently unsupported

Return values

- **::HSA_STATUS_SUCCESS** – Memory mapped successfully
- **::HSA_STATUS_ERROR_INVALID_ARGUMENT** – va, size or memory_handle are invalid
- **::HSA_STATUS_ERROR_OUT_OF_RESOURCES** – Insufficient resources
- **::HSA_STATUS_ERROR** – Unexpected internal error

hsa_status_t **hsa_amd_vmem_unmap**(void *va, size_t size)

Unmap a virtual memory handle.

Unmap previously mapped virtual address range

Parameters

- **va** – [in] virtual address range where memory will be mapped
- **size** – [in] of memory mapping

Return values

- **::HSA_STATUS_SUCCESS** – Memory backing unmapped successfully
- **::HSA_STATUS_ERROR_INVALID_ALLOCATION** – memory_handle is invalid
- **::HSA_STATUS_ERROR_INVALID_ARGUMENT** – size is invalid
- **::HSA_STATUS_ERROR** – Unexpected internal error

hsa_status_t **hsa_amd_vmem_set_access**(void *va, size_t size, const hsa_amd_memory_access_desc_t *desc, size_t desc_cnt)

Make a memory mapping accessible.

Make previously mapped virtual address accessible to specific agents. **size** must be equal to size of previously mapped virtual memory handle. Calling **hsa_amd_vmem_set_access** multiple times on the same **va** will overwrite previous permissions for all agents

Parameters

- **va** – [in] previously mapped virtual address
- **size** – [in] of memory mapping
- **desc** – [in] list of access permissions for each agent
- **desc_cnt** – [in] number of elements in desc

Return values

- **::HSA_STATUS_SUCCESS** –
- **::HSA_STATUS_ERROR_INVALID_ARGUMENT** – va, size or memory_handle are invalid
- **::HSA_STATUS_ERROR_INVALID_ALLOCATION** – memory_handle is invalid
- **::HSA_STATUS_ERROR_OUT_OF_RESOURCES** – Insufficient resources
- **::HSA_STATUS_ERROR_INVALID_AGENT** – Invalid agent in desc
- **::HSA_STATUS_ERROR** – Unexpected internal error

`hsa_status_t hsa_amd_vmem_get_access(void *va, hsa_access_permission_t *perms, hsa_agent_t agent_handle)`

Get current access permissions for memory mapping.

Get access permissions for memory mapping for specific agent.

Parameters

- **va** – [in] previously mapped virtual address
- **perms** – [in] current permissions
- **agent_handle** – [in] agent

Return values

- `::HSA_STATUS_SUCCESS` –
- `::HSA_STATUS_ERROR_INVALID_AGENT` – Invalid agent
- `::HSA_STATUS_ERROR_INVALID_ALLOCATION` – va is not mapped or permissions never set for this agent
- `::HSA_STATUS_ERROR` – Unexpected internal error

`hsa_status_t hsa_amd_vmem_export_shareable_handle(int *dmabuf_fd, hsa_amd_vmem_alloc_handle_t handle, uint64_t flags)`

Get an exportable shareable handle.

Get an exportable shareable handle for a memory_handle. This shareabl handle can then be used to re-create a virtual memory handle using `hsa_amd_vmem_import_shareable_handle`. The shareable handle can be transferred using mechanisms that support posix file descriptors Once all shareable handles are closed, the memory_handle is released.

Parameters

- **dmabuf_fd** – [out] shareable handle
- **handle** – [in] previously allocated virtual memory handle
- **flags** – [in] Currently unsupported

Return values

- `::HSA_STATUS_SUCCESS` –
- `::HSA_STATUS_ERROR_INVALID_ALLOCATION` – Invalid memory handle
- `::HSA_STATUS_ERROR_OUT_OF_RESOURCES` – Out of resources
- `::HSA_STATUS_ERROR` – Unexpected internal error

`hsa_status_t hsa_amd_vmem_import_shareable_handle(int dmabuf_fd, hsa_amd_vmem_alloc_handle_t *handle)`

Import a shareable handle.

Import a shareable handle for a memory handle. Importing a shareable handle that has been closed and released results in undefined behavior.

Parameters

- **dmabuf_fd** – [in] shareable handle exported with `hsa_amd_vmem_export_shareable_handle`
- **handle** – [out] virtual memory handle

Return values

- `::HSA_STATUS_SUCCESS` –

- `::HSA_STATUS_ERROR_INVALID_ALLOCATION` – Invalid memory handle
- `::HSA_STATUS_ERROR_OUT_OF_RESOURCES` – Out of resources
- `::HSA_STATUS_ERROR` – Unexpected internal error

`hsa_status_t hsa_amd_vmem_retain_alloc_handle(hsa_amd_vmem_alloc_handle_t *memory_handle, void *addr)`

Returns memory handle for mapped memory.

Return a memory handle for previously mapped memory. The handle will be the same value of handle used to map the memory. The returned handle must be released with corresponding number of calls to `hsa_amd_vmem_handle_release`.

Parameters

- **memory_handle** – [out] memory handle for this mapped address
- **mapped** – [in] address

Return values

- `::HSA_STATUS_SUCCESS` –
- `::HSA_STATUS_ERROR_INVALID_ALLOCATION` – Invalid address

`hsa_status_t hsa_amd_vmem_get_alloc_properties_from_handle(hsa_amd_vmem_alloc_handle_t memory_handle, hsa_amd_memory_pool_t *pool, hsa_amd_memory_type_t *type)`

Returns the current allocation properties of a handle.

Returns the allocation properties of an existing handle

Parameters

- **memory_handle** – [in] memory handle to be queried
- **pool** – [out] memory pool that owns this handle
- **memory** – [out] type

Return values

- `::HSA_STATUS_SUCCESS` –
- `::HSA_STATUS_ERROR_INVALID_ALLOCATION` – Invalid memory_handle

HIP DEPRECATED RUNTIME API FUNCTIONS

Several of our API functions have been flagged for deprecation. Using the following functions results in errors and unexpected results, so we encourage you to update your code accordingly.

20.1 Context management

CUDA supports cuCtx API, which is the driver API that defines “Context” and “Devices” as separate entities. Context contains a single device, and a device can theoretically have multiple contexts. HIP initially added limited support for these APIs in order to facilitate porting from existing driver codes. These APIs are now marked as deprecated because there are better alternate interfaces (such as `hipSetDevice` or the stream API) to achieve these functions.

- `hipCtxCreate`
- `hipCtxDestroy`
- `hipCtxPopCurrent`
- `hipCtxPushCurrent`
- `hipCtxSetCurrent`
- `hipCtxGetCurrent`
- `hipCtxGetDevice`
- `hipCtxGetApiVersion`
- `hipCtxGetCacheConfig`
- `hipCtxSetCacheConfig`
- `hipCtxSetSharedMemConfig`
- `hipCtxGetSharedMemConfig`
- `hipCtxSynchronize`
- `hipCtxGetFlags`
- `hipCtxEnablePeerAccess`
- `hipCtxDisablePeerAccess`
- `hipDevicePrimaryCtxGetState`
- `hipDevicePrimaryCtxRelease`
- `hipDevicePrimaryCtxRetain`
- `hipDevicePrimaryCtxReset`

- `hipDevicePrimaryCtxSetFlags`

20.2 Memory management

- `hipMallocHost` (replaced with `hipHostMalloc`)
- `hipMemAllocHost` (replaced with `hipHostMalloc`)
- `hipMemcpyToArray`
- `hipMemcpyFromArray`

20.3 Profiler control

- `hipProfilerStart` (use `roctracer/rocTX`)
- `hipProfilerStop` (use `roctracer/rocTX`)

20.4 Texture management

- `hipGetTextureReference`
- `hipTexRefSetAddressMode`
- `hipTexRefSetArray`
- `hipTexRefSetFilterMode`
- `hipTexRefSetFlags`
- `hipTexRefSetFormat`
- `hipTexRefGetAddress`
- `hipTexRefGetAddressMode`
- `hipTexRefGetFilterMode`
- `hipTexRefGetFlags`
- `hipTexRefGetFormat`
- `hipTexRefGetMaxAnisotropy`
- `hipTexRefGetMipmapFilterMode`
- `hipTexRefGetMipmapLevelBias`
- `hipTexRefGetMipmapLevelClamp`
- `hipTexRefGetMipMappedArray`
- `hipTexRefSetAddress`
- `hipTexRefSetAddress2D`
- `hipTexRefSetMaxAnisotropy`
- `hipTexRefSetBorderColor`
- `hipTexRefSetMipmapFilterMode`

- `hipTexRefSetMipmapLevelBias`
- `hipTexRefSetMipmapLevelClamp`
- `hipTexRefSetMipmappedArray`
- `hipTexRefGetBorderColor`
- `hipTexRefGetArray`
- `hipBindTexture`
- `hipBindTexture2D`
- `hipBindTextureToArray`
- `hipGetTextureAlignmentOffset`
- `hipUnbindTexture`
- `hipBindTextureToMipmappedArray`

FP8 NUMBERS

FP8 numbers were introduced to accelerate deep learning inferencing. They provide higher throughput of matrix operations because the smaller size allows more of them in the available fixed memory.

HIP has two FP8 number representations called *FP8-OCP* and *FP8-FNUZ*.

Open Compute Project(OCP) number definition can be found [here](#).

Definition of FNUZ: fnuz suffix means only finite and NaN values are supported. Unlike other types, Inf are not supported. NaN is when sign bit is set and all other exponent and mantissa bits are 0. All other values are finite. This provides one extra value of exponent and adds to the range of supported FP8 numbers.

21.1 FP8 Definition

FP8 numbers are composed of a sign, an exponent and a mantissa. Their sizes are dependent on the format. There are two formats of FP8 numbers, E4M3 and E5M2.

- E4M3: 1 bit sign, 4 bit exponent, 3 bit mantissa
- E5M2: 1 bit sign, 5 bit exponent, 2 bit mantissa

21.2 HIP Header

HIP header defined the FP8 ocp/fnuz numbers [here](#).

21.3 Supported Devices

Table 21.1: Supported devices for fp8 numbers

Device Type	FNUZ FP8	OCP FP8
Host	Yes	Yes
gfx940/gfx941/gfx942	Yes	No
gfx1200/gfx1201	No	Yes

21.4 Usage

To use the FP8 numbers inside HIP programs.

```
#include <hip/hip_fp8.h>
```

FP8 numbers can be used on CPU side:

```
__hip_fp8_storage_t convert_float_to_fp8(
    float in, /* Input val */
    __hip_fp8_interpretation_t interpret, /* interpretation of number E4M3/E5M2 */
    __hip_saturation_t sat /* Saturation behavior */
) {
    return __hip_cvt_float_to_fp8(in, sat, interpret);
}
```

The same can be done in kernels as well.

```
__device__ __hip_fp8_storage_t d_convert_float_to_fp8(
    float in,
    __hip_fp8_interpretation_t interpret,
    __hip_saturation_t sat) {
    return __hip_cvt_float_to_fp8(in, sat, interpret);
}
```

An important thing to note here is if you use this on MI300 GPU, it will be fnuz number but on any other GPU it will be an OCP number.

The following code example does roundtrip FP8 conversions on both the CPU and GPU and compares the results.

```
#include <hip/hip_fp8.h>
#include <hip/hip_runtime.h>
#include <iostream>
#include <vector>

#define hip_check(hip_call) \
{ \
    auto hip_res = hip_call; \
    if (hip_res != hipSuccess) { \
        std::cerr << "Failed in hip call: " << #hip_call \
        << " with error: " << hipGetErrorName(hip_res) << std::endl; \
        std::abort(); \
    } \
}

__device__ __hip_fp8_storage_t d_convert_float_to_fp8(
    float in, __hip_fp8_interpretation_t interpret, __hip_saturation_t sat) {
    return __hip_cvt_float_to_fp8(in, sat, interpret);
}

__device__ float d_convert_fp8_to_float(float in,
    __hip_fp8_interpretation_t interpret) {
    __half hf = __hip_cvt_fp8_to_halfraw(in, interpret);
    return hf;
}
```

(continues on next page)

(continued from previous page)

```

}

__global__ void float_to_fp8_to_float(float *in,
                                     __hip_fp8_interpretation_t interpret,
                                     __hip_saturation_t sat, float *out,
                                     size_t size) {
    int i = threadIdx.x;
    if (i < size) {
        auto fp8 = d_convert_float_to_fp8(in[i], interpret, sat);
        out[i] = d_convert_fp8_to_float(fp8, interpret);
    }
}

__hip_fp8_storage_t
convert_float_to_fp8(float in, /* Input val */
                    __hip_fp8_interpretation_t
                    interpret, /* interpretation of number E4M3/E5M2 */
                    __hip_saturation_t sat /* Saturation behavior */) {
    return __hip_cvt_float_to_fp8(in, sat, interpret);
}

float convert_fp8_to_float(
    __hip_fp8_storage_t in, /* Input val */
    __hip_fp8_interpretation_t
    interpret /* interpretation of number E4M3/E5M2 */) {
    __half hf = __hip_cvt_fp8_to_halfraw(in, interpret);
    return hf;
}

int main() {
    constexpr size_t size = 32;
    hipDeviceProp_t prop;
    hip_check(hipGetDeviceProperties(&prop, 0));
    bool is_supported = (std::string(prop.gcnArchName).find("gfx94") !=
↳ std::string::npos) || /* gfx94x
                        (std::string(prop.gcnArchName).find("gfx120") !=
↳ std::string::npos); /* gfx120x
    if(!is_supported) {
        std::cerr << "Need a gfx94x or gfx120x, but found: " << prop.gcnArchName <<
↳ std::endl;
        std::cerr << "No device conversions are supported, only host conversions are
↳ supported." << std::endl;
        return -1;
    }

    const __hip_fp8_interpretation_t interpret = (std::string(prop.gcnArchName).find(
↳ "gfx94") != std::string::npos)
                                                ? __HIP_E4M3_FNUZ /* gfx94x
                                                : __HIP_E4M3;      /* gfx120x

    constexpr __hip_saturation_t sat = __HIP_SATFINITE;

```

(continues on next page)

(continued from previous page)

```

std::vector<float> in;
in.reserve(size);
for (size_t i = 0; i < size; i++) {
    in.push_back(i + 1.1f);
}

std::cout << "Converting float to fp8 and back..." << std::endl;
// CPU convert
std::vector<float> cpu_out;
cpu_out.reserve(size);
for (const auto &fval : in) {
    auto fp8 = convert_float_to_fp8(fval, interpret, sat);
    cpu_out.push_back(convert_fp8_to_float(fp8, interpret));
}

// GPU convert
float *d_in, *d_out;
hip_check(hipMalloc(&d_in, sizeof(float) * size));
hip_check(hipMalloc(&d_out, sizeof(float) * size));

hip_check(hipMemcpy(d_in, in.data(), sizeof(float) * in.size(),
                    hipMemcpyHostToDevice));

float_to_fp8_to_float<<<1, size>>>(d_in, interpret, sat, d_out, size);

std::vector<float> gpu_out(size, 0.0f);
hip_check(hipMemcpy(gpu_out.data(), d_out, sizeof(float) * gpu_out.size(),
                    hipMemcpyDeviceToHost));

hip_check(hipFree(d_in));
hip_check(hipFree(d_out));

// Validation
for (size_t i = 0; i < size; i++) {
    if (cpu_out[i] != gpu_out[i]) {
        std::cerr << "cpu round trip result: " << cpu_out[i]
                   << " - gpu round trip result: " << gpu_out[i] << std::endl;
        std::abort();
    }
}
std::cout << "...CPU and GPU round trip convert matches." << std::endl;
}

```

There are C++ style classes available as well.

```

__hip_fp8_e4m3_fnuz fp8_val(1.1f); // gfx94x
__hip_fp8_e4m3      fp8_val(1.1f); // gfx120x

```

Each type of FP8 number has its own class:

- `__hip_fp8_e4m3`
- `__hip_fp8_e5m2`

- `__hip_fp8_e4m3_fnuz`
- `__hip_fp8_e5m2_fnuz`

There is support of vector of FP8 types.

- `__hip_fp8x2_e4m3`: holds 2 values of OCP FP8 e4m3 numbers
- `__hip_fp8x4_e4m3`: holds 4 values of OCP FP8 e4m3 numbers
- `__hip_fp8x2_e5m2`: holds 2 values of OCP FP8 e5m2 numbers
- `__hip_fp8x4_e5m2`: holds 4 values of OCP FP8 e5m2 numbers
- `__hip_fp8x2_e4m3_fnuz`: holds 2 values of FP8 fnuz e4m3 numbers
- `__hip_fp8x4_e4m3_fnuz`: holds 4 values of FP8 fnuz e4m3 numbers
- `__hip_fp8x2_e5m2_fnuz`: holds 2 values of FP8 fnuz e5m2 numbers
- `__hip_fp8x4_e5m2_fnuz`: holds 4 values of FP8 fnuz e5m2 numbers

FNUZ extensions will be available on MI300 only.

LICENSE

Copyright © 2008 - 2024 Advanced Micro Devices, Inc.

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the “Software”), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED “AS IS”, WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

H

- hipArray3DCreate (C++ function), 181
- hipArray3DGetDescriptor (C++ function), 183
- hipArrayCreate (C++ function), 181
- hipArrayDestroy (C++ function), 181
- hipArrayGetDescriptor (C++ function), 183
- hipArrayGetInfo (C++ function), 182
- hipDestroyExternalMemory (C++ function), 192
- hipDestroyExternalSemaphore (C++ function), 191
- hipDeviceCanAccessPeer (C++ function), 159
- hipDeviceDisablePeerAccess (C++ function), 160
- hipDeviceEnablePeerAccess (C++ function), 159
- hipDeviceGetStreamPriorityRange (C++ function), 145
- hipDrvMemcpy3D (C++ function), 189
- hipDrvMemcpy3DAsync (C++ function), 189
- hipDrvPointerGetAttributes (C++ function), 162
- hipExternalMemoryGetMappedBuffer (C++ function), 192
- hipExternalMemoryGetMappedMipmappedArray (C++ function), 192
- hipExtMallocWithFlags (C++ function), 163
- hipExtStreamCreateWithCUMask (C++ function), 148
- hipExtStreamGetCUMask (C++ function), 149
- hipFree (C++ function), 167
- hipFreeArray (C++ function), 182
- hipFreeAsync (C++ function), 150
- hipFreeHost (C++ function), 167
- hipGetProcAddress (C++ function), 174
- hipGetSymbolAddress (C++ function), 174
- hipGetSymbolSize (C++ function), 174
- hipHostAlloc (C++ function), 164
- hipHostGetDevicePointer (C++ function), 164
- hipHostGetFlags (C++ function), 165
- hipHostMalloc (C++ function), 163
- hipHostRegister (C++ function), 165
- hipHostUnregister (C++ function), 166
- hipImportExternalMemory (C++ function), 191
- hipImportExternalSemaphore (C++ function), 190
- hipMalloc (C++ function), 162
- hipMalloc3D (C++ function), 181
- hipMalloc3DArray (C++ function), 182
- hipMallocArray (C++ function), 180
- hipMallocAsync (C++ function), 149
- hipMallocFromPoolAsync (C++ function), 155
- hipMallocPitch (C++ function), 166
- hipMemAllocPitch (C++ function), 166
- hipMemcpy (C++ function), 167
- hipMemcpy2D (C++ function), 183
- hipMemcpy2DArrayToArray (C++ function), 186
- hipMemcpy2DAsync (C++ function), 184
- hipMemcpy2DFromArray (C++ function), 187
- hipMemcpy2DFromArrayAsync (C++ function), 187
- hipMemcpy2DToArray (C++ function), 185
- hipMemcpy2DToArrayAsync (C++ function), 185
- hipMemcpy3D (C++ function), 188
- hipMemcpy3DAsync (C++ function), 189
- hipMemcpyAsync (C++ function), 176
- hipMemcpyAtoA (C++ function), 171
- hipMemcpyAtoD (C++ function), 170
- hipMemcpyAtoH (C++ function), 188
- hipMemcpyAtoHAsync (C++ function), 172
- hipMemcpyDtoA (C++ function), 170
- hipMemcpyDtoD (C++ function), 169
- hipMemcpyDtoDAsync (C++ function), 172
- hipMemcpyDtoH (C++ function), 169
- hipMemcpyDtoHAsync (C++ function), 172
- hipMemcpyFromSymbol (C++ function), 175
- hipMemcpyFromSymbolAsync (C++ function), 175
- hipMemcpyHtoA (C++ function), 188
- hipMemcpyHtoAAsync (C++ function), 173
- hipMemcpyHtoD (C++ function), 168
- hipMemcpyHtoDAsync (C++ function), 171
- hipMemcpyParam2D (C++ function), 184
- hipMemcpyParam2DAsync (C++ function), 184
- hipMemcpyPeer (C++ function), 160
- hipMemcpyPeerAsync (C++ function), 160
- hipMemcpyToSymbol (C++ function), 174
- hipMemcpyToSymbolAsync (C++ function), 175
- hipMemcpyWithStream (C++ function), 168
- hipMemGetAddressRange (C++ function), 160
- hipMemGetInfo (C++ function), 180
- hipMemPoolCreate (C++ function), 154
- hipMemPoolDestroy (C++ function), 155

`hipMemPoolExportPointer` (C++ function), 158
`hipMemPoolExportToShareableHandle` (C++ function), 156
`hipMemPoolGetAccess` (C++ function), 154
`hipMemPoolGetAttribute` (C++ function), 152
`hipMemPoolImportFromShareableHandle` (C++ function), 157
`hipMemPoolImportPointer` (C++ function), 158
`hipMemPoolSetAccess` (C++ function), 153
`hipMemPoolSetAttribute` (C++ function), 151
`hipMemPoolTrimTo` (C++ function), 151
`hipMemPtrGetInfo` (C++ function), 180
`hipMemset` (C++ function), 177
`hipMemset2D` (C++ function), 179
`hipMemset2DAsync` (C++ function), 179
`hipMemset3D` (C++ function), 179
`hipMemset3DAsync` (C++ function), 179
`hipMemsetAsync` (C++ function), 178
`hipMemsetD16` (C++ function), 177
`hipMemsetD16Async` (C++ function), 177
`hipMemsetD32` (C++ function), 178
`hipMemsetD32Async` (C++ function), 178
`hipMemsetD8` (C++ function), 177
`hipMemsetD8Async` (C++ function), 177
`hipModuleGetGlobal` (C++ function), 173
`hipPointerGetAttribute` (C++ function), 162
`hipPointerGetAttributes` (C++ function), 161
`hipPointerSetAttribute` (C++ function), 161
`hipSignalExternalSemaphoresAsync` (C++ function), 190
`hipStreamAddCallback` (C++ function), 149
`hipStreamCallback_t` (C++ type), 144
`hipStreamCreate` (C++ function), 144
`hipStreamCreateWithFlags` (C++ function), 144
`hipStreamCreateWithPriority` (C++ function), 145
`hipStreamDestroy` (C++ function), 145
`hipStreamGetDevice` (C++ function), 148
`hipStreamGetFlags` (C++ function), 147
`hipStreamGetPriority` (C++ function), 147
`hipStreamQuery` (C++ function), 146
`hipStreamSynchronize` (C++ function), 146
`hipStreamWaitEvent` (C++ function), 147
`hipWaitExternalSemaphoresAsync` (C++ function), 190
`hsa_amd_vmem_address_free` (C++ function), 197
`hsa_amd_vmem_address_reserve` (C++ function), 197
`hsa_amd_vmem_export_shareable_handle` (C++ function), 200
`hsa_amd_vmem_get_access` (C++ function), 199
`hsa_amd_vmem_get_alloc_properties_from_handle` (C++ function), 201
`hsa_amd_vmem_handle_create` (C++ function), 198
`hsa_amd_vmem_handle_release` (C++ function), 198
`hsa_amd_vmem_import_shareable_handle` (C++ function), 200
`hsa_amd_vmem_map` (C++ function), 198
`hsa_amd_vmem_retain_alloc_handle` (C++ function), 201
`hsa_amd_vmem_set_access` (C++ function), 199
`hsa_amd_vmem_unmap` (C++ function), 199

U

`USE_PEER_NON_UNIFIED` (C macro), 161